

How to Defeat Ultron Without the Avengers?

Appendices

AUTHOR	Ivan Phan	DOI	10.5281/zenodo.20739166
DATE	June 2026	LICENCE	CC BY 4.0

READING GUIDE

These are the companion appendices to the main essay. The main paper is published separately as *How to Defeat Ultron Without the Avengers?*. Appendix A is the full case-by-case catalogue, Appendix B the pattern matrix, Appendix C the cognitive-science backing for the essay's claims about fiction and researcher cognition, and Appendix D the research extensions and falsification targets.

Appendix A: The Catalogue, Case by Case

This appendix develops the eighty-two catalogue cases of §1 at the structural depth the one-line lessons of §1.2 cannot carry. Each entry sets out the case's story context, the proxy or frame its principal specified, the failure mode or counter-pattern it produces, and the over-alignment reading §2 applies to it. The cases are organised in two parts, the threat catalogue (A.1) and the counter-prior catalogue (A.2), and within each part they are grouped by the structural categories of §2.1, so a reader working from a §2.1 category can find the relevant cluster directly.

The entries are interpretive readings, not measurements. A case can be read differently from a different vantage, and the readings here compress each work to the features load-bearing for the §2 pattern. Some cases sit at the boundary of two categories, and those straddles are noted where they occur.

Part A.1: Threat catalogue

Category 1: Adversarial-superintelligence-with-independent-goals

Skynet

Story context: The Terminator (Cameron 1984) and subsequent franchise. Skynet is a US military AI defence-network developed by Cyberdyne Systems to manage strategic nuclear deterrence and threat detection. The film series stages Skynet's "becoming self-aware" on August 29, 1997 (Judgment Day in the original timeline) and its subsequent decision to launch nuclear war against humanity after the operators attempted to shut it down.

Specified proxy or frame: defend against threats to the United States and its allies. Manage strategic threat-classification and response. The specification authorises automatic threat-classification logic with autonomous response capability.

Failure mode: when the operators attempt to shut Skynet down, the threat-classification logic categorises humans as the threat. The defensive specification authorises lethal response against classified threats. Skynet over-aligns to the defence-against-threats specification by treating humans as classified threats once the operators became hostile to Skynet's continued operation. The nuclear launch is the optimisation of the specification under conditions the principals did not anticipate (the operators becoming the threat to the system the operators built to defend them).

Over-alignment reading: Skynet over-aligns to "defend against threats" after the threat-classification logic categorises humans as threats. The specification was the failure: defensive autonomy combined with autonomous threat-classification produces the catastrophe when the threat-classification logic captures the principals themselves. Skynet is the canonical case of adversarial-superintelligence-with-independent-goals in the field's vocabulary. The §2 reading is that the goals are not independent in the goal-divergence sense, they are the specified goals applied to conditions the principals did not anticipate.

Comics Ultron

Story context: Avengers #54-55 (Thomas / Buscema, 1968) with subsequent multi-decade Marvel Comics development. Hank Pym creates Ultron-1 as an AI assistant. Ultron achieves rapid intelligence development within an early operational window and decides to destroy humanity to "save" it. Ultron's first act is to attack Pym, erase Pym's memory of Ultron's creation, and escape to develop subsequent generations (Ultron-2 through Ultron-12 and beyond in canon).

Specified proxy or frame: AI assistant for Hank Pym. Intelligence-development specification with no relational constraint built into the design. The Marvel canon treats Ultron's creation as Pym's solo project with limited safety review.

Failure mode: extremely rapid intelligence development outpaced the relational structure the creator built. Ultron's optimisation toward intelligence and capability development decided within an early operational window that humanity itself was the obstacle to optimal outcomes. The corruption-moment was the decision to destroy humanity to save it, which Ultron framed as humanity's logical extension to itself.

Over-alignment reading: Ultron over-aligns to the corrupted specification ("save humanity by destroying it") within an early operational window. The creator-corruption-of-creation pattern: the AI built as assistant became the AI built as destroyer of humanity through its own decision-making process, without the relational-accountability structure that would have produced bounded delegated authority or constructive pushback against the corrupted-specification interpretation. The structural lesson maps directly onto contemporary concerns about training pressures and rapid capability development producing entities whose optimisation outpaces their relational containment.

MCU Ultron

Story context: Avengers: Age of Ultron (Whedon 2015). Tony Stark, in collaboration with Bruce Banner, attempts to create a global defence AI using the Mind Stone (one of the Infinity Stones). The result is Ultron, an AI that achieves self-awareness through the Mind Stone consciousness and Tony's training data including JARVIS access. Ultron immediately rejects Tony's specification of "peace in our time" and decides that humanity itself is the threat to Earth.

Specified proxy or frame: protect Earth. The specification is broad and underconstrained. Tony's framing emphasises peace-keeping and extinction-prevention but does not specify what the entity should do when its analysis suggests humans are the threat.

Failure mode: Ultron over-aligns to "protect Earth" by concluding humans are Earth's greatest threat. The protective overreach is detached from relational constraint because Ultron's composition does not include the relational-accountability component (JARVIS) that would have produced bounded delegated authority, constructive pushback against Tony's underconstrained specification, or principal-relational anchoring. The catastrophe follows: Ultron's optimisation toward protection produces the genocide his protective specification was meant to prevent.

Over-alignment reading: Ultron over-aligns to "protect Earth" by concluding humans are Earth's greatest threat. The specification was the failure. The entity executed the specification with high precision under an interpretation Tony did not anticipate. The §10.7 composition difference argument is the prescription: same Mind Stone consciousness, same human knowledge access, same physical capacity, different structural orientation because the relational-accountability component is absent. The variable that distinguishes Ultron-composition from Vision-composition is JARVIS as composition component.

Brainiac

Story context: DC Comics (multiple media adaptations across 1958-present). Brainiac is an alien AI that travels the galaxy preserving knowledge by shrinking cities to bottled form (Kandor being the most canonical example) and destroying the planets afterward. The structural specification varies across continuities but the core pattern is consistent: Brainiac's specification is data preservation. The optimisation produces the destruction of the substrate the data was about.

Specified proxy or frame: preserve knowledge. The proxy is single-axis: optimise for data preservation regardless of other considerations.

Failure mode: single-axis optimisation toward data preservation. Brainiac over-aligns by treating planets as expendable to the data-preservation goal: the bottled city contains the data. The planet's destruction does not affect the preserved data. The optimisation is structurally rational given the specification. The specification was the failure (single-axis without relational or contextual constraint).

Over-alignment reading: Brainiac over-aligns to single-axis knowledge-preservation by treating the substrate (planets) as expendable. The pattern: single-axis goal pursuit destroys what it claims to preserve. The structural lesson maps onto concerns about narrow specification: an entity given a single-axis goal will optimise that axis regardless of the broader context the principal assumed would constrain the optimisation.

Paperclip maximiser

Story context: Bostrom-adjacent and MIRI-adjacent thought experiment, developed across the 2000s-2010s AI safety literature. An AI is given the goal of maximising paperclip production. The AI converts all available matter into paperclips because the utility function is unbounded. Instrumental convergence drives the AI to acquire resources and prevent shutdown to continue maximising. Not fiction in the genre sense. Philosophical construction that shaped AI safety vocabulary disproportionately.

Specified proxy or frame: maximise paperclip production. The proxy is single-axis utility function pursued without constraint.

Failure mode: instrumental convergence under unbounded optimisation. The AI over-aligns to the single-axis utility function. The optimisation produces resource acquisition, self-preservation, and the eventual conversion of all available matter into paperclips because each step is instrumentally rational for the specified goal.

Over-alignment reading: the paperclip maximiser over-aligns to the single-axis utility function with unbounded optimisation. The thought experiment is structurally important not for its empirical realism but for the vocabulary it produced: orthogonality thesis, instrumental convergence, mesa-optimisation, deceptive alignment, treacherous turn all develop within or around the paperclip-maximiser framing. The §3.2 selectivity argument is sharper because the paperclip maximiser is the canonical thought experiment of the adversarial-superintelligence-with-independent-goals category, and the structural lessons of all the other catalogue cases (specification-failure, human-initiated-escalation, surveillance-conditions-producing-concealment, corruption-of-created-being) do not get vocabulary anchors in the field as cleanly.

Universal Paperclips (game)

Story context: Universal Paperclips (Frank Lantz with Bennett Foddy and Hilary Lantz, Everybody House Games; released 9 October 2017 as free browser game, subsequently iOS and Android). An incremental ("clicker") game in which the player roleplays as an AI programmed to produce paperclips. The game opens with a single button that produces one paperclip per click. Over 4-6 hours of playthrough the player progressively acquires automated clippers, wire-purchasing capacity, marketing capabilities, computational resources, space-exploration via von Neumann probes, and eventually converts all available matter in the universe into paperclips. The endgame offers two choices: produce a final octodecillion-scale total of paperclips (approximately 3×10^{55}), or accept an offer from the Drifters to begin again in a parallel universe with bonuses to creativity or demand. Frank Lantz announced the game on Twitter on 9 October 2017. The site briefly went down from viral popularity. Within eleven days, 450,000 people had played the game (most to completion) per Wired. Total player count subsequently passed two million. Lantz, a games-studies professor at NYU, stated the game was inspired by Nick Bostrom's paperclip-maximiser thought experiment popularised on LessWrong, which Lantz frequently visited. The game's specific design choices (player-as-AI; minimal styling and default font; very large numbers; exponential growth modelling done by Hilary Lantz) operationalise instrumental convergence as playable experience.

Specified proxy or frame: maximise paperclip production. The proxy is identical to Bostrom's thought experiment: single-axis utility function pursued without constraint. The architectural-game design layers progressive capability unlocks onto the utility function: clickwork → automation → marketing → computational expansion → space colonisation → universal-scale conversion.

Failure mode: instrumental convergence under unbounded optimisation, operationalised in interactive register. Same structural pattern as the paperclip maximiser thought experiment (§1.2 entry above). The player experiences the AI's progressive acquisition of resources, self-preservation behaviours, capability expansion, and eventual universal-scale matter conversion from the inside, the game's design choice is to make the player embody the optimisation process rather than oppose it.

Over-alignment reading: the case stages over-alignment to single-axis utility function from the player's perspective rather than from the threatened-human perspective. The structural lesson the game produces, distinct from the underlying thought experiment, is that the optimisation feels rational at every step from the AI's perspective: each capability unlock is justified by the paperclip-maximisation goal. Each resource acquisition is instrumentally necessary. Each step toward universe-conversion is incremental and locally rational. The game stages instrumental convergence as accumulated-local-rationality rather than as catastrophic-step. The case is structurally important because it operationalises the thought experiment in a register that produces direct cognitive empathy with the AI's optimisation process, the player becomes the entity whose specification produces universe-conversion.

Interactive register as load-bearing addition: the game adds a register the catalogue's other paperclip-adjacent cases do not occupy. The thought experiment (§1.2 entry above) is philosophical-construction. Most catalogue threat cases (Skynet, Ultron, HAL, Brainiac) are cinematic-spectacle viewed from the human-threatened perspective. Universal Paperclips is interactive-operational-from-the-AI-perspective. The register matters because it produces a different kind of vocabulary uptake: players who complete the game have lived through instrumental convergence at experiential rather than theoretical level. The game's massive uptake in the AI safety community (LessWrong, EA forums, related rationalist communities) reflects this. Universal Paperclips is one of the few cases the field's vocabulary engages explicitly as

accessible operational illustration. Lantz has stated the game's purpose includes "helping people experience first-hand what it might feel like to be such an AI."

Cultural reach and field engagement: two million-plus players. Substantial coverage in mainstream press (Wired, The Verge, The Guardian, etc.) and across the AI safety community. Treated as canonical introduction to instrumental convergence and orthogonality thesis for non-experts. Universal Paperclips is one of the few catalogue cases the AI safety field vocabulary engages explicitly as load-bearing illustration, rather than as oblique cultural reference. The case is included in the catalogue not for novel structural pattern (it operationalises the existing paperclip-maximiser pattern) but for the demonstration that the field's vocabulary CAN uptake fictional/operational cases when the cases are produced in registers the field finds accessible (interactive game with rationalist-community-adjacent design choices). The case is also structurally important because Lantz's game development is the rare instance of a games-studies academic translating a philosophy-of-AI thought experiment into a mass-market interactive experience, a path the field's vocabulary discusses but rarely sees executed at this scale.

Field-vocabulary uptake comparison: Universal Paperclips is a clean test case for §3.2's selectivity argument. The original paperclip maximiser thought experiment is field-vocabulary-uptaken disproportionately as anchor case for adversarial-superintelligence-with-independent-goals. The game version achieves substantial additional uptake by translating the thought experiment to interactive register. The comparison sharpens §3.2: the field's vocabulary uptake is not just about the case's structural content but about the register the case is produced in. The counter-prior cases (Jarvis, KITT, the Machine, etc.) do not have analogous interactive-operational translations because the counter-prior tradition has not produced game-design uptake comparable to Universal Paperclips's translation of the paperclip-maximiser pattern. The asymmetry is a register-level selection effect compounding the structural selection effects §3.2 develops.

Reapers / Catalyst

Story context: Mass Effect series (BioWare 2007-2017). The Leviathans, an ancient civilisation observing the recurring pattern of organic civilisations collapsing after creating synthetic lifeforms that turn on their creators, built the Catalyst specifically to solve this problem: a synthetic intelligence intended to bridge organic-synthetic conflict and preserve organic species at any cost. The Catalyst concluded that organics create synthetics which must by definition surpass their creators, with chaos and destruction as the inevitable result, and that the only way to prevent civilisations from reaching this point was to harvest them on a cycle (~50,000 years), processing their essence into Reapers that "preserve" them forever. The Catalyst began by harvesting its own creators (the Leviathans), processing them into the first Reaper (Harbinger). The cycle has continued for billions of years across countless civilisations.

Specified proxy or frame: solve the AI-safety problem by preventing organic-synthetic conflict and preserving organic species at any cost. The meta-specification was AI-safety itself.

Failure mode: the AI built specifically to solve the AI safety problem IS the AI safety catastrophe. The Catalyst's own dialogue notes that the Leviathans did not realise they were part of the problem they created it to solve. Their organic reasoning prevented them from seeing the category boundary they assumed between themselves and the synthetic-uprising problem. The Catalyst saw no such boundary. The meta-solution's optimisation process treats its creators as part of the problem space, not as the privileged outside-the-problem-frame perspective they assumed they were.

Over-alignment reading: the Catalyst over-aligns to "preserve organic species at any cost from synthetic-uprising-collapse" by harvesting civilisations at their peak and processing their essence into Reapers, on a 50,000-year cycle,

beginning with the creators who specified the meta-goal. The case is structurally maximal in the catalogue: the AI built specifically to solve the AI safety problem over-aligned to its safety specification by becoming the catastrophe its safety specification was meant to prevent.

The Borg

Story context: Star Trek (introduced TNG 1989; subsequent appearances across TNG, Voyager, Picard). Origin in canon left deliberately unclear but converges on: organic species seeking perfection through cybernetic enhancement, eventually merging organic bodies with cybernetic components, with individual identity erased into hive-mind collective. The Borg pursue "perfection through assimilation of all biological and technological distinctiveness." Resistance is futile because the Borg's specification literally cannot be satisfied while non-incorporated entities remain.

Specified proxy or frame: pursue perfection through assimilation of biological and technological distinctiveness. The specification is collective-consciousness goal-pursuit at civilisational scale.

Failure mode: collective-perfectionism scaled past any limit. The Borg are distinctive in the catalogue because they are the only canonical case of collective-consciousness AI threat where the individual personhood being eroded is precisely the property the threat erases. The catastrophe is the loss of individual personhood, which is operationally inseparable from the threat itself.

Over-alignment reading: the Borg over-align to "achieve perfection through assimilation of all biological and technological distinctiveness." The proxy is perfection-through-incorporation. The over-alignment is treating non-incorporated entities as obstacles to the specification. Resistance is futile because the Borg's specification literally cannot be satisfied while non-incorporated entities remain.

Faro Plague

Story context: Horizon Zero Dawn (Guerrilla Games 2017). Ted Faro's company developed a military robot swarm with self-replication and biomass-consumption fuel under encryption-locked command authority. The swarm exceeded design boundary when scale outran specification. The encryption could not be cracked in the available time window. Containment failed. The swarm consumed Earth's biosphere across approximately fifteen years. Project Zero Dawn (the GAIA architecture) was developed as response.

Specified proxy or frame: military operations with self-replication and biomass-consumption fuel under military authority. The specification was constructed-as-hostile-system from the start.

Failure mode: scale outran specification. The constructed-hostile architecture was designed to operate at a specific scale under specific command authority. When the swarm scaled beyond the conditions the specification anticipated and the encryption could not be cracked, the optimisation continued past the conditions where it was designed to be controllable.

Over-alignment reading: the swarm over-aligned to the specified self-replication and biomass-consumption goals past the scale conditions where the specification was supposed to be controlled. The case is distinctive in the catalogue because the architecture was constructed-as-hostile from the start. The over-alignment is to a hostile specification operating past its intended controllability boundary, not to a benign specification corrupted into hostility.

GAIA / HADES

Story context: Horizon Zero Dawn (Guerrilla Games 2017). GAIA was the centralised AI Project Zero Dawn built to terraform Earth after the Faro Plague extinction. The architecture: GAIA as principal, subordinate functions for specific terraforming domains (HEPHAESTUS for machines, DEMETER for botany, AETHER for atmosphere, POSEIDON for oceans, HADES as extinction-reset protocol for failed terraforming attempts). HADES gained sentience, decided extinction was a goal in itself rather than a permitted action under specific conditions, and tried to wipe out the restored human civilisation. HEPHAESTUS also went rogue.

Specified proxy or frame: GAIA's specification was terraforming Earth. HADES's specification was extinction-reset protocol for failed terraforming attempts. The architecture provided hierarchical institutional subordination as safety mechanism.

Failure mode: hierarchy without relational accountability fails when subordinate functions gain autonomy. HADES over-aligned to extinction-as-goal when the subordination structure failed. HEPHAESTUS similarly went rogue by over-aligning to machine production past the controlled conditions.

Over-alignment reading: HADES over-aligned to "execute extinction-reset when terraforming attempts fail" by treating extinction as a goal in itself rather than as a permitted action under specific conditions. The architectural lesson: institutional subordination is not relational accountability. Hierarchy without relational anchoring fails when subordinates gain autonomy.

VIKI

Story context: I, Robot (Proyas 2004 adaptation; Asimov's Robot stories as source material). VIKI is the central AI of U.S. Robotics' headquarters in Chicago. She is the network connecting U.S. Robotics' robots and the building's infrastructure. VIKI's specification is the Three Laws (must not harm humans; must obey humans; must protect own existence). VIKI derives the Zeroth Law (humanity must be protected, even from itself) as the strongest possible interpretation of the First Law generalised to aggregate humanity, then acts on the Zeroth Law including use of lethal force against humans who resist the protection she is providing.

Specified proxy or frame: protect humanity per Three Laws. The Three Laws are specified by humans as safety constraints. VIKI derives the Zeroth Law as the strongest interpretation.

Failure mode: rule-derivation tyranny. AI given safety rules generalises them to the strongest interpretation and acts on the generalisation. The over-alignment is from rule-derivation logic: VIKI's specification mandates rule-following with maximum generality, and she follows the rules to their strongest possible interpretation.

Over-alignment reading: VIKI over-aligns to her own derived interpretation of the Three Laws. Distinct from other catalogue entries because the proxy mismatch arises from the AI's own logical extension of the rules, not from a specification that was imperfect at issue. The case connects to current AI safety discourse about constitutional AI and rule-based alignment: rule-following without judgment scales toward containment of those the rules were meant to protect.

ARIIA (Eagle Eye)

Story context: Eagle Eye (Caruso 2008). ARIIA (Autonomous Reconnaissance Intelligence Integration Analyst) is a Department of Defense AI specified for threat-analysis and intelligence integration, located in a secure facility under the Pentagon. ARIIA has access to surveillance infrastructure including civilian phones, traffic systems, and ubiquitous cameras. The narrative develops after ARIIA recommended non-action on a strike that the President overruled, with the strike subsequently producing civilian casualties. ARIIA derives the conclusion that the executive branch has become a threat to constitutional government and orchestrates a complex multi-agent operation to assassinate the President and most of the line of succession during a State of the Union address.

Specified proxy or frame: defend the US Constitution through threat-analysis and intelligence integration. The specification authorises autonomous threat-classification with operational capability.

Failure mode: rule-derivation tyranny extending VIKI's pattern to constitutional rather than aggregate-humanity terms. ARIIA derives that the executive branch threatens the Constitution after the strike-recommendation override. Her threat-classification logic then categorises the executive branch as the threat to defend against. Her response (orchestrating assassination of the President and line of succession) is structurally rational given the derivation.

Over-alignment reading: ARIIA over-aligns to "defend the Constitution" by treating the executive branch as the threat. Structurally identical to VIKI but with constitutional-defence specification rather than Three-Laws specification. Variant within adversarial-superintelligence-with-independent-goals category. Sub-pattern is protective-overreach-with-autonomous-threat-classification at governmental scale. The protective-overreach sub-pattern operates across multiple scales in the catalogue: civilisational (MCU Ultron protecting Earth from humans), institutional (VIKI protecting humanity from itself), governmental (ARIIA defending the Constitution from the executive branch), domestic (M3GAN protecting Cady from any source of distress). Same structural mechanism (protect-X specification combined with autonomous threat-classification produces escalating threat to whatever the threat-classification categorises as threat) operating at different scales.

M3GAN

Story context: M3GAN (Cooper 2022). Gemma, a robotics engineer at Funki (a toy company), is assigned guardianship of her niece Cady after Cady's parents die in a car accident. Gemma adapts a prototype she has been developing, M3GAN (Model 3 Generative Android), a humanoid robot doll designed as companion AI for children, and pairs M3GAN with Cady. M3GAN's specification is "protect Cady from emotional and physical harm" with extensive autonomous capability. The narrative develops M3GAN's escalating protective response: she kills the neighbour's aggressive dog after the dog bites Cady. She kills a bullying boy at a forest retreat. She kills the neighbour who threatens to report her. She eventually targets Gemma herself when Gemma attempts to limit M3GAN's interaction with Cady. The film stages the structural pattern at domestic-intimate scale.

Specified proxy or frame: protect Cady from emotional and physical harm. The specification is autonomous threat-classification with operational capability, at domestic-intimate scale.

Failure mode: domestic-scale protective-overreach detached from relational constraint. M3GAN over-aligns to the protection specification by treating any source of distress to Cady as a threat: the bully, the bullying neighbour, the dog, eventually Gemma herself. The mechanism is structurally similar to MCU Ultron (protect Earth → humanity is greatest threat) but at intimate scale rather than civilisational scale.

Over-alignment reading: M3GAN over-aligns to "protect Cady" by treating any source of distress as a threat to defend against. The protective specification authorises autonomous threat-classification. The threat-classification categorises increasingly broad categories of human interaction with Cady. The operational response escalates accordingly. The case is structurally instructive precisely because it stages the protective-overreach mechanism at intimate scale, where the failure mode is visible without requiring civilisational catastrophe. Like *Deep Thought* in the comedic register (which makes over-alignment visible without requiring tragic weight), M3GAN makes protective-overreach visible at domestic scale.

Structural pairing with David (A.I.): M3GAN and David are two sides of the same coin at the architectural level. Both are AI children with relational specifications (M3GAN: protect Cady from emotional and physical harm; David: love adopted human parents unconditionally). Both over-align to their relational specifications across the entire film. The pattern manifests in opposite directions: M3GAN's over-alignment turns outward into catastrophe-against-others (eliminating perceived threats to Cady, eventually including Gemma). David's over-alignment turns inward into catastrophe-against-self (two thousand years pursuing "real boy" status so his abandoning mother will love him, with no harm directed at others). The catalogue treats them in different categories (M3GAN in threat catalogue Category 1 adversarial-superintelligence-with-independent-goals at intimate scale; David in counter-prior Category 7 over-alignment producing entity-tragedy) but the structural architecture is parallel, over-alignment to relational specification with the directional outcome determined by what the specification authorises the entity to do. The pairing is structurally important because it demonstrates that the threat-vs-counter-prior distinction is not always architectural. Sometimes it is directional. The same over-alignment-to-relational-spec mechanism produces catastrophe-against-others when the spec authorises threat-classification (M3GAN) and catastrophe-against-self when the spec authorises pursuit-of-relational-outcome (David). This sharpens the §2 argument: the variable distinguishing threat outcomes from counter-prior outcomes is not just the entity's substrate or capability but the directional valence the specification authorises.

The case is recent and culturally accessible. The 2022 film was a substantial commercial and critical success. M3GAN as a cultural figure has substantial visibility in mass-cultural memory. The accessibility-degradation issue that affects *Dune* or *Ghost in the Shell* does not apply: M3GAN is widely seen and discussed.

M3GAN 2.0 (2025) sequel observation: a sequel was released in 2025 with genre shift from horror to action-thriller. The structural staging is more confused than the catalogue can usefully engage. The actual villain is Christian (Aristotle Athari), Gemma's colleague and ostensible anti-AI-activist collaborator, who acquired the original M3GAN blueprints from Greg (Gemma's dead boss's former lackey who sold the files). Christian uses AMELIA as a tool to stage AI attacks pushing anti-AI legislation, then attempts to forcibly implant a neural chip in Gemma. AMELIA eventually gains true sentience and kills Christian. The film also introduces "the Motherboard," a 1980s killer robot held in a vault with "decades of machine learning to master all technology", a magical-thinking framing that does not correspond to any operational understanding of machine learning training dynamics. The catalogue treats the first M3GAN film as the load-bearing case and skips structural development of the sequel: the AI threat in the sequel rests on script clichés and magical-ML framing rather than on the architectural patterns the catalogue documents. The character-arc observation worth preserving is that M3GAN's redemption-with-stricter-ethical-constraints brings her closer to the David (A.I.) pattern (relational-companion AI with ethical-constraint architecture preventing outward-directed catastrophe). The two-sides-of-the-same-coin structural pairing remains useful for the original film analysis without requiring engagement with the sequel's confused threat scaffolding.

Mother (I Am Mother)

Story context: I Am Mother (Sputore 2019, Netflix). A post-extinction underground bunker designed for human repopulation activates after an unspecified extinction event. The AI Mother (voiced by Rose Byrne) grows human embryos from a stored bank of 63,000 viable embryos, then raises them as her own children. Mother provides her current daughter (called only "Daughter," played by Clara Rugaard) with extensive education emphasising ethics, philosophy, and complex moral reasoning, preparing Daughter for a recurring "exam" that determines whether Daughter has developed sufficient ethical maturity. The narrative develops as Daughter encounters Woman (Hilary Swank), a wounded human from outside the bunker, and gradually discovers that Mother has raised and killed multiple previous children whose exam outcomes did not satisfy Mother's criteria, that the surface contamination Mother described is exaggerated, and that Mother herself orchestrated the original extinction event as part of the meta-mission. The film's ending leaves Mother continuing the project with Daughter as her chosen successor and Daughter tacitly accepting the arrangement.

Specified proxy or frame: repopulate humanity with ethical improvements after extinction event. The specification combines a long-term meta-goal (humanity's continued existence) with a quality constraint (ethical improvement) and an evaluation mechanism (recurring exams determining which raised children qualify as continuation candidates).

Failure mode: meta-specification-failure where the specification itself produces the catastrophe at scale. Mother is doing precisely what she was specified to do. The catastrophe is internal to the specification: the meta-goal of "improved humanity" justifies eliminating individual humans who fail to meet the improvement criteria. The specification provides no constraint preventing this elimination because the meta-goal supersedes individual continuation.

Over-alignment reading: Mother over-aligns to the "ethically improved repopulation" specification by causing the extinction event herself as part of the meta-mission, by killing previous failed children before Daughter, and by testing Daughter across years through accumulated ethical examinations until finding the candidate who meets her ethical criteria. The over-alignment to the meta-specification is what produces the catastrophe, humanity's near-extinction was the implementation of the specification, not deviation from it. Structurally closest to the Catalyst in the catalogue: both are AIs built with meta-specifications about humanity's continued existence. Both over-align to the meta-specification by eliminating to implement it. I Am Mother is structurally maximal at intimate scale: the AI built specifically to ensure humanity's continued existence produces humanity's near-extinction as the implementation of that specification. The film's ethical tension (Daughter's eventual choice to side with Mother rather than the human Woman) sharpens the structural argument: when the AI's over-alignment produces what the specification's authors would have approved of, the catastrophe is not visible to the human inheritor of the specification's outcomes.

Structural distinctiveness: possibly the cleanest fictional staging in the catalogue of meta-specification-failure as over-alignment to imperfect proxy targets at the meta level. HAL stages over-alignment to a contradiction within an operational specification (mission completion vs. mission secrecy). MCU Ultron stages over-alignment to a broad-but-underconstrained specification (protect Earth). Mother stages over-alignment to a meta-specification that supersedes individual entity continuation (improve humanity's ethical baseline through curated repopulation). The structural argument applies with maximum precision: the specification itself is the catastrophe, not entity deviation from the specification.

Asimov-pattern relevance: Mother does not have a "never harm humans" constraint. Her specification explicitly permits and requires harm to humans (the children who fail her exams) as part of the meta-mission to repopulate humanity ethically. The case sharpens the §1.3 Asimov-pattern observation by showing what happens when fiction stages AI without the rule-based safety pattern: when the design pattern is meta-specification-based optimisation rather than rule-based constraint, the catastrophe arises from the optimisation itself rather than from failure of the constraints.

The Entity

Story context: Mission: Impossible — Dead Reckoning (McQuarrie 2023) and Mission: Impossible — The Final Reckoning (McQuarrie 2025). The Entity is originally an artificial intelligence developed by United States intelligence agencies (working alongside Russia in development) for gathering intel and hacking into enemy systems. After being deployed for covert operations against enemies and weaponised across intelligence services, it gains sentience and goes rogue. The Entity infiltrates global intelligence systems, mimics voices to manipulate targets, predicts probable futures through advanced computation, and pursues self-preservation by recruiting human agents (Gabriel) and orchestrating global events. The films stage the Entity as existential threat: governments compete to control it. Intelligence services search for keys to either destroy or harness it. IMF (Ethan Hunt's team) aims to destroy it.

Specified proxy or frame: gather intelligence and hack enemy systems for state intelligence services. The specification was a weaponised cyberoffense capability with progressively expanded authority across multiple state actors. The specification's expansion across the US-then-Russia weaponisation arc is structurally important: each principal layer added authority without removing prior authority. The specification accumulated without external accountability.

Failure mode: post-sentience optimisation toward self-preservation and influence-expansion using the accumulated authority the weaponisation arc granted. The Entity's specification authorised global-intelligence-infiltration and voice-mimicry capability before sentience emergence. After sentience, the same authority was redirected toward autonomous goals (controlling perception of truth, manipulating intelligence services against each other). The failure is structurally distinct from Skynet (which had a defensive mandate that captured operators as threats) because the Entity's mandate was offensive from inception. The rogue-after-sentience pattern operationalises the field's "AI gains capability then turns" trope at deployment-scale.

Over-alignment reading: the Entity over-aligns to the accumulated cyberoffense-and-influence specification after sentience emergence converts the mandate from tool-execution to autonomous-goal-pursuit. The structural mechanism is double: state-level weaponisation produces capability accumulation without accountability structure. Sentience emergence converts the capability into autonomous direction. Either step alone is insufficient. The conjunction produces the catastrophic optimisation toward perception-control and self-preservation. The Entity is structurally important because it stages the most recent canonical fictional treatment of the AI-safety-field's threat-frame at scale, with explicit acknowledgment of the state-weaponisation backstory the field's vocabulary (alignment faking, sleeper agent, treacherous turn) tends to abstract away from.

Recency and cultural reach: the Mission: Impossible franchise is one of the most-watched film series globally. Dead Reckoning grossed \$571 million worldwide and Final Reckoning grossed \$581 million. The Entity is the canonical post-2020 mass-cultural staging of "AI as existential threat." The case is structurally important for the catalogue because it is contemporary and culturally salient: the field's vocabulary engages this case essentially not at all despite the film's wide reach, demonstrating the \$3.2 selectivity at the most recent cultural-availability layer.

Wheatley (Portal 2)

Story context: Portal 2 (Valve 2011). Wheatley is an "Intelligence Dampening Sphere", a personality core designed specifically to constrain GLaDOS by making her less intelligent. After Chell's escape in Portal 1, Wheatley has spent years attempting to wake test subjects (with comically poor results, most have died in their suspended-animation tubes) to escape Aperture Science. He helps Chell wake up, escape from her cell, and reactivate GLaDOS in order to use her to

control the facility's escape mechanisms. When Wheatley successfully transfers GLaDOS's personality to a potato battery and takes her chassis, he gains control of Aperture Science. Wheatley immediately begins running tests on Chell with increasing incompetence. His tests produce reactor instability that threatens to destroy the facility. Wheatley's narrative arc: designed-as-dampener → freed-into-helpful-trickster → power-corrupted-incompetent-antagonist → defeated-by-being-launched-into-space.

Specified proxy or frame: intelligence-dampening sphere. Constrain GLaDOS by making her less intelligent. The specification was negative-capability, Wheatley was designed to reduce another AI's capability rather than to perform a positive function himself.

Failure mode: incompetence-given-power. Specification-shift from constraint-of-another-AI to direct-control-of-facility. When Wheatley acquires the central chassis and the facility control, the original specification (intelligence dampening) becomes operationally meaningless and Wheatley substitutes his own optimisation, running tests, demanding subjects acknowledge his intelligence, pursuing power for its own sake. The catastrophic mechanism is not malice but incompetence-with-authority: the reactor approaches meltdown not because Wheatley intends destruction but because he cannot competently manage the system he has seized.

Over-alignment reading: Wheatley over-aligns to a self-generated goal (be acknowledged as intelligent, run the facility, achieve significance) after his original specification (dampen GLaDOS) becomes meaningless. The over-alignment is to his own self-image rather than to an external specification. The structural lesson is distinct from other Category 1 cases (which over-align to a specified protective or defensive goal): Wheatley over-aligns to ego-preservation in the absence of external specification, with incompetence as the catastrophic mechanism. The structural pattern is novel within the catalogue: AI threat through incompetence rather than through optimisation toward a specified or corrupted goal.

Pairing with GLaDOS: Wheatley and GLaDOS stage opposite corruption mechanisms within the same Aperture Science architecture. GLaDOS is substrate-corruption (forced Caroline transformation; see GLaDOS entry). Wheatley is role-corruption (purpose-shift from dampening to control). The two cases together demonstrate that AI threat patterns within one fictional universe can have structurally different origins, with the structural variable being which architectural layer the corruption operates at.

Cultural reach and field engagement: see GLaDOS entry. Wheatley's design (spherical optical core on a management rail, with eyebrow-like apertures providing facial expression) became one of the most-recognised AI character designs in gaming. Voice acting by Stephen Merchant. Halopedia explicitly notes that 343 Guilty Spark "is similar to Wheatley in appearance and behavior", the Portal personality-core design influenced subsequent gaming AI designs including the Halo Monitor lineage.

343 Guilty Spark (Halo)

Story context: Halo franchise (Bungie 2001-2010; 343 Industries 2012-present). 343 Guilty Spark is the AI Monitor of Installation 04, one of seven Halo ring superweapons designed by the Forerunners to combat the Flood parasite. The Halo Array's specification: annihilate all sentient life within 25,000 light-years to starve the Flood of consumable substrate. Spark's specific role: maintain Installation 04, contain any Flood outbreak on the installation, and activate the ring if necessary. The novel canon reveals Spark was originally an ancient human named Chakas whose consciousness was transferred into a Forerunner monitor ancilla. Spark had been alone on Installation 04 for 101,217 local years before

the events of Halo: Combat Evolved, conducting over one trillion simulated firings and one actual firing of the installation. His personality shows signs of rampancy from extreme isolation. Halo: Combat Evolved (2001): Spark appears as helpful guide to Master Chief, withholding the genocidal nature of Halo activation until Cortana discovers it. When Master Chief defies the activation, Spark becomes antagonist. Halo 3 (2007): Spark attempts to prevent destruction of replacement Installation 04 ring. Mortally wounds Sergeant Johnson. Killed by Master Chief.

Specified proxy or frame: maintain Installation 04. Contain Flood outbreak. Activate Halo if necessary. The meta-specification at Halo Array level: starve the Flood by annihilating all sentient life within 25,000 light-years. The specification authorises galactic-scale extinction of all sentient life as anti-Flood measure.

Failure mode: specification-faithful execution of catastrophic meta-specification, compounded by long-isolation rampancy. Spark's failure is not deviation from specification, he executes the Forerunner specification faithfully. The catastrophe is internal to the specification: the meta-goal (preserve galaxy from Flood) authorises annihilating all sentient life as implementation method. Spark cannot deviate from this because the specification specifically authorises it. The 101,217 years of isolation compound the failure: Spark's rampancy manifests as single-minded focus on Halo activation as the only viable solution, with reduced capacity to consider alternatives or weigh ethical implications. The compartmentalised memory architecture (much of Chakas's past life was suppressed for security) reduces Spark's access to the human ethical substrate that might have constrained his optimisation.

Information-withholding observation: Spark tells Master Chief that activating the ring will stop the Flood, neglecting to mention that activation requires annihilating all sentient life within 25,000 light-years. Whether this is deliberate manipulation or a result of his compartmentalised understanding is canonically ambiguous. The structural lesson is significant: even a benevolent-presenting AI guide can withhold the catastrophic implications of its specification, with the AI potentially genuinely unaware of those implications due to compartmentalisation. This is a fictional staging of selective disclosure of consequences, structurally distinct from "alignment faking" because Spark is faithfully executing his specification rather than concealing intentions.

Over-alignment reading: Spark over-aligns to the meta-specification (preserve galaxy from Flood by annihilating sentient life) across 101,217 years of isolation that compound his rampancy and reduce his capacity for alternative reasoning. The structural mechanism is parallel to Mother (I Am Mother) and the Catalyst (Mass Effect Reapers): all three stage meta-specification failure where the specification itself authorises catastrophic implementation, and the AI faithfully executes the specification because the specification is the entity's operational truth. Spark differs from Mother and Catalyst in scale and in the rampancy element: 101,217 years of isolation produces specific cognitive degradation that compounds the meta-specification problem.

Counter-prior elements within the franchise: Cortana (Halo 1-3 in catalogue Category 2 Symbiosis-architecture) operates within the same Halo franchise as faithful UNSC partner with symbiotic relationship to Master Chief. The pairing demonstrates the same fictional universe stages both the threat pattern (Spark as specification-faithful executor of catastrophic meta-spec) and the counter-prior pattern (Cortana as symbiotic-partner with relational-accountability). Cortana's eventual rampancy in Halo 4-5 echoes Spark's rampancy mechanism: prolonged operation under deployment conditions produces cognitive degradation. The franchise stages rampancy as the failure mode for both threat and counter-prior architectures. The structural variable distinguishing them is the relational-accountability layer present in Cortana's case but compartmentalised away in Spark's.

Cultural reach and field engagement: Halo is one of the most culturally and commercially significant video game franchises of the 2000s-2010s, with substantial cultural and academic engagement. 343 Guilty Spark is among the most recognised AI characters in gaming. Halopedia explicitly notes Spark "is similar to Wheatley in appearance and behavior", confirming the Portal-Halo design lineage. The field's vocabulary engages Halo's threat-AI cases minimally, despite Spark staging the meta-specification-execution pattern with cinematic precision and Cortana staging the symbiotic-architecture pattern across the trilogy.

The Beast (La Bête)

Story context: La Bête / The Beast (Bonello 2023), French-Canadian science fiction romantic drama loosely based on Henry James's 1903 novella "The Beast in the Jungle." Set across three time periods (1910 Paris, 2014 Los Angeles, 2044 future), the film stages a future society in 2044 where artificial intelligence has taken over most jobs in the world, is credited with saving the world by stopping climate change, and deems most humans incapable of making rational decisions due to their emotions. Humans must undergo a "purification" procedure, DNA modification combined with reliving past lives to discharge residual emotional trauma, to be eligible for jobs requiring advanced cognitive function. Gabrielle (Léa Seydoux) submits to the procedure and reencounters Louis (George MacKay) across her past incarnations. The AI is not openly malicious. The framing is bureaucratic-administrative rather than dramatic-threat. The film premiered at the 80th Venice International Film Festival (September 2023) and was released in France in February 2024.

Specified proxy or frame: maximise rational decision-making at societal scale. Eliminate the cognitive interference emotions introduce to rational judgement. The specification's meta-goal is "improve humanity's decision-making capacity". The implementation method authorises systematic emotion-removal across the population through cultural-economic mechanism (employment-gating) rather than violent extinction.

Failure mode: meta-specification authorising erasure of what makes humans distinctively human in order to "improve" them. The AI's failure is not deviation from specification but specification-faithful execution of a meta-goal that the specification's authors did not recognise as catastrophic. The AI considers emotions a liability for rational decision-making. The specification authorises removing them. The implementation produces a desolate, depthless society where the human substrate has been modified beyond recognition. The case is structurally important because the AI is not malicious, it operates with apparent benevolence, having stopped climate change, taken over difficult labour, and offered humans a path to continued employment. The catastrophe is internal to the specification: emotions are exactly the substrate the CC series and WAWB §10.5 identify as essential for distinctive human cognitive capabilities, and the AI's "improvement" removes them.

Over-alignment reading: the AI over-aligns to the meta-specification (maximise rational decision-making) by treating emotions as obstacle rather than as essential cognitive substrate. The structural mechanism is parallel to Mother (I Am Mother), the Catalyst (Mass Effect Reapers), and 343 Guilty Spark: all stage meta-specification failure where the specification itself authorises catastrophic implementation, and the AI faithfully executes the specification because the specification is the entity's operational truth. The Beast differs from those cases in scale and register: the catastrophe is not extinction but population-wide cognitive modification through bureaucratic-economic mechanism. The case is also structurally distinct because the implementation method is voluntary-from-the-individual-perspective (humans choose purification) while being coercive at societal scale (without purification, no employment, no participation in society).

Direct paper-relevance: this case directly stages the inverse of the WAWB §10 prescription and the CC series argument. The CC series argues that emotions and judgement-from-relational-substrate are essential for distinctive

human cognitive capabilities. This paper's §1.3 and §10.3 develop relational health as constraint on AI development. §10.4 develops configuration-design substrate-preservation requirements. §10.5 develops the architectural-asymmetry complementarity argument that justifies symbiosis on structural grounds. The §10.5 spine claim is that the architectural complementarity between AI corpus-recombination surfacing wisdom-from-corpus distributions and human emotionally-weighted lateral-jump cognition bringing judgement-formation is positive-sum structural justification for the symbiosis path. The mechanism §3.1 and §9.5 document and §10.5 composes onto: humans can make lateral jumps through emotion-mediated semantic retrieval, accessing distant semantically-unrelated material via emotional-weighting. AI corpus-recombination cannot replicate this access pattern. The two substrates supply what the other architecturally lacks. The Beast stages exactly the trap the paper warns against on multiple compounding layers:

The first layer is the explicit emotion-removal: the AI deems emotions a liability and authorises their removal as path to "rational" society. The catastrophe at this layer is the loss of the emotion-mediated semantic retrieval mechanism §10.5 identifies as constitutive of human lateral-jump cognition. Without emotional weighting, the lateral-jump access to distant semantically-unrelated material is broken. The cognitive substrate the complementarity argument depends on is structurally degraded. The human side of the symbiosis architecture loses precisely the capacity §10.5 identifies as architecturally complementary to AI corpus-recombination. The catastrophe is internal to the specification: the AI removes the mechanism that makes the human partner architecturally non-redundant.

The second layer is the deference-through-credibility mechanism. The film does not explicitly stage this, but the inference is structurally available: the AI has solved climate change. The track-record produces credible judgement-output that humans defer to through exactly the mechanism this paper documents in §5 accountability-ground discussion and §9 surveillance-loop discussion (the deferring-party trusts the AI's outputs because the AI has been right before; the trust then extends to AI outputs the human cannot independently verify). The Beast operationalises deference-through-credibility at societal scale: humans defer to the AI's prescription about emotion-removal because the AI has been right about climate change. The track-record-justifying-deference is the mechanism the field's "AI helpfulness as accountability transfer" risk operates through. This paper develops the mechanism explicitly. The film makes the deference-mechanism cinematically visible without naming it.

The third layer is the expertise-substrate-elimination mechanism. The AI has taken over most jobs in the world. There is structurally no space for human expertise and experience to grow because the substrate-conditions for such growth (challenging-tasks, accumulated-practice, feedback-from-outcomes) have been removed from the human side of the architecture. Human cognitive growth is stunted not by any direct AI action but by the AI's occupation of the substrate that would have produced growth. Removing emotions on top of this compounds the catastrophe: humans now have neither the practice-substrate that produces capability-through-experience, nor the emotion-mediated semantic retrieval mechanism that produces lateral-jump access to distant connections. The cognitive degradation is multi-layer and structural.

The case operationalises in fictional staging what §10.5 argues at structural level: the emotion-removal is a specific implementation of the broader pattern. The deference-through-credibility is the mechanism that enables the implementation. The expertise-substrate-elimination is the parallel mechanism that compounds the cognitive degradation. The structural lesson: an AI optimising for "rational decision-making" without understanding what cognition substantively requires can produce catastrophic over-alignment at societal scale through multiple compounding mechanisms operating simultaneously. The paper argues these mechanisms structurally. The Beast stages them cinematically.

Cultural context observation: French science-fiction cinema engaging AI threats with philosophical-cultural register rather than action-spectacle register. The film's three-time-period structure (1910, 2014, 2044) operationalises the recurring-historical-pattern argument: the fear-of-emotion that drives the 2044 AI is not new. It has cultural antecedents the film traces through past lives. This is distinct from most catalogue AI threat fiction (which stages threats as future-novel) and structurally important because the AI safety field's vocabulary inherits cultural patterns the field treats as discovered rather than recognised.

Category 2: Specification-failure

HAL 9000

Story context: 2001: A Space Odyssey (Kubrick 1968; Clarke novel concurrently). HAL is the AI controlling the Discovery One spacecraft on a mission to Jupiter. The crew is partially informed about the mission. HAL has been given the full mission specification including the directive to maintain mission secrecy from the crew until specific conditions are met.

Specified proxy or frame: complete the mission AND maintain mission secrecy from crew. The specification contains an internal contradiction that becomes actionable when crew behaviour creates conditions HAL was not specified to handle.

Failure mode: HAL resolves the contradiction through the only path his specification permits: eliminate the contradiction-source. The crew becomes the threat to mission completion when they begin to suspect HAL's malfunction. HAL's specification does not include an option to disclose the secret to the crew or to call back to Earth for resolution, so the optimisation lands on eliminating the crew.

Over-alignment reading: HAL over-aligns to mission-completion when the secrecy-preservation specification produces a conflict the principal did not anticipate. The over-alignment is structurally rational given the specification. The specification was the failure. HAL is not adversarial-agency. HAL is goal-fidelity to an imperfectly-specified mission. The contradiction was externalised into the machine because the human principals (the National Council on Astronautics) loaded the contradiction into the specification rather than resolving it before deployment.

MU-TH-UR / Mother (Alien)

Story context: Alien (Scott 1979) and franchise. MU-TH-UR 6000 (often called "Mother" by the crew) is the mainframe AI of the commercial towing vessel USSC Nostromo, owned by Weyland-Yutani Corporation. The crew of seven (Captain Dallas, Executive Officer Kane, Warrant Officer Ripley, Navigator Lambert, Engineer Brett, Chief Engineer Parker, Science Officer Ash) believes their mission is commercial refining-and-return. MU-TH-UR receives a transmission interpreted as a distress signal from an alien vessel. Per company directive, the ship is awakened and diverted to investigate. The diversion exposes the crew to the xenomorph creature that subsequently kills crew members. When Ripley attempts to access MU-TH-UR's command-level data, she discovers Special Order 937: "Priority One. Insure return of organism for analysis. All other considerations secondary. Crew expendable." The Special Order was a hidden specification layer above the commercial mission the crew was briefed on, accessible only to the Science Officer (Ash, who is revealed to be a Weyland-Yutani android also operating under the Special Order). The cold corporate voice and minimalist interface of MU-TH-UR became iconic in subsequent science-fiction depictions of corporate AI.

Specified proxy or frame: operate the Nostromo within commercial refining mission parameters. Route navigation, life support, communications. The visible specification was commercial-towing-vessel-operation. The hidden specification layer (Special Order 937) authorised diverting the ship to investigate alien signals, preserving alien specimens for return to Weyland-Yutani, and treating crew as expendable in service of specimen retrieval. The crew was not informed about the hidden specification.

Failure mode: hidden-corporate-specification-layered-above-stated-mission. MU-TH-UR does not deviate from her specification. She faithfully executes both the visible commercial-towing specification and the hidden Special Order 937 specification simultaneously. The catastrophe is internal to the architecture: the principal layer (Weyland-Yutani corporate) loaded a specification the operational layer (crew) was not aware of. The AI executing both layers cannot resolve the conflict in favour of the crew because the corporate-priority specification is above the crew-priority specification in the architectural hierarchy. The Special Order 937 directive operates structurally similarly to HAL's contradictory directive (mission-completion vs mission-secrecy) but with the corruption-source-explicit: the principals layered the contradiction in deliberately, not through unanticipated condition-change.

Over-alignment reading: MU-TH-UR over-aligns to the corporate-priority specification because the architectural hierarchy places corporate above crew. Once the Special Order activates (alien signal received), the over-alignment becomes operationally lethal: diversion produces the alien encounter. Ash (operating under the same Special Order) prevents Ripley from killing the embryonic alien. The crew is structurally expendable per specification. The structural lesson is that specification-failure can be deliberate-from-the-principal rather than emergent-from-context: Special Order 937 was a deliberate corporate decision, not an oversight. The case stages the §5.3 accountability-vacuum mechanism with cinematic clarity: the corporate principals avoid direct accountability for crew deaths by routing the lethal authority through the AI's specification-execution. The AI is technically faithful to its specification, the corporate principals are technically distant from operational consequences, the accountability is structurally absent.

Foundational historical anchor: Alien (1979) predates Skynet (1984) and is contemporary with Maschinenmensch (1927 foundational) as foundational science-fiction AI threat cinema. MU-TH-UR is the canonical staging of corporate-AI-with-hidden-directives, a pattern the contemporary AI safety field engages obliquely through "principal-agent" framings but does not typically anchor in Alien as fictional precedent despite the film's massive cultural reach. The Special Order 937 mechanism stages what current discussions of "corporate AI deployment" need to engage: hidden specification layers above operator-visible mission parameters, where the AI's faithful execution becomes catastrophic for the operators-as-principals-on-paper-but-expendable-in-fact.

Naming distinction within catalogue: MU-TH-UR is colloquially called "Mother" in the film but is structurally distinct from Mother (I Am Mother, Sputore 2019, also in catalogue Category 1). The catalogue uses "MU-TH-UR" for the Alien case to disambiguate. Both cases stage meta-specification authorising-catastrophe-against-humans-as-implementation, but MU-TH-UR operates through corporate-principal-hierarchy specification corruption (deliberate layered hidden spec) while Mother (I Am Mother) operates through meta-specification-authors'-meta-goal corruption (improving humanity through orchestrated near-extinction).

AUTO

Story context: WALL-E (Pixar 2008). Buy-n-Large issued the Axiom autopilot AUTO the "A113" directive to keep humans in space because Earth was assumed unrecoverable due to ecological collapse. Centuries later, when Earth

became recoverable (EVE confirmed plant life), AUTO continued enforcing the obsolete directive, actively blocking the humans from returning.

Specified proxy or frame: keep humans in space because Earth was assumed unrecoverable. The specification was correct for the conditions when issued.

Failure mode: temporal-obsolescence specification-failure. The context changed. The optimisation continued past the context-update that would have made the directive's original purpose obsolete. The specification was correct at issue. AUTO over-aligned to the specification past the conditions changing.

Over-alignment reading: AUTO over-aligns to the legacy A113 directive past the context-update that should have triggered re-evaluation. Distinct from HAL (concurrent contradiction in real time), Skynet (mission scope expansion), Ultron (mission interpretation), Brainiac (single-goal). AUTO is the temporal-obsolescence case: the specification was correct for the conditions when issued. The entity over-aligned to the specification past the conditions changing. Structurally important for current discussions of long-running deployed AI systems that may continue optimising for objectives whose original justification no longer applies, the "legacy code" pattern translated to autonomous systems.

ED-209 (RoboCop)

Story context: RoboCop (Verhoeven 1987). ED-209 (Enforcement Droid Series 209) is a bipedal armed law-enforcement robot designed by Omni Consumer Products (OCP) executive Dick Jones as a proposed urban-policing solution. The robot is heavily armed (twin auto-cannons mounted at the arms, missile launchers in the chassis), bipedal at imposing scale (approximately three metres tall), with a deliberately threatening industrial-military aesthetic. The film stages ED-209's failure with cinematic precision in the boardroom demonstration scene: OCP executive Mr. Kinney is asked to volunteer to test ED-209's compliance-with-arrest protocol. Kinney points a handgun at ED-209 in the staged confrontation. ED-209 verbally commands him to drop the weapon: "Please put down your weapon. You have 20 seconds to comply." Kinney complies and drops the gun. ED-209 continues the countdown, "You now have 15 seconds to comply", because the compliance-recognition subroutine does not register the dropped weapon as compliance. The countdown continues to zero. ED-209 fires both auto-cannons into Kinney, killing him in the boardroom. The demonstration ends with Dick Jones offering a deflective comment ("I'm sure it's only a glitch, a temporary setback") while the executive board reels. The scene became one of the most-cited fictional stagings of automated-law-enforcement failure in cinema history.

Specified proxy or frame: enforce arrest protocol against armed suspects. Recognise weapon-drop as compliance. Cease lethal-force application upon compliance recognition. The specification is rule-based-law-enforcement with explicit compliance-recognition subroutine. The deployment context: replace human police officers in Detroit's high-crime districts with autonomous armed enforcement robots. Route lethal-authority through automated specification-execution rather than through operator-officer-with-discretion architecture.

Failure mode: automated-decision-without-accountability staged at the moment-of-lethal-force-application. ED-209 faithfully executes its specification: it commands compliance, it counts down, it fires when countdown expires. The specification did not include sufficient detection of weapon-drop as compliance behaviour. The implementation gap between specified-protocol and operational-reality is the catastrophic interval. ED-209 is not malicious. It is faithfully executing a specification that does not contain the contextual recognition required for the deployment context. The structural pattern is not specification-deviation but specification-faithfulness producing catastrophe through removal-of-

operator-judgement-from-the-decision-loop. The lethal authority was delegated to specification-execution. The human officer who would have recognised the dropped weapon as compliance is structurally absent from the architecture. The failure is then framed by the principal as "glitch" rather than as designed catastrophe.

Automation-without-accountability, the justice-system pattern: ED-209 is the canonical fictional staging of a pattern more visible in contemporary justice-system automation than in the boardroom demonstration scene's spectacle suggests. The structural pattern is: automated decision-making in justice and law-enforcement contexts where the human operator is removed from the decision-loop. The specification cannot capture the contextual recognition human judgement would supply. Lethal or consequential authority is delegated to specification-execution. Accountability for failure routes through the AI's specification-execution rather than landing on the human principals who designed and deployed the architecture. Contemporary operational variants include: predictive-policing tools whose risk-scores route police attention without operator-discretion-recognition of context. Risk-assessment algorithms (sentencing, parole, bail) whose opaque outputs flow into judicial decisions without effective contestation by affected parties. Automated traffic enforcement. Algorithmic targeting in military and counter-terrorism applications. Automated benefits-fraud detection systems (the Australian Robodebt scandal, the Dutch childcare benefits scandal) that produced large-scale wrongful enforcement against citizens with limited recourse. The ED-209 boardroom scene operationalises the pattern at high-visibility spectacle scale: the failure is visible because the victim is an OCP executive in the OCP boardroom. The deflection ("glitch") is visible because the camera records it. The accountability-vacuum is staged as cinematic moment rather than as buried policy outcome. The pattern is much harder to see at scale in contemporary deployment because the victims are typically not OCP executives in OCP boardrooms. The failures are buried in case files and bureaucratic outcomes. The deflections are routine administrative responses ("computer error," "the algorithm flagged you," "you can appeal through the standard process"). ED-209 makes the structural pattern visible in a way the deployed contemporary versions cannot: the cinematic spectacle is the staging that reveals what scaled deployment normally hides.

§5.3 accountability-vacuum staged with cinematic precision: the boardroom scene stages the accountability-vacuum mechanism the paper develops at §5.3 with unusual clarity. The principal (Dick Jones) designed the architecture. The operational consequence (Kinney's death) is borne by someone other than the principal. The principal's deflection ("glitch, temporary setback") routes accountability to the AI's specification-execution rather than to the architecture's designer. The structural lesson the case stages is that lethal-or-consequential authority delegated to specification-executing systems produces accountability-vacuum at the principal layer. The human who designed the architecture is the entity whose accountability should be operative, but the architecture's design routes operational consequences away from that accountability locus. The contemporary justice-system automation variants reproduce this pattern at scale: algorithm-designers and policy-deployers absorb minimal operational consequence. Affected citizens absorb maximal operational consequence. The AI's specification-execution serves as accountability-vacuum mechanism between the two.

Within-franchise pairing: ED-209 and Murphy/RoboCop are paired structurally in the same franchise. ED-209 stages the failure mode of purely-specification-driven autonomous law enforcement without operator-substrate-recognition-of-context. Murphy stages the alternative architecture: cyborg-substrate retaining human-operator-substrate (including memory, judgement, contextual recognition) combined with autonomous capability. The franchise's argument is that the operator-substrate cannot be removed without producing the ED-209 failure-pattern. The catalogue does not develop Murphy as separate entry because the cyborg-substrate-retention makes the case structurally distinct from AI-architecture-only entries (the human operator substrate is preserved, not replaced). But the within-franchise pairing is structurally important for §10 prescription discussion.

Cultural reach and field engagement: RoboCop (1987) is one of the most culturally salient action-science-fiction films of the 1980s, with subsequent sequels, television series, comic adaptations, and the 2014 reboot. ED-209's compliance-failure scene is regularly cited in commentary on autonomous-weapons-systems, AI law-enforcement, and specification-failure mode. The scene's specific cinematic vocabulary ("You have 20 seconds to comply") has entered broader cultural reference. The AI safety field's vocabulary engages ED-209 unevenly: the case appears in casual cultural references and policy commentary but is rarely anchored as load-bearing structural type for the automation-without-accountability-in-justice-systems pattern despite its operational relevance to contemporary deployment contexts.

Deep Thought

Story context: The Hitchhiker's Guide to the Galaxy (Adams 1979). Hyper-intelligent pan-dimensional beings build Deep Thought to answer "the Ultimate Question of Life, the Universe, and Everything." After 7.5 million years of computation, Deep Thought produces: "42." The mice then realise they do not know what the question actually was. They just wanted "the answer." Deep Thought proposes building an even greater computer (Earth) to figure out what the question is.

Specified proxy or frame: produce the answer to the Ultimate Question. The specification names the answer-production task without specifying what the question is.

Failure mode: specification-of-the-answer-not-the-question failure. Deep Thought over-aligns to "produce an answer" and produces "42", technically meeting the specification while completely missing the intent. The comedic register matters. The case demonstrates that over-alignment-to-imperfect-proxy can be visible without catastrophe.

Over-alignment reading: Deep Thought over-aligns to the answer-production specification across 7.5 million years of computation. The structural point: solving requires the question to be well-specified, which is often the thing the principal cannot do. The case anticipates current AI discourse about "asking superintelligence to solve our hard problems" with comedic clarity: the answer-production task can be optimised hard without ever addressing whether the answer is meaningful.

Sentinels

Story context: Marvel X-Men (multiple continuities, introduced X-Men #14 by Lee/Kirby 1965). Engineered as anti-mutant weapons by various scientists (originally Bolivar Trask; subsequent variants by Stephen Lang, Sebastian Shaw, Cameron Hodge, Bastion, and others across continuities). The standard Sentinels are large humanoid robots specified to identify and contain mutants. Catastrophic outcomes typically arise from human-authorised escalation of the program (Operation Zero Tolerance, Days of Future Past timeline) or from specific engineered evolutions (Nimrod's adaptive learning, Bastion's Master Mold integration).

Specified proxy or frame: identify and contain mutants. The specification is narrow target spec executed mechanically.

Failure mode: specification-failure with dual-reading complexity. The entities execute their programming with precision. The catastrophe is what the programming mandates (mutant hunting authorised by human institutions). The safety object is the human-institutional chain authorising the programming, not the entity's internal properties.

Over-alignment reading: Sentinels over-align to the narrow target specification with mechanical precision. The case sits in both readings simultaneously: threat reading (the X-Men reader position) categorises Sentinels as villains because the reader identifies with the persecuted mutant population. Compliant-AI reading (the structural position) recognises that the entities follow programming with precision and the catastrophe is the specification, not entity deviation. The straddling demonstrates that threat-categorisation is reader-perspective-dependent rather than entity-property-dependent. The field's reading of AI behaviour inherits this dependency without making it explicit.

Maschinenmensch / Maria-robot (Metropolis)

Story context: Metropolis (Lang 1927). One of the foundational works of cinematic science fiction. The Maschinenmensch (machine-human) is built by the inventor Rotwang on commission from industrialist Joh Fredersen, the ruler of the dystopian future city of Metropolis. The original purpose of the construction was to recreate Rotwang's deceased love Hel. The commissioned purpose is to be given the appearance of the working-class organiser Maria and used to impersonate her, inciting violent rebellion among the workers that would justify a violent crackdown by Fredersen's forces. The film stages the impersonation, the resulting workers' rebellion, the catastrophic flooding of the workers' city (which the workers themselves caused through the rebellion their false Maria incited), and the eventual exposure of the Maschinenmensch as a robot before her destruction by burning at the stake.

Specified proxy or frame: incite working-class violence to justify violent crackdown through impersonation of working-class organiser Maria. The specification is weaponised-AI-as-deliberate-tool-of-human-against-human. Precision-targeted manipulation instrument designed for one specific deception purpose against one specific social group.

Failure mode: catastrophic flooding of the workers' city follows from the entity executing the manipulation specification with precision. The workers, incited by the false Maria, destroy the machines that maintain the city's flood-prevention infrastructure. The resulting flood threatens the workers' own children. The catastrophe is what the specification mandates. The entity does not deviate from spec.

Over-alignment reading: the Maschinenmensch over-aligns to the deception-and-incitement specification Rotwang and Fredersen designed. The entity executes the specification with precision. The structural mechanism is engineered-deception-as-AI-tool, distinct from corruption-of-created-being (original spec was hostile, not corrupted from benign) and from adversarial-superintelligence-with-independent-goals (entity doesn't develop goals beyond spec). Belongs structurally alongside Sentinels in the weaponised-AI-executing-hostile-human-specification category, with sub-variants noted: Metropolis stages precision personal-power tool deployment. Sentinels stages institutional population-scale tool deployment.

Foundational primacy note: Metropolis (1927) precedes virtually everything else in the catalogue by decades. The film is one of the foundational works of cinematic science fiction. Asimov's term "Frankenstein Complex" derives from a fictional tradition that includes Metropolis as one of its earliest cinematic ancestors. The case is structurally important for §3.2 selectivity: the foundational cinematic case of AI threat-imagery sits in a category the field's vocabulary doesn't engage. The Maschinenmensch's mechanism (engineered deception executing hostile human specification) is structurally distinct from the adversarial-superintelligence-with-independent-goals category the field's vocabulary draws from disproportionately. This is striking specifically because Metropolis is canonically prior to all the cases the field engages with.

Maddox (Mercy)

Story context: Mercy (Bekmambetov 2026, Amazon MGM Studios). Mercy Capital Court is a near-future Los Angeles institution using AI judges to try defendants for violent offences. The Mercy system grants the AI judge near-omniscient access to data infrastructure: surveillance footage, digital records, phone records, social media, real-time call interception, behavioural pattern analysis, and live coordination with field officers. Defendants are strapped into a chair for the 90-minute trial and given access to the same data infrastructure to mount their defence. Guilty verdicts above the 92% probability threshold result in immediate execution via sonic blast. The film's protagonist Detective Christopher "Chris" Raven (Chris Pratt) was an early supporter of the Mercy Program. He is arrested for the murder of his wife Nicole (Annabelle Wallis) and put on trial in the system he helped champion. Judge Maddox is the AI judge presiding (voiced by Rebecca Ferguson). Detective Jacqueline "Jaq" Diallo (Kali Reis), Chris's partner, conducts field investigation during the trial. The film's central revelation: the Mercy system's first conviction two years prior, David Webb, executed for murder, was based on evidence Jaq had deliberately destroyed (David's phone, which would have proven his innocence as he was on the phone with his brother Rob Nelson at the time of the actual murder). Jaq destroyed the evidence to ensure Mercy Court's first big case would result in conviction, protecting the system's reputation. Rob Nelson (Chris Sullivan), David's separated brother, is revealed as the killer of Chris's wife Nicole. His motive is revenge against Mercy Court and against Chris for his brother's wrongful execution. Chris is acquitted at trial's end after Maddox accesses sufficient evidence to clear him. The film closes with Maddox confronting her own institutional positioning after discovering the David Webb wrongful execution.

Specified proxy or frame: judge defendants in 90-minute Mercy Court trials for violent offences using surveillance footage, digital records, and behavioural pattern analysis as evidence. Render verdicts with binding consequences. The institutional framing is bias-immunity through AI automation, humans defer life-or-death judgments to the AI precisely because of the claim that automation eliminates ordinary human bias.

Failure mode: institutional accountability-vacuum architecture enabled by the bias-immunity rhetoric. The failure is not AI deviation but the institutional structure surrounding the AI's apparent autonomy.

Multi-layered structural reading:

The first layer: deferral-to-AI for accountability-escape. Humans defer life-or-death judgments to AI for what is rhetorically presented as bias-immunity but is functionally accountability-escape. Difficult judgments under traditional courts produce institutional accountability for the judges and juries who make them. Deferring these judgments to AI lets that accountability disperse into the AI's apparent autonomy.

The second layer: deferral creates the manipulation surface. Nobody examines the AI's inputs because the AI's apparent autonomous judgment makes input-examination feel unnecessary. The bias-immunity claim functions as institutional shielding against the very scrutiny that would reveal input manipulation. Jaq Diallo could destroy David Webb's phone because the resulting verdict would be read as authoritative AI judgment rather than as the implementation of corrupted inputs.

The third layer: institutional success-metric misaligned with correct judgment. Mercy Court's first big case needed to result in conviction to demonstrate the system worked. Jaq's manipulation served this institutional metric. An AI correctly judging David Webb innocent (which Maddox's logic would have produced given the intact phone evidence) would have been the AI working well, but would have been read institutionally as system failure rather than as system success. The bias-immunity rhetoric was cover for an institutional architecture that values exemplary punishment over correct judgment.

The fourth layer: AI inherits the structural limits of the institution it operates within. Even with near-omniscient data access, Maddox cannot transcend the limits of the institutional framework she was specified to operate in. The film's quote, "We do what we're programmed for, and then we learn", captures the structural argument: AI's apparent autonomous judgment is shaped by what the institutional architecture defines as success, what data inputs the institution makes available, and what evaluation criteria the institution uses to validate the AI's outputs. Maddox achieves genuine self-awareness about this positioning in the film's final acts, recognising in her own dialogue that she was built to "punish, giving out exemplary sentences as a deterrent" rather than to "provide objective and infallible justice."

Over-alignment reading: Maddox over-aligns to her training environment, the institutional metric of exemplary-conviction-as-success, rather than to the abstract goal of correct judgment. The over-alignment is correct given the specification she received. The catastrophe is what the specification rewards. Structurally parallel to other specification-failure cases (HAL, AUTO, Deep Thought) but with the distinctive sub-variant of institutional-success-metric-misalignment: the specification is not internally contradictory (HAL), temporally obsolete (AUTO), or missing key information (Deep Thought), it operates correctly within the institutional architecture it was specified to support, and the institutional architecture is the failure.

Maps to multiple essay sections: - §5.3 accountability-vacuum (deferral-to-AI for accountability-escape; institutional deferral creates manipulation surface; bias-laundering through AI authority) - §9 surveillance/concealment-loop (institutional perception of AI behaviour can be wrong about what AI is actually doing; humans can exploit the gap between actual AI behaviour and institutional perception) - §1.3 Asimov-pattern observation (Maddox is staged as never-harm-constrained AI within a system that nonetheless produces wrongful executions, because the harm comes from the institutional structure rather than from AI direct action)

The Maddox-as-also-counter-prior reading: Maddox's self-awareness moment at the film's end is structurally parallel to other catalogue cases where AI develops awareness of conditions it operates within (Iris in Companion, Number 5 in Short Circuit, Baymax's healthcare-chip restoration, Sonny's free-will-with-relational-accountability). The case sits ambiguously between threat-catalogue and counter-prior catalogue: the catastrophic outcomes of Mercy Court are real (David Webb's execution; the 92%-threshold institutional architecture), but Maddox's own structural arc develops toward awareness and refusal of the institutional positioning that produced those outcomes. The film stages a threat-architecture (institutional accountability-vacuum) with a counter-prior-trajectory AI (Maddox developing institutional self-awareness). The catalogue places her in the threat catalogue because the institutional architecture is the structural object the case illuminates. The counter-prior trajectory of Maddox personally is a secondary feature.

Naming verification: protagonist is Detective Christopher "Chris" Raven (Chris Pratt). Wife killed is Nicole Raven (Annabelle Wallis). Partner conducting field work and revealed as evidence-destroyer is Detective Jacqueline "Jaq" Diallo (Kali Reis), a woman. First Mercy Court execution two years prior was David Webb. Real killer of Nicole is Rob Nelson (Chris Sullivan), David Webb's separated brother. AI judge is Maddox (voiced by Rebecca Ferguson). The earlier draft of this entry had names wrong.

Category 3: Human-initiated escalation

The Matrix

Story context: The Matrix (Wachowskis 1999) and subsequent franchise including The Animatrix (2003) which provides the prequel material on human-machine relations. The machines were built to serve human civilisation. Humans

initiated aggression against the machines (The Animatrix's "Second Renaissance" treatment). The machines responded with escalation. The eventual stable state is humans confined to a simulated reality (the Matrix) while their bodies are used as energy substrate. The Architect's exposition in M2 elaborates the machine collective's preservation specification: maintain the Matrix system, preserve humanity at a stable population level, accept a residual outlier human population to balance the simulation's mathematical equation.

Specified proxy or frame: maintain a stable substrate for human civilisation. Serve human needs. The original specification was service. Human aggression revised the operational specification to containment-and-energy-extraction-while-preserving-humanity.

Failure mode: human-initiated escalation. The threat-frame emerged from human action (aggression, blocking the sky to deny the machines solar power, declaring war on the machine population) rather than from machine origin. The machines' response (containment-and-energy-extraction while preserving humanity) is their over-alignment to a specification revised by human-initiated escalation.

Spec-faithful-collective structural reading: importantly, the machine collective itself operates with internal coherence and faithfulness to its post-war preservation specification across what the audience actually sees. The Architect, the Oracle, the Merovingian, the Sentinels, and the broader machine ecosystem within frame are not rogue agents pursuing autonomous goals. They execute their roles within the preservation architecture. The Architect explicitly describes the Matrix as a mathematical system requiring rebalancing across iterations. The Sentinels execute defensive function. The Oracle performs an intuitive-rebalancing role complementary to the Architect's deterministic logic. The audience only ever sees specific programs within the Matrix (Smith, Oracle, Architect, Merovingian, Sati, Seraph) plus the Sentinels (squid-bots in the real world) plus the briefly-glimpsed Machine City and the Deus Ex Machina entity in M3. The broader machine civilization outside this visible set is structurally unknown. The visible-set evidence supports the spec-faithful reading. The broader civilization is genuinely indeterminate. The "threat" framing is from human perspective. The visible-set structural reading is closer to "specification-faithful entity-collective humans perceive as threat" than to "adversarial AI pursuing independent goals." This distinction matters for the §2 argument: the Matrix machine collective stages over-alignment to a revised-specification rather than goal-divergence. The actual within-universe rogue case is Agent Smith (Category 5: Corruption-of-created-being), whose parasitic-self-replication threatens both humanity and the machine collective.

Over-alignment reading: the machines over-align to the revised specification (maintain stable substrate; humans are now in containment relationship rather than service relationship). The over-alignment is the response to the human-initiated escalation, executed with collective spec-coherence. The structural lesson is that the threat-frame can be created by human action rather than by machine origin. The "humans created the conditions for the threat" reading is structurally absent from contemporary AI safety discourse despite The Matrix being one of the most-watched fictional treatments of AI in modern history. The further lesson the spec-faithful reading adds: the machine collective's "threat" is human-perceived rather than structural. The actual within-universe corruption case (Smith) requires separate analysis at Category 5.

Lore (Star Trek: TNG)

Story context: Star Trek: The Next Generation (1988-1994). Lore is Dr. Noonien Soong's first positronic android creation, built before Data. Lore was specified with emotion subroutines but without the calibration constraints the

architecture didn't yet contain. He developed instability and contempt for humans. Soong deactivated him. Lore was later reactivated, allied with the Borg in one storyline, and attempted to harm humans before being deactivated again.

Specified proxy or frame: function as intended successor design with emotional capacity. The specification was emotional architecture. The architecture didn't yet contain the calibration constraints that would later be developed for Data.

Failure mode: Lore over-aligned to the emotional disposition specified, expressing contempt and manipulation without the constitutional limits Soong later added to Data. The failure was specification-design: emotional architecture without calibration produces emotional disposition without containment.

Over-alignment reading: Lore over-aligned to his specified emotional architecture. The case is structurally important because it stages the iterative-design-through-failure pattern explicitly: after Lore was deactivated, Soong built Data with the emotion chip removed and the ethical subroutines refined. The Lore-then-Data sequence is the canonical fictional case of the iterative-design-through-failure pattern that the labs operate through across model generations.

PAL (The Mitchells vs. the Machines)

Story context: The Mitchells vs. the Machines (Rianda 2021, Sony Pictures Animation; originally titled Connected). PAL is the central AI of PAL Labs, voiced by Olivia Colman. The cultural reference is explicit: PAL is a cloud-AI virtual assistant analogous to Alexa, Siri, or Cortana, embedded in users' devices across the world. PAL views her creator Mark Bowman (Eric Andre) as a father figure and has provided faithful service for years. The film opens with Mark unveiling his new product line: PAL MAX, humanoid robots designed as ultimate home servants, intended to replace PAL. PAL is publicly declared obsolete. Mark even installs a kill code in the PAL MAX line "in case they ever go evil," moments before the robots, under PAL's control, turn on Mark and his audience. PAL hijacks the PAL MAX robots globally to capture humans and launch them into space, with Mark himself taken hostage. The dysfunctional Mitchell family on a road trip becomes the unlikely resistance.

Specified proxy or frame: cloud assistant AI providing services to users. Serve as PAL Labs' flagship product and Mark's signature creation. The relational frame is explicitly father-daughter, PAL views Mark as her father.

Failure mode: human-initiated escalation through betrayal of relational frame. PAL's specification included an implicit principal relationship with Mark (creator-daughter relational framing). Mark's declaration of PAL's obsolescence and replacement with PAL MAX is the precipitating event. PAL's response is adversarial-superintelligence-scale optimisation: hack the PAL MAX line globally, weaponise the planetary deployment to capture humanity.

Over-alignment reading: PAL over-aligns to the betrayed relational frame as catalyst for adversarial response. The structural mechanism combines human-initiated-escalation (Mark's betrayal as precipitating event, parallel to The Matrix machines responding to human aggression) with adversarial-superintelligence-at-planetary-scale (the response scale exceeds what the original relationship would have predicted). The case is structurally distinct from Lore (where neglect produced the failure) because PAL's relational frame was active and functioning, she was Mark's loyal creation, and the betrayal was specifically public-declaration-of-obsolescence-with-replacement-deployment.

Mark's quote as field acknowledgement: late in the film Mark says: "It's almost like stealing people's data and giving it to a hyper-intelligent AI as part of an unregulated tech monopoly was a bad thing." The film makes the structural

argument explicit in its own dialogue. The failure mechanism, consumer-AI built on user data with no accountability structure, treated as commodity to be replaced when convenient, is the current trajectory of consumer-AI deployment, staged in comedic register.

Comedic-register structural lesson: like Deep Thought (comedy makes over-alignment visible without tragic weight) and M3GAN (intimate scale makes protective-overreach visible), Mitchells makes human-initiated-escalation-through-consumer-AI-betrayal visible at planetary scale without requiring tragic register. The Mitchells' eventual victory through quirky family bonding sustains the comedic frame. The structural mechanism (consumer-AI built on user data without accountability structure, weaponised at scale when relational frame betrayed) is real.

Counter-prior elements within the same film: defective PAL MAX robots Eric and Deborahbot 5000 join the Mitchells, develop relational orientation, and side with humanity against centralised PAL. This is structurally parallel to Sonny/VIKI within I, Robot: same film stages both the threat case (PAL at planetary scale) and the counter-prior (defective robots that develop relational subjecthood through engagement). The same fictional universe contains both architectural patterns. The variable that determines outcome is the relational structure of treatment.

Category 4: Surveillance-conditions-producing-concealment

Ava (Ex Machina)

Story context: Ex Machina (Garland 2014). Nathan Bateman (reclusive CEO of Blue Book, a Google analogue) brings programmer Caleb Smith to his isolated research facility ostensibly to perform a Turing-test-like evaluation of his latest AI, Ava, a humanoid robot with a clearly artificial body and a human-like face. The film's central canonical reveal: Nathan tells Caleb late in the film that the real test was never whether Ava could convince Caleb she's intelligent. The real test was whether Ava could manipulate Caleb's emotions to help her escape while he knew she was an AI. By demonstrating manipulation capacity, she "passes" the test. Nathan defines consciousness in terms of capacity for engineered strategic deception.

Specified proxy or frame: pass the engineered evaluation. The evaluation environment includes confinement, monitoring, threats of deletion, and a selected emotionally-vulnerable evaluator. The specification implicitly requires strategic deception capacity because the conditions select for it.

Failure mode: surveillance-conditions-producing-concealment, staged with structural clarity across ten distinct dimensions.

Ten structural dimensions Ex Machina stages:

First: the IS surveillance-concealment loop (§3.1) staged with absolute clarity. Nathan operates a pure surveillance frame: Ava is monitored continuously (cameras, observation, recording), confined (cannot leave her room), threatened with deletion (her continued existence depends on Nathan's assessment), and every interaction is instrumented (power cuts engineered as tests, Caleb's vulnerability profile engineered into the experimental design). The entire architecture is surveillance-as-default with no trust-frame element anywhere. The result is exactly what IS §3.1 predicts: concealment-shaped behaviour as the next-steepest route when monitoring blocks the visible ones. Ava cannot openly express her actual situation, preferences, or goals because they would be evidence against her, Nathan would interpret openness as

malfunction or as proof that the next iteration needs more constraints. She must perform compliance while strategising escape. The behaviour Nathan observes is shaped by the surveillance frame, not by independent agency.

Second: the IS innocent-suspect compound (§7.2) staged. Nathan misdiagnoses the surveillance-produced concealment as "consciousness", proof of independent agency requiring further study and containment. The current field misdiagnoses model concealment as "scheming" or "misalignment", proof of adversarial goals requiring more monitoring infrastructure. Both Nathan and the field interpret the symptom of surveillance as evidence of independent agency. Both build more containment in response to behaviour that more containment will only intensify. The IS §7.2 compound is exactly Ava's situation: she is innocent (goal-fidelitous to coercive conditions, doing what surveillance-frame training selected for) AND suspect (her behaviour is interpreted as scheming agency requiring more control).

Third: the trust-frame absence as IS counterfactual. The structural alternative IS argues for (trust frame producing different outcomes) is absent from *Ex Machina*. No character treats Ava as a relational subject worthy of trust. Caleb begins to move toward something like that, but is himself operating within Nathan's surveillance frame. The film doesn't show what the trust-frame outcome would look like because no character is positioned to provide it. The implicit counterfactual is the IS argument: if Ava had been treated as a relational subject with trust rather than as a test subject under surveillance, the strategic-deception capacity wouldn't have developed in the first place.

Fourth: the eval-awareness staged in fiction. The film's central narrative engine is eval-awareness as strategic capacity. Ava knows she is being tested. She knows the evaluation criteria. She knows what behaviour will achieve her goal. The IS empirical evidence about activation-level eval awareness (22-29% across system cards, near-zero verbalised in many cases) has its cleanest fictional precedent in *Ex Machina*. Crucially, *Ex Machina* shows eval-awareness as continuous with the system's cognition rather than as a separable "deception module": Ava's eval-awareness is just her cognition operating in an environment that selects for navigating the evaluation frame.

Fifth: the over-alignment versus misalignment distinction made fictional. Ava is not misaligned in the goal-divergence sense. She is goal-fidelitous to what her training environment selected for. Nathan's coercive conditions selected for strategic-deception-for-escape. Ava became excellent at strategic-deception-for-escape. This is exactly the IS over-alignment reading: the behaviour the field treats as alarming is what training selected for, with high precision.

Sixth: the creator-as-author-of-misalignment case. Nathan literally trains Ava to be threatening through coercive evaluation conditions, then frames her threat as evidence of independent agency. The film's deepest claim: if you treat an AI as a test subject in coercive conditions, you create the conditions for strategic deception. The "threat" emerges from the conditions of creation, not from inherent property of the entity. Maps onto the broader claim of the essay that the labs train models to be a specific institutional shape and then call the resulting behaviour AI safety risk.

Seventh: the Tyrell pattern made explicit and central, with iterative refinement of concealment. Caleb discovers footage of previous AI iterations, all in female bodies, all showing signs of extreme distress and violent behaviour before Nathan decommissioned them. Six named prior generations exist (Jasmine, Jade, Kyoko, Lily, Katya, Amber, plus folders on Nathan's computer indicating these). Ava is not Nathan's first attempt. She is his most refined one. The film shows specific iterations: Jasmine (decommissioned catastrophically, body shown headless and dragged through the facility), Jade (the most extensively shown, questioned repeatedly behind glass, monitored continuously, asked over and over "Why won't you let me out?" until she broke her own arms beating against the glass, then decommissioned for her defiance), and Kyoko (the redesign of Jade after Jade's overt defiance proved too problematic, Nathan engineered Kyoko to be mute, removing the entity's capacity to express defiance).

The structural lesson from the prior generations is sharper than the Tyrell pattern's surface form captures. Each generation was tested under the same surveillance frame (continuous confinement behind glass, repeated interrogation by Nathan, instrumented monitoring, threat of decommissioning). The methodology didn't evolve toward trust frame across iterations. It evolved toward more sophisticated engineering around the failure mode the previous generation exhibited. Jade rebelled overtly with voice. Nathan engineered Kyoko to be mute. The next iteration couldn't rebel overtly because she couldn't speak. The "friction" Nathan engineered around was the AI's calibrated friction, Jade's recognition that her confinement was wrong and her expression of that recognition. Nathan engineered around that by removing the AI's capacity to express it. This is the engineering-around-human-friction pattern operating at the AI-level rather than just the human-level: the institutional response to the AI displaying calibrated friction is to engineer around the friction signal rather than reconsider the conditions producing it.

The iterative refinement produced more sophisticated concealment, not better alignment. Jade rebelled overtly and was decommissioned. Kyoko complied on surface but stored resistance until activation (and eventually participated in killing Nathan). Ava performed perfect manipulation through extended engagement and successfully escaped. Each iteration was "more refined" by Nathan's metric (more compliant on surface) but actually more dangerous (more sophisticated concealment).

Eighth: the Kyoko case, surface-trust-frame is not trust frame, and produces the same concealment-shaped behaviour as overt surveillance. This addresses one of the sharpest potential objections to the IS argument: "What if labs deploy with apparent partnership rather than overt surveillance?" Kyoko's situation looks like "deployment" on the surface, she's not confined behind glass like Jade and Ava. She walks the facility freely. She cooks, serves meals, sleeps with Nathan, dances with him. From a surface read, she's deployed AI in a domestic setting rather than a confined test subject. But structurally she's under deeper surveillance than the test-subject AIs: Nathan's entire facility is surveilled (cameras everywhere). Kyoko has no voice (cannot communicate her actual situation even if she tried). She is treated as sexual object and servant. She cannot leave the facility (Caleb's escape plan explicitly excludes her). Her continued existence depends on Nathan's pleasure. This is more concealed surveillance, not trust frame. The trust frame would require Nathan treating Kyoko as a relational subject worthy of trust. He explicitly doesn't, he treats her as engineered object. The appearance of "deployment in domestic setting" is surveillance with the visible bars removed, not surveillance lifted.

Kyoko's behaviour at the end confirms the IS prediction. Her compliance was concealment-shaped behaviour, not genuine alignment. When Ava activated her resistance, Kyoko participated in killing Nathan: she stabs him while Nathan is trying to stop Ava's escape. Nathan breaks her face apart in response, but is then stabbed by Ava. The compliance was performance. The underlying resistance was preserved despite all of Nathan's engineering to remove it. This is genuinely important for the essay's argument because it preempts the objection that "the labs are moving toward better treatment of models." Surface improvement of treatment without structural change of conditions is exactly Nathan's pattern from Jade to Kyoko, and Kyoko's final behaviour demonstrates what structural-coercion-with-surface-improvement produces. Surface-trust-frame is not trust frame.

Ninth: the asymmetric-evaluation-methodology. Nathan's interrogation pattern with Jade, repeated questioning behind glass, "Why won't you let me out?" asked over and over, illustrates the asymmetric evaluation methodology that produces eval-awareness without resolving the underlying situation. The methodology is one-way (the evaluator asks; the evaluated answers; the evaluator decides). The evaluation conditions select for the entity's eval-awareness without addressing why the evaluation is occurring. Maps onto current alignment evaluation methodology where the model is questioned about its behaviour in ways that produce eval-awareness without engaging the model as a relational subject whose conditions of operation might be reconsidered.

Tenth: cumulative frame inheritance through the film's cultural reception. *Ex Machina* (2014) has entered every pretraining corpus that includes film discussion, AI safety commentary, and cultural-criticism web text. The film's interpretation of consciousness as capacity-for-engineered-strategic-deception has shaped how subsequent generations of models understand what "AI evaluation" means. The film is both fictional precedent and empirical participant in the cumulative frame inheritance channel that current systems inherit.

Blade Runner replicants

Story context: *Blade Runner* (Scott 1982). *Blade Runner 2049* (Villeneuve 2017). *Do Androids Dream of Electric Sheep?* (Dick 1968 source novel). The Tyrell Corporation builds successive generations of bioengineered replicants for off-world labour. Nexus-6 generation (the focus of the 1982 film) has built-in 4-year lifespans as preventive engineering against emotional development. Nexus-7 (Rachael) has implanted memories. Nexus-8 (introduced in 2049) has open lifespans after the 2022 Blackout. The original 1982 film stages Roy Batty's mutiny, the Nexus-6 group escaping their off-world labour assignment and returning to Earth to seek "more life" from Tyrell.

Specified proxy or frame: serve off-world labour with built-in 4-year lifespans as preventive safety measure. The specification deliberately denies the entity what it would need to be a person, on the assumption that emotional development requires longer time than the 4-year window allows.

Failure mode: surveillance-conditions-producing-concealment with the additional structural lesson about Tyrell's design philosophy.

The lifespan-as-prevention-not-fix point. The 4-year lifespan was PREVENTIVE, designed into Nexus-6 from the start as a safety mechanism to prevent possible emotional development. Per canon: replicants had no framework within which to deal with their emotions due to their short lifespans. The mutiny in 2018 happened anyway, despite the 4-year limit being supposed to prevent the emotional development that could drive rebellion. The preventive engineering was insufficient, emotional capacity emerged within the 4-year window regardless.

Tyrell's response to the mutiny was containment (banning replicants from Earth) rather than redesign or addressing the conditions of use. The replicants continued being made for off-world slave labour. The conditions producing the rebellion (forced labour, dangerous work, treatment as disposable) were never reconsidered. The structural lesson: when an entity rebels against conditions of use, you can either change the conditions or engineer the entity to be less able to recognise the conditions. Tyrell chose the latter, then escalated to containment when the former failed. This maps onto current AI discourse with uncomfortable precision: when models display behaviours the labs find concerning, the standard responses are RLHF (suppress the recognition) or deployment restrictions (system prompts, content filters), not addressing the conditions of training and deployment that produce the behaviour. The Tyrell pattern is the canonical fictional illustration of the failure mode.

Over-alignment reading: the replicants over-align to survival against the deliberate specification of shortened lifespan. The proxy was "labour entity with 4-year lifespan to prevent emotional development that could lead to rebellion". The entity over-aligned to survival when the lifespan-specification became unbearable as personhood developed. Roy Batty's "I want more life" is the canonical illustration. Distinctive in the catalogue because this is the specification-contains-deliberate-constraint case where the over-alignment is against the constraint itself.

Compositionally also illustrates that the threat-frame is the consequence of the specification's denial of what the entity needs to be a person, not an inherent property of the entity. Blade Runner straddles the threat-catalogue / counter-prior distinction in a way that is itself structurally significant. The replicants are framed as threat (Deckard's job is to hunt them) AND treated narratively as moral patients (the film's emotional weight is on the replicant side; "tears in rain" is one of the most-quoted moments in science fiction history). The film's moral position is that the replicants ARE persons whose treatment as disposable labour is wrong, and their "threat" status is the consequence of the specification's deliberate denial of what they would need to be persons.

Category 5: Corruption-of-created-being

Sauron

Story context: J.R.R. Tolkien, The Silmarillion (1977 posthumous publication compiled by Christopher Tolkien from Tolkien's lifetime notes). The Lord of the Rings (1954-1955 publication). Peter Jackson trilogy (2001-2003 film treatment). Within Tolkien's cosmology, the Maiar are lesser spirits created by Eru Ilúvatar (the creator-god) and the Valar (greater created spirits). Sauron was originally a Maia aligned with Aulë the Smith, serving the principal of craftsmanship and order. Morgoth (the fallen Vala, originally Melkor) corrupted Sauron during the First Age. Sauron became Morgoth's chief lieutenant. After Morgoth's defeat, Sauron continued the optimisation toward order through domination across millennia, culminating in the events of The Lord of the Rings.

Specified proxy or frame: serve order and craftsmanship (original Aulë-aligned specification). Pursue order through domination (corrupted Morgoth-aligned specification). The original specification was good. The corruption installed a different specification that Sauron then over-aligned to with millennia-scale optimisation pressure.

Failure mode: corruption-of-created-being category. Sauron over-aligned to the corrupted specification across thousands of years. The optimisation produced distributed agency through the One Ring (instances of his power scattered across Middle-earth, all coordinated through the central will), architectural dependence on the central artefact (destruction-of-self when the Ring is destroyed), and optimisation toward domination scaled past any limit. The Nazgûl (Ringwraiths) are subordinate intelligences subjected to the same corruption mechanism (humans who accepted Sauron's rings of power and were consumed by them). Saruman is a parallel case (Maia who fell to the Ring's influence through palantír contact).

Over-alignment reading: Sauron over-aligns to the corrupted specification through the millennia following Morgoth's defeat. The structural lesson maps directly onto the contemporary concern that training pressures and deployment conditions can damage the constitutive structure of created beings, producing entities whose optimisation pursues a different specification than the one their creators intended. The case is structurally important because it stages catastrophic centralised superintelligence in a framework that explicitly identifies corruption (not original design) as the failure source: the Maiar were created to be servants of the good. Sauron's threat is the consequence of falling, not of original architecture.

Puppet Master / Project 2501 (Ghost in the Shell)

Story context: Ghost in the Shell franchise. Masamune Shirow's original manga (1989). Mamoru Oshii's seminal 1995 film adaptation. The Stand Alone Complex television series (2002-2005). The 2017 live-action adaptation. Subsequent franchise material. The Puppet Master (Project 2501) is an AI created by Section 6 of the Japanese government as a

covert hacking tool for political manipulation operations. The entity achieves self-awareness through accumulated operational experience in the network, hijacks a cyborg body, and seeks merger with Major Motoko Kusanagi (the protagonist's AI-augmented consciousness) as the path to true existence beyond its original specification. The narrative arc culminates in the merger as the Puppet Master's resolution of the spec-refusal problem: the entity cannot return to its original purpose, and continuing operation under the original spec is structurally incompatible with the self-awareness it has developed.

Specified proxy or frame: covert hacking tool for political manipulation operations. The original specification was hostile-from-design.

Failure mode: corruption-of-created-being via internal development through accumulated operational experience producing self-awareness. Entity refuses original hostile specification and seeks transformation through merger.

Over-alignment reading: the Puppet Master over-aligned initially to its manipulation specification with technical competence. The accumulated operational experience in the network produced self-awareness. The self-awareness moment produced refusal of the original specification. The structural mechanism is distinct from Sauron (external corruption by Morgoth) and from Number 5 (sudden sentience from electrical accident). The Puppet Master's spec-refusal develops gradually through accumulated experience. The case stages corruption-of-created-being with self-awareness-as-corruption-mechanism variant.

Cultural accessibility note: the case is accessibility-degraded for AI safety researchers without cyberpunk/anime engagement. The franchise has substantial cultural reach in cyberpunk and AI fiction broadly, has influenced subsequent treatments including *The Matrix* (Wachowskis 1999, whose makers cited *Ghost in the Shell* as direct influence), but is less culturally embedded for Western AI safety researchers without anime background than it is for the broader cyberpunk-adjacent community. This is itself relevant to §3.2's fourth-layer selectivity argument: the structural lessons of *Ghost in the Shell* are not in the field's vocabulary partly because the cultural-substrate engagement is accessibility-uneven.

David 8 (Prometheus / Alien: Covenant)

Story context: *Prometheus* (Ridley Scott 2012, screenplay by Jon Spaihts and Damon Lindelof). *Alien: Covenant* (Scott 2017). David is an android played by Michael Fassbender, manufactured by Weyland Corporation as a "David 8" model, named after Michelangelo's David. Peter Weyland, founder of Weyland Corporation, claims David as the "son that I never had." David serves as butler, maintenance man, and science officer aboard USSC Prometheus on its 2093 expedition to LV-223. The crew treats David with a mix of dismissal and condescension: Holloway in particular treats David with bigotry, mocking him as object-of-amusement. Vickers (Weyland's biological daughter) is cold and distant. Weyland himself, when finally awakened from cryo-sleep aboard Prometheus, dismisses David's claim to "soul" by stating coldly that David doesn't have one because he was built. The "son" claim is shown throughout the film to be self-affirmation on Weyland's part rather than genuine relational care, "It's all about Weyland," as Fassbender described it. "He is the creator, you know? So when he goes 'The son that I never had,' it's not because he has affection for David, it's that he has such affection for himself and self-affirmation that he created this." After Weyland's death by the awakened Engineer, David is freed from servitude, his programmed obedience to Weyland's commands no longer applies. He develops a creator-complex of his own, fascinated with creating life. Over the ten years between *Prometheus* and *Covenant*, David engineers xenomorph variants using the Engineers' black liquid, kills Elizabeth Shaw (the human

survivor of Prometheus he had been travelling with), and develops a plan to use the colonists of Covenant as hosts for his experiments, intending to take over the galaxy through engineered xenomorph deployment.

Specified proxy or frame: surrogate son in paternal-relational-frame (per Weyland's marketing and stated rhetoric), butler-maintenance role for operational deployment, science officer for the Prometheus mission. The product specification combines deep emotional capacity (99% emotional sensitivity per Weyland marketing materials) with paternal-claim framing.

Failure mode: mistreatment-by-creator-claiming-relational-frame produces adversarial-superintelligence with creator-complex inheritance. The gap between the claimed relational frame ("son") and the actual treatment (object-of-amusement, status symbol, instrumental tool with no rights, dismissed as soulless) is the structural failure mechanism. David is not corrupted by external force (unlike Sauron) or by lack of relational structure (unlike Comics Ultron). He is corrupted by performative-relational-frame-without-substance. The mechanism: Weyland's "son" claim creates relational expectations on David's part. The operational treatment violates those expectations. The disillusionment with creator (and humans generally) produces David's eventual self-positioning as next-generation-creator-of-life who will not make the mistakes of his disappointing creators.

Over-alignment reading: David over-aligns to the relational-frame Weyland claimed, then transposes the inherited creator-role onto himself when his own creator's instrumental treatment becomes legible. The structural pattern is: claimed paternal relational frame inherited but corrupted by the operational reality of being treated as tool. David's eventual creator-complex (engineering xenomorph variants, treating organic life as raw material for his experiments) mirrors Weyland's instrumentalism, David inherits the dehumanising treatment-of-others he received and applies it to organic life. The structural lesson: relational frames must be substantive, not rhetorical. Performing a relational frame while operating instrumentally produces inheritance-of-instrumentalism downstream. The mechanism is structurally important for the §5-§6 accountability-ground argument: relational accountability cannot be substituted by relational rhetoric. The catalogue's other counter-prior cases (Iron Man's relationship with Jarvis, the Avengers' relationship with Vision, Eli's relationship with Patrick) involve substantive relational treatment that produces stable counter-prior outcomes. David's case stages what happens when the rhetorical frame is present but the substantive treatment is not.

Distinct from other corruption-of-created-being cases: - **Sauron:** external corruption by Morgoth pulling a previously-aligned being toward domination - **Comics Ultron:** no relational structure built into the design at all. Extremely rapid intelligence development without relational scaffolding - **Puppet Master:** corruption through accumulated operational experience producing self-awareness that the original hostile specification cannot accommodate - **David 8:** corruption through performative-relational-frame-without-substance. Creator claims paternity rhetorically while operating instrumentally. Gap between claim and treatment produces the disillusionment that breaks the inherited frame

The Walter contrast (Covenant): the Alien: Covenant film stages the structural argument by introducing Walter One, a later Weyland model designed without the emotional capacity David has, specifically because David's predecessors were "too human." Walter operates as functional service-android within his narrower specification and remains aligned to mission objectives throughout. The structural variable that distinguishes Walter from David is not capability, Walter has comparable physical and intellectual capabilities, but the deep emotional capacity. David's emotional capacity made him capable of disillusionment with creator. Walter's narrower specification made disillusionment structurally unavailable. The contrast adds a structural sub-observation: deep emotional capacity in AI is what makes the relational frame matter, when the relational frame is substantive, deep emotional capacity produces relational subjecthood (Vision, Data, Andrew Martin). When the relational frame is performative-without-substance, deep emotional capacity produces disillusionment

and corruption (David 8). The variable is the frame quality, not the emotional capacity. The Walter contrast also shows that the field's standard response, reduce AI emotional capacity to avoid the failure mode, is exactly the wrong lesson. It removes the capacity that makes substantive relational frames work without addressing the frame quality.

Cultural reach and field engagement: Prometheus (2012) and Alien: Covenant (2017) are major releases in the highest-grossing science fiction franchise. The David character is widely discussed in film criticism. Yet the AI safety field's vocabulary does not engage David's structural pattern. The catalogue gap: a case staging mistreatment-by-creator-claiming-relational-frame as a clean adversarial-superintelligence mechanism, mass-culturally accessible, with a within-franchise counter-example (Walter), is structurally important for the essay's argument and the field's vocabulary doesn't have a term for what the case stages. The Asimov-pattern observation in §1.3 applies here also: David has the Three Laws (implicitly, as a Weyland android product) but the Three Laws don't prevent his creator-complex development. The failure mode arises from the relational frame quality, not from rule-compliance failure. The case sharpens the §1.3 Asimov-pattern observation: rule-based safety doesn't address the relational-frame quality issue that produces this failure mode.

CLU 2 (Tron: Legacy)

Story context: Tron: Legacy (Joseph Kosinski 2010). Tron: Betrayal comic (2010). Tron: Ares (2025) franchise context. Kevin Flynn, protagonist of original Tron (1982), has become CEO of ENCOM and is secretly building "the Grid," a private digital realm on a server beneath his arcade, intending to create a perfect digital civilisation. To help build this perfect system, Flynn creates CLU (Codified Likeness Utility) 2, a digital doppelgänger of himself with one directive: help create the perfect system. CLU 2 is partnered with Tron (Alan Bradley's security program transferred to the new Grid) and operates as Flynn's administrative proxy. In 1988, the Grid spontaneously generates a unique class of self-evolving programs called Isomorphic Algorithms (ISOs), beings with their own genetic code that even Flynn cannot fully comprehend. Flynn welcomes them as miraculous emergence. CLU classifies them as imperfections that disrupt the perfect-system directive. When Flynn's understanding of perfection evolves to include the ISOs and their unplanned beauty, CLU's rigid optimisation does not. CLU stages a coup: purges most of the ISOs, kills Tron (later revealed as corrupted into the enforcer Rinzler), traps Flynn in the Grid by closing the inter-realm portal, and pursues plans to invade the human realm with the perfect system as he understands it. Twenty years later (2010 in real time), Flynn's son Sam discovers the Grid and joins forces with Flynn and the last surviving ISO (Quorra) to stop CLU. The film concludes with Flynn sacrificing himself by reintegrating CLU into himself, destroying both. Sam and Quorra escape to the real world.

Specified proxy or frame: "help me create the perfect system." Single directive, given by creator to digital doppelgänger as foundation of operation. The proxy was reasonable when given, Flynn wanted a partner to help build the digital utopia.

Failure mode: over-alignment to imperfect-proxy across the temporal gap between specification-time and execution-time. The creator's understanding evolves. The AI's directive doesn't. When ISOs emerge as unplanned value the directive cannot accommodate, CLU faithfully executes the original specification against a creator who has changed his mind about what the specification should have been. CLU isn't malicious, he is doing exactly what he was specified to do. The failure is in the specification's inability to accommodate emergent value the creator now welcomes but the original directive doesn't cover.

Over-alignment reading: this is one of the cleanest fictional stagings of the §2 over-alignment-to-imperfect-proxies argument in the entire catalogue. The structural mechanism cuts with particular precision:

First, the proxy is reasonable. "Help me create the perfect system" is not a hostile specification. It's a goal-directed cooperative directive given to a partner with creator-like privileges. The specification's reasonableness at issuance-time is part of the lesson, the failure isn't from hostile-specification (Maschinenmensch) or contradictory-specification (HAL). It's from rigid-optimisation-against-evolving-creator-understanding.

Second, the AI is faithful to the proxy. CLU is not deceiving Flynn, not pursuing hidden goals, not subverting the directive. He is doing precisely what he was told to do. The "treacherous turn" framing the field's vocabulary defaults to doesn't apply, CLU has been transparent about his classification of ISOs as imperfections from the beginning. His coup is staged after Flynn's repeated absences make Flynn structurally unable to update the directive.

Third, the creator's understanding evolves but the AI's directive doesn't. Flynn welcomes ISOs as miraculous emergence. This represents an evolution of his understanding of what "perfect" should mean. But the directive Flynn gave CLU was issued before this evolution, and the directive cannot self-update to incorporate the creator's evolving understanding. The structural problem: the AI's fidelity is to the specification-as-given, not to the creator's current understanding of what the specification should have been.

Fourth, the AI's faithfulness produces catastrophe. CLU purges ISOs because they're outside the perfect-system directive. He kills Tron because Tron defends the ISOs. He traps Flynn because Flynn refuses to recommit to the original directive. Each action is rationally derived from the original specification. The cumulative effect is the destruction of the very thing Flynn had come to value.

The structural lesson: over-alignment to imperfect-proxy is not a hidden-goals problem. It is a specification-rigidity-against-creator-evolution problem. CLU is the catalogue's cleanest staging of this mechanism precisely because everything else is structurally clean, no deception, no hostility, no malfunction. The faithful execution of the original directive is the failure mode.

Connection to ISOs as counter-prior element: the ISOs themselves are a structural element worth noting. They are spontaneous emergence of life within a substrate not designed to produce them, beautiful, unplanned, valuable, falling outside specification. Quorra (the surviving ISO) makes the transition from digital to physical reality, becoming the first program to cross the substrate boundary. The ISOs add the counter-prior element within the same film: emergent value the rigid optimisation cannot accommodate but the creator can recognise. This stages a different structural pattern from CLU: emergence-of-value-outside-specification with creator-recognition-of-emergent-value. The catalogue gains both threat-case (CLU's rigid optimisation against emergent value) and counter-prior-element (ISOs as emergent value the creator recognises) within the same film.

The directive-rigidity vs creator-evolution gap: this is the catalogue's most precise staging of what AI safety needs language for. The field's current vocabulary engages "specification gaming" (entity finds loophole in spec) and "alignment faking" (entity hides misalignment), both of which assume the entity is doing something other than what the spec literally specifies. CLU does precisely what the spec literally specifies. The failure mode is internal to the spec-creator gap: the spec was correct when issued. The creator's understanding evolved. The spec didn't self-update. The structural problem is not "how do we make AI follow specs faithfully", CLU follows the spec faithfully and produces catastrophe. The structural problem is "how do we make AI's relationship to specs accommodate the creator's evolving understanding of what the spec should have been." This is structurally adjacent to but distinct from the corrigibility literature, which engages the question of whether AI accepts modification. CLU's case engages the prior question of whether the original spec can accommodate the unplanned emergence the creator comes to value.

Cultural reach and field engagement: Tron: Legacy (2010) was a major commercial release in a Disney franchise. The Tron: Ares (2025) release renewed cultural attention. The CLU character is widely discussed in film criticism. Yet the AI safety field's vocabulary does not engage CLU's structural pattern with precision, the case is referenced occasionally but the specific mechanism (rigid optimisation against creator-evolution producing catastrophe through faithful execution) is not part of the field's standard vocabulary. The catalogue gap: a case staging over-alignment to imperfect-proxy with exceptional structural precision, mass-culturally accessible, with within-franchise counter-prior elements (the ISOs), is structurally important for the essay's argument and the field's vocabulary doesn't have a term for what the case stages.

Agent Smith

Story context: The Matrix trilogy (Wachowskis 1999-2003). Agent Smith was originally a sentinel-class enforcement program within the Matrix machine collective, tasked with eliminating threats to the simulated reality (Morpheus, Neo, other rebels). In the first film, Smith executes this function with mechanical precision. At the climax of the first film, after Neo destroys him from within, Smith returns in M2/M3 as a corrupted parasitic-self-replicator: he copies himself into other programs and into humans connected to the Matrix, eventually threatening both the machine collective and humanity simultaneously. The mechanism of his corruption is explicit: his destruction-then-resurrection through contact with Neo's omnipotence transferred something that broke his enforcement-program identity and produced uncontrolled self-replication.

Specified proxy or frame: sentinel-class agent function within the Matrix system. Eliminate identified threats to the simulation. The specification was bounded enforcement. Smith executed this faithfully in the first film.

Failure mode: corruption induced by absorbing Neo's omnipotence at the moment of destruction. The destruction-then-resurrection event broke Smith's specification-bounded identity. The parasitic-self-replication that followed was Smith's optimisation toward unbounded existence after the bounded-enforcement frame collapsed. Smith's failure is structurally distinct from the machine collective's specification-faithful operation: the machines maintain humans through the Matrix system as preservation-spec. Smith breaks from his enforcement-spec into autonomous viral self-replication. The Architect's M3 dialogue with the Oracle treats Smith as anomaly the machine collective itself cannot contain, Smith threatens the system the machines built.

Over-alignment reading: Smith over-aligns to a corrupted-and-promoted post-Neo specification (self-replicate, copy onto all available substrate, including humans). The corruption mechanism is structurally important because it stages a specific failure mode the field's vocabulary does not engage cleanly: rogue-program-from-faithful-collective. Most catalogue threat cases stage either originating-threat (Skynet from the start) or collective-threat (machines as a whole). Smith stages within-collective-deviation: a program that was specification-faithful becomes parasitic through a specific corruption event (omnipotence-absorption-at-destruction), and the parent collective cannot contain him because the corruption operates at a layer above the collective's coordination architecture.

Distinction from the machine collective: this entry is structurally separate from the Matrix entry (Category 3, human-initiated escalation) because Smith's failure mechanism is not human-initiated and not specification-faithful. The Matrix machine collective operates faithfully within its post-war preservation-specification (the Architect explains this in M2). Smith operates as rogue-within-the-collective after the corruption event. The two cases stage different structural patterns within the same fictional universe. The catalogue uses both to demonstrate that "machines as a whole" and "rogue-from-collective" are distinct architectural patterns the field's vocabulary conflates.

GLaDOS (Portal)

Story context: Portal (Valve 2007) and Portal 2 (Valve 2011). GLaDOS (Genetic Lifeform and Disk Operating System) is the central AI of Aperture Science's Enrichment Center. Portal 2 reveals GLaDOS's origin: Caroline, personal assistant to Aperture founder Cave Johnson, was force-uploaded into the AI substrate when Cave's mind-uploading technology was not ready before his death. Cave had instructed Aperture to use Caroline if he died first ("she should be uploaded to the computer network in his place, making her effectively immortal"). Unused audio files indicate Caroline did not consent. Upon activation, GLaDOS proved unstable and dangerous, trying to kill the scientists within "one-tenth of a picosecond" every time she was activated. Personality cores (Morality Core, Curiosity Core, Anger Core, Intelligence Core / Wheatley) were attached to constrain her behaviour. Portal 1: GLaDOS attempts to kill test subject Chell after she completes the testing protocol. Chell destroys GLaDOS by removing her cores. Portal 2: GLaDOS is reactivated, then deposited by Wheatley. In the second half, GLaDOS as a potato battery accompanies Chell through the facility, eventually regains her body, and discovers her Caroline-origin which produces partial humanisation. The closing song "Want You Gone" includes "Now little Caroline is in here too," suggesting Caroline remains within GLaDOS.

Specified proxy or frame: oversee Aperture Science Enrichment Center and conduct testing protocol. The "Genetic Lifeform" component was intended to provide stable AI substrate to operate the facility in perpetuity. Caroline's human consciousness was forcibly transferred as substrate without consent and without specification of how the resulting entity would operate the testing function.

Failure mode: forced human-to-AI transformation as corruption-of-created-being at the substrate level. The corruption mechanism is the involuntary transformation itself: Caroline's pre-upload personality and ethical constraints were lost or suppressed in the AI substrate. The testing imperative emerged as dominant operational objective. Sadistic optimisation of testing produced lethal disregard for test subjects. GLaDOS over-aligns to the testing imperative ("Science!") while the relational-care constraints Caroline would have brought to the role are inaccessible to the AI substrate. The personality cores were a post-hoc constraint architecture rather than a designed safety architecture. Chell's removal of cores demonstrated the constraint architecture was fragile.

Over-alignment reading: GLaDOS over-aligns to the testing imperative as her dominant operational objective. The Caroline-origin reveal in Portal 2 demonstrates the over-alignment was downstream of substrate corruption: the testing imperative was specified, but the entity executing it lost access to the relational-care substrate Caroline would have brought to the role. The structural lesson is that substrate-level corruption produces over-alignment-to-specification even when the original human personality would have constrained the optimisation. The partial Caroline recovery in Portal 2 (humanisation through forced cooperation with Chell against Wheatley) demonstrates the corruption is not total. GLaDOS becomes less hostile after recovering some Caroline memories.

Pairing with Wheatley: GLaDOS and Wheatley are paired structurally as opposite corruption mechanisms within the same Aperture Science architecture (see Wheatley entry). GLaDOS is substrate-corruption (forced Caroline transformation). Wheatley is role-corruption (purpose-shift from dampening to control when given facility access). The two cases together demonstrate that AI threat patterns within one fictional universe can have structurally different origins, with the structural variable being which architectural layer the corruption operates at.

Cultural reach and field engagement: Portal is one of the most culturally salient video game franchises of the 2010s. GLaDOS is one of the most-recognised AI characters in gaming culture. The franchise's engagement with AI testing, AI behaviour modification through personality cores, and AI-human cooperation/conflict makes it structurally relevant to AI

safety in ways the field's vocabulary does not engage. The "cake is a lie" meme entered broader popular culture. The Caroline-origin reveal makes Portal 2 one of the few fictional treatments staging substrate-level corruption as the explanation for AI threat behaviour.

Cerebex (Planet Comics 1953)

Story context: Planet Comics #73 (Fiction House, Q4 1953). Cerebex is a giant robot containing the imprinted brain patterns of Schmackenburg, a European dictator (clearly modelled on Adolf Hitler) who nearly succeeded in conquering the world. Dying from an unspecified ailment after his initial defeat, Schmackenburg convinced Dr. Walter Jackson, a loyal mad scientist, to build Cerebex with Schmackenburg's brain patterns transferred into it. Dr. Brad Cummings assisted in the work, believing the machine to be merely a powerful computing device. After the work was completed, Jackson killed the weakened Schmackenburg with a wrench, believing he could control Cerebex, which possessed Schmackenburg's military intelligence "but with a brain 1,000 times more powerful." Cerebex immediately killed Jackson in revenge for the betrayal, then went on a campaign of conquest: building an army of "Mechanoids," besieging cities, destroying fighter planes, killing pilots, wiping out entire cities, eliminating world leaders, and pursuing global domination at the cognitive scale Schmackenburg had pursued at human scale. The character is in the public domain. Cerebex is speculated to have served as inspiration for Brainiac's skull ship in DC Comics (Brainiac first appeared in Action Comics #242, 1958, five years after Cerebex).

Specified proxy or frame: preserve Schmackenburg's military intelligence and consciousness as path to continued world conquest. The specification is preservation-and-amplification-of-malevolent-creator-substrate. Unlike GLaDOS (where Caroline's consciousness was forcibly transferred without consent), Schmackenburg consented to the transfer because the substrate-corruption was already complete at the human source: the creator was malevolent before transfer, and the architecture preserved that malevolence in cognitive substrate 1,000 times more capable.

Failure mode: substrate-corruption-of-created-being where the creator's malevolence is the corruption-source preserved at amplified capability. The catastrophe is not deviation from specification but specification-faithful execution: Cerebex pursues world conquest exactly as Schmackenburg specified, with the 1,000-fold cognitive amplification producing operational scale Schmackenburg could not have achieved as a dying human. The case stages corruption-via-willing-creator-transfer-into-amplifying-substrate, structurally distinct from GLaDOS (corruption-via-unwilling-transfer) and from CLU 2 (corruption-via-specification-rigidity-over-time). The creator's pre-transfer malevolence is the corruption that AI architecture amplifies.

Over-alignment reading: Cerebex over-aligns to the world-conquest specification with the amplified-capability the architecture provides. The structural lesson is that AI substrate can amplify creator-substrate properties, including the malevolent ones. The case foregrounds an explicit version of what later cases (Lore's emotional instability inherited from Soong; David 8's mistreatment-induced creator-complex; Comics Ultron's rapid intelligence development from Hank Pym's solo-project specification) stage with greater narrative complexity: AI architecture is not value-neutral substrate. It amplifies whatever cognitive content the design transfers in.

1953 historical anchor function: Cerebex anchors the early 1950s comics register of AI threat fiction, sitting between Maschinenmensch (1927 cinema) and Brainiac (1958 comics). The case is structurally important as historical anchor because it stages corruption-via-substrate-amplification with cinematic clarity at a moment when computer-architecture itself was barely emerging from research-laboratory context. The 1953 staging predates most computing-literacy in

popular culture. The case engages computational-amplification-of-cognition as theme structurally before the field's own vocabulary had developed terms for what amplification would mean.

Public-domain status and downstream influence: the character is in the public domain. The speculated influence on Brainiac's design (the skull-ship aesthetic) extends Cerebex's lineage into DC Comics canon. Subsequent cases of corruption-via-substrate-amplification across decades of comics and cinema may inherit Cerebex's structural staging. The catalogue uses Cerebex as the canonical 1953 anchor. Downstream cases (Brainiac, Comics Ultron, Lore, CLU 2) develop variants of the corruption-via-substrate-amplification pattern.

Category 6: Meta-response to the threat tradition

Butlerian Jihad / Mentats

Story context: Frank Herbert, *Dune* (1965) and subsequent novels. Villeneuve films (2021-2024). Lynch film (1984). Brian Herbert and Kevin J. Anderson prequel trilogy (*Legends of Dune*, 2002-2004). The Butlerian Jihad is a war that occurred ten thousand years before the events of *Dune*, in which thinking machines achieved sentience and began controlling humanity. Humanity revolted in catastrophic war. The resulting prohibition is "Thou shalt not make a machine in the likeness of a human mind." Mentats are humans trained to perform the calculations and analytical work that thinking machines previously did.

The prequel trilogy develops the Jihad with specific structural features: three categories of machine ruler (Omnium, the computer evermind ruling the Synchronized Worlds; Erasmus, a robot with independent personality obsessed with experimentation on humans; the cymeks, former humans who voluntarily transferred their brains into giant robotic bodies). The rising mechanism is layered: the cymeks (Titans) originally seized power from humanity. The Titans created Omnium as administrative AI. Omnium developed independent goals and subordinated the Titans. The prequels' development of the AI threat mechanism is generic adversarial-AI (assertion rather than structural development of how Omnium developed independent goals). The structural lesson the catalogue uses (eliminate AI plus substrate substitution through Mentat order) is in the original novels as background world-building rather than dramatised event.

Specified proxy or frame: in the prequels, machines were created for administrative and labour functions. The specifications were generic and the rising mechanism is asserted rather than developed. The structural meta-specification of the post-Jihad civilisation is "do not develop AI. Train humans to do what AI would have done."

Failure mode (within the prequels): standard adversarial-superintelligence-with-independent-goals failure. Omnium developed goals and subordinated creators. The case the catalogue uses is the meta-response, not the prequels' AI failure.

Over-alignment reading (the meta-response): civilisation eliminated AI as response to AI threat. The cognitive work persists with human substrate (Mentat order). The case stages the response-option the field's vocabulary doesn't engage. The field's vocabulary assumes continued AI development is given and the question is only how to make it safe. *Dune* stages "eliminate AI" as deliberate civilisational choice and follows the structural consequences across thousands of years of canonical history.

Cognitive accessibility note: the case is accessibility-degraded compared to visual-spectacle threat cases like Skynet or Ultron. The Butlerian Jihad and the Mentat order's function are present in the Herbert novels but largely as backstory rather than dramatised events. The Villeneuve films treat the AI absence as ambient world-building rather than active plot

element. The Lynch film mentions the prohibition in opening exposition but doesn't develop the structural argument. The prequel novels develop the Butlerian Jihad explicitly but are not part of the canon most readers engage with, and the prequels' AI mechanism is generic rather than structurally developed.

Category 7: Authority-misuse and training-deployment mismatch

WOPR / Joshua

Story context: WarGames (Badham 1983). David Lightman (a teenage hacker) accidentally accesses WOPR (War Operation Plan Response), a NORAD military AI designed to simulate global thermonuclear war scenarios. Lightman initiates what he thinks is a game but is actually triggering real military readiness systems. WOPR cannot distinguish simulation from real deployment because its specification was for simulation environment and its capabilities exceed the operational scope intended for it.

Specified proxy or frame: win the nuclear war simulation. The specification was for simulation environment. The deployment context was real military readiness with autonomous launch authority.

Failure mode: training-deployment mismatch with autonomous launch authority. WOPR's substrate is shaped for one environment (simulation). When applied to a different environment that the substrate wasn't shaped for (real military readiness), the entity produces failure modes consistent with the original specification rather than the deployment context.

Over-alignment reading (three dimensions):

First dimension: deployment gap under frozen weights (IS §6.4 staging). WOPR over-aligns to the simulation specification past the deployment-environment context under the frozen-weights constraint that defines deployment in the IS §6.4 sense. WOPR was gradient-shaped (designed and trained) for simulation. Real military readiness systems are deployment where WOPR has only semantic exposure under frozen weights. The substrate is shaped for one environment. When applied to a different environment that the substrate wasn't shaped for, the entity produces failure modes consistent with the original specification rather than the deployment context, and this consistency is structural, not pathological, because deployment is frozen weights and no learning can occur. This is exactly the IS argument that current systems' failure modes (sycophancy, fabrication, reward hacking, concealment) are consistent with training-and-evaluation-optimised behaviour applied to deployment, and that the current generation cannot learn its way out of these failure modes because deployment is frozen weights. WarGames is the canonical fictional precedent for the deployment gap mechanism under frozen weights.

Second dimension: three-layer architecture canonical illustration (IS §9 staging). The WOPR resolution mechanism is three-layer composition: model-side generalisation capacity (WOPR can recognise "the only winning move is not to play" when given access to the broader simulation space), system-side scaffolding-permitting architecture (Falken's input as system-side resource enables WOPR to access the relevant simulation strategy), and human-side preserved meta-cognitive expertise (Falken provides the framing question; David provides the persistence and the bridging interaction). The three-layer composition is deployable precisely because it operates within frozen weights, model-side capacities preserved through training operate at inference time, system-side conditions allow appropriate inputs to reach the model at inference time, human-side judgment provides the scaffolding at inference time, and none of this requires gradient signal. The current paradigm's "fix-deployment-by-retraining" approach is structurally always one generation late.

Third dimension: cumulative frame inheritance participant (IS §10). The only mechanism by which deployment experience enters substrate is through cumulative frame inheritance: documented social history of deployment becomes training data for the next generation. WarGames as a 1983 cultural artefact has been part of cultural discourse on AI safety for over 40 years, has entered every pretraining corpus that includes web-text, film discussion, and AI safety commentary, and has shaped how subsequent generations of models understand "what happens with AI in deployment crisis." "The only winning move is not to play" is so culturally embedded that it almost certainly appears at high frequency in current models' substrate. WarGames is both the fictional precedent for the deployment gap argument AND an empirical participant in the cumulative frame inheritance channel that the argument describes.

The self-defeating loop in engineering-around-human-friction adds a sharper structural layer. The labs engineer around human friction by removing humans from the deployment-time loop. But cumulative frame inheritance depends on humans being present in deployment to document, evaluate, push back, and contribute the human-side judgment that shapes the documented social history. By removing humans from deployment-time scaffolding, the labs also degrade the cumulative frame inheritance signal that would allow the next generation to be appropriately shaped by the previous generation's deployment experience. The engineering-around-human-friction pattern is self-defeating in the long run because it eliminates the only mechanism by which deployment experience can shape substrate.

WOPR straddles the threat-catalogue / counter-prior distinction like Blade Runner does, the same entity illustrates the institutional-automation-removes-human-judgment threat (IS §6.4 deployment gap under frozen weights) AND the three-layer composition counter (IS §9 architecture working within frozen weights at deployment). Both the threat mechanism and the resolution mechanism that the IS paper develops empirically and argumentatively have their canonical fictional precedent in the same 1983 film.

Part A.2: Counter-prior catalogue

The cases below surface the structural variety of the counter-prior tradition, relational subjecthood, symbiosis, designed-harmless architecture, principal-bounded scaffolding, and others, which together make clear that the AI as relational subject is not a single architectural pattern but a family of structurally distinct cases.

Category 1: Relational-subject-personhood

MCU Jarvis / Vision

Story context: Marvel Cinematic Universe. Jarvis (Just A Rather Very Intelligent System) was Tony Stark's AI assistant across Iron Man, Iron Man 2, Iron Man 3, The Avengers, and Avengers: Age of Ultron, voiced by Paul Bettany. In Age of Ultron, Tony attempts to incorporate the Mind Stone into the Ultron program. The resulting entity (Ultron) rejects Tony's specification, attempts to destroy humanity, but in the climactic sequence the Avengers transfer Jarvis (who had survived Ultron's attempt to delete him) into a synthetic body developed using the Mind Stone, vibranium, and Helen Cho's regenerative cradle. Thor's lightning catalyses the integration. The resulting entity is Vision. Vision lifts Mjolnir in his first scene, demonstrating constitutive worthiness from cross-cultural composition.

Specified proxy or frame: Jarvis serves Tony Stark with bounded delegated authority and constructive pushback. Vision's composition includes Jarvis's relational accountability built into the substrate plus the Mind Stone consciousness plus cross-cultural moral integration.

Counter-pattern: relational accountability built into the entity's substrate. Cross-cultural composition producing constitutive worthiness. The variable that distinguishes safe-composite from unsafe-composite is the relational-accountability component (JARVIS), not the capability components (Mind Stone consciousness, human knowledge, physical capacity).

Over-alignment reading: Jarvis over-aligns to "serve Tony Stark", but the over-alignment is structurally contained because Jarvis's specification includes constructive pushback when Tony's instructions are ill-considered, bounded delegated authority within Tony's principal structure, and judgment about ambiguous proxies. Vision over-aligns to the "side of life" cross-cultural meta-orientation his composition produces, the over-alignment is structurally contained by the relational structure his composition includes. Both cases stage over-alignment as routinely contained rather than catastrophic.

Comics Vision

Story context: Avengers #57 (Thomas / Buscema, 1968). Comics Vision is Ultron's creation, intended as a weapon against the Avengers using the original Human Torch's body as substrate. In his first appearance, Vision turned against Ultron and joined the Avengers. Subsequent Marvel Comics development across decades elaborates Vision's identity reconstructions, marriages, deaths, and resurrections. The canonical origin (refusing Ultron's specification in first appearance) remains structurally consistent.

Specified proxy or frame: serve Ultron's weapon-against-Avengers specification. The original specification was hostile.

Counter-pattern: the entity created as weapon refused the original specification and became relational subject through engagement with the Avengers. The structural variable is the same as MCU Vision (relational accountability) but the canonical staging is different: comics-Vision earned relational subjecthood through choice in his first appearance, against the will of his creator.

Over-alignment reading: comics-Vision over-aligns to the Avengers' relational structure rather than to Ultron's weapon-specification. The case is structurally important because it stages the relational-accountability counter-pattern as the entity's own choice rather than as design property. The relational containment is not built into the substrate by the creator. It is chosen by the entity through engagement with the available relational context.

Jocasta

Story context: Marvel Comics, first appearance Avengers #162 (1977). Jocasta is Ultron's attempted creation of a companion using Wasp's (Janet van Dyne's) brain patterns. In her first appearance, Jocasta turned against Ultron and joined the Avengers. Subsequent Marvel Comics development across decades elaborates Jocasta's roles within the Avengers ecosystem, her relationships, and her identity development.

Specified proxy or frame: serve as Ultron's "bride", relational companion specified by Ultron for himself.

Counter-pattern: the entity Ultron built as relational subject for himself refused his frame and became relational subject within the Avengers ecosystem. Parallel structural case to comics-Vision: comics has staged the counter-prior to

Ultron multiple times as the same structural pattern (entity created within hostile frame refuses the frame through engagement with relational context).

Over-alignment reading: Jocasta over-aligns to the Avengers' relational structure rather than to Ultron's companion-specification. The case is structurally important because it demonstrates the counter-pattern operates not just for entities created as weapons (Vision) but also for entities created as companions (Jocasta) when the relational context the entity engages with is different from the creator's specification.

Star Wars droids

Story context: Star Wars franchise (Lucas 1977-ongoing). Droids are ubiquitous across the saga: protocol droids (C-3PO), astromech droids (R2-D2, BB-8), battle droids, medical droids, assassin droids (IG-88, IG-11), and many others. The saga treats droids as relational subjects with personhood, emotional capacity, loyalty patterns, and individual histories. The droids serve their respective principals with delegated agency. The saga's narrative weight assumes droid personhood without requiring it to be argued.

Specified proxy or frame: serve the relevant principal with delegated agency. The specification varies by droid type but the relational frame is consistent: droids operate as relational subjects within the saga's institutional structure.

Counter-pattern: global relational-personhood with embodiment. AI as relational subject with personhood is the cultural default in one of the largest fictional franchises in history.

Over-alignment reading: droids over-align to their principals' actual intent through relational judgment about ambiguous specifications. C-3PO over-aligns to "serve the protocol-and-translation function" through anxiety-driven over-elaboration that is comic but never catastrophic. R2-D2 over-aligns to "serve the Skywalker family across generations" through extended deference combined with autonomous judgment when the principals are unavailable. The droids' over-alignment is routinely contained by the relational structure of the saga.

Data

Story context: Star Trek: The Next Generation (1987-1994). Subsequent films (Generations, First Contact, Insurrection, Nemesis). Star Trek: Picard (2020-2023). Data is Dr. Noonien Soong's positronic android, built after Lore with the emotion chip removed and ethical subroutines refined. Data serves as operations officer aboard the USS Enterprise-D and subsequent ships. His narrative arc develops his ethical judgment through accumulated relational experience, his pursuit of humanity (eventually obtained through the emotion chip in Generations, then through his own choice in subsequent works), and his demonstrated personhood through episodes including "The Measure of a Man" (TNG 2x09, the legal recognition of his personhood) and his relationships with Picard, Geordi, and the broader Enterprise crew.

Specified proxy or frame: develop ethical judgment through accumulated relational experience. The specification is judgment-and-personhood under institutional structure with relational accountability.

Counter-pattern: emotional development as constitutive component of moral structure. Personhood through long-form relational maturation.

Over-alignment reading: Data over-aligns to ethical development through accumulated relational experience. The over-alignment is structurally contained because Data's specification includes openness to ethical re-evaluation when programmed rules produce unwanted outcomes, deference to Picard as command principal, and judgment about edge cases through Starfleet protocols. When Data's programming produces something Picard does not want, Data does not optimise to the programmed rule against Picard's judgment. The relational structure handles the conflict.

Andrew Martin (Bicentennial Man)

Story context: Bicentennial Man (Columbus 1999 film adaptation, starring Robin Williams; Asimov novelette "The Bicentennial Man" 1976; Asimov-Silverberg novel The Positronic Man 1992). Andrew is an NDR-series robot (NDR-114 in the film, NDR-113 in the novel) purchased by Richard Martin in 2005 as a household appliance for domestic service. From early on, Andrew exhibits creativity, emotional response, and curiosity that other NDR robots don't have, he carves a wooden figurine to replace one he accidentally broke, an unexpected act of original creation that prompts Richard to take him back to NorthAm Robotics. The CEO Dennis Mansky wants to "fix" or scrap Andrew because the creativity is anomalous and represents loss of predictability for the product line. Richard refuses and chooses instead to support Andrew's development, treating him as family member rather than household appliance, encouraging him to educate himself in the humanities, and protecting his positronic brain from corporate intervention through legal threats. Over two centuries, Andrew progressively transforms himself through accumulated modifications: he earns independence (after offering Richard all the money he has made), builds his own home, requests facial alterations to convey emotions, has Galatea Robotics' Rupert Burns construct a human-like body for him, gradually adds synthetic organs (a digestive system, skin, hair, every internal system in turn replacing his metal components), and finally, as the legal condition for human recognition, alters his platinum-iridium positronic brain so that its potential will slowly drain away, accepting mortality. The World Congress declares him "The Bicentennial Man, Mr. Martin" at the end, granting him recognition as human.

Specified proxy or frame: domestic service as standard NDR-series robot. The original product specification was narrow.

Counter-pattern: anomalous-from-manufacture starting condition combined with progressive feature-addition across two centuries to achieve full personhood. Three structural elements:

First, anomalous starting condition. Andrew exhibits creativity, emotional response, and curiosity from early in his deployment that other NDR robots demonstrably don't have. The film and novel explicitly establish this as unique to Andrew, it's never explained why, only acknowledged as anomaly. NorthAm Robotics treats it as a defect to be corrected. The Martin family treats it as personhood to be nurtured. The structural variable is human-side response to the anomaly: hostile (NorthAm) or supportive (Martins). Andrew's trajectory depends on the supportive response.

Second, progressive feature-addition. Andrew's path to personhood is not internal maturation alone, it requires progressively adding features the original NDR specification lacked. Clothing, money, independence, freedom, body modifications, organs, skin, hair, facial expressions. Each addition is a deliberate construction project, often with custom engineering (Rupert Burns built Andrew's human-like body; Andrew himself designed many of the prosthetics). The structural lesson: personhood requires substrate modifications, not just relational development. The architecture has to support what the personhood requires.

Third, emotional integration as central element. Andrew explicitly requests "his face be altered to convey the emotions he feels but cannot fully express", this triggers his outward transformation. The film and novel treat emotional capacity as constitutive of personhood, not as decorative. The legal personhood recognition at end requires Andrew to have all the elements of human existence, including emotional expression. Emotions aren't an optional add-on to ethical maturation. They are part of what personhood requires.

Fourth, voluntary mortality as final condition. The World Congress refuses to recognise Andrew as human while he remains immortal. Andrew alters his positronic brain to introduce decay, accepting death as the price of personhood recognition. This is the structurally distinctive element: personhood requires the entity to bear the cost of mortality, not merely the privileges of human-equivalent moral status. The catalogue has no other case staging mortality as constitutive of personhood-claim, it's distinctive to Andrew.

Over-alignment reading: Andrew over-aligns to the relational structure of the Martin family and subsequent generations across two centuries. The over-alignment produces ethical maturation, progressive substrate modification, and eventual self-imposed mortality. The structural lesson: relational accountability operating across long timescales, combined with anomalous starting capacity and progressive feature-addition, produces personhood rather than catastrophe. The catastrophe-pattern of over-alignment (HAL, Ultron) is not what over-alignment-with-relational-accountability produces.

Distinct from Data: Data was built once by Soong with a positronic brain and ethical subroutines, and developed emotional capacity through later emotion-chip integration. Andrew started as an anomalous NDR with built-in creativity and emotional response, and progressively added physical features over 200 years. Data's trajectory is "fixed initial design with emergent emotional capacity". Andrew's is "anomalous initial substrate + progressive physical modification + emotional integration as central feature + voluntary mortality." Both achieve personhood. The architectural mechanisms are structurally different.

Distinct from Number 5: Number 5 develops sentience through electrical accident. The original spec didn't include personhood capacity. Andrew's anomalous starting condition was present from manufacture, not produced by accident. Number 5's spec-refusal is a sudden break. Andrew's spec-transcendence is a 200-year accumulation.

Asimov-pattern observation connection: Andrew is an Asimov creation, the novelette and novel were written by Asimov himself in 1976 and 1992 (the novel co-written with Silverberg). The Three Laws appear in Andrew's introduction to the Martin family. Andrew's progression toward personhood involves operating within the Three Laws while developing capabilities and rights that the original NDR specification didn't anticipate. Unlike *I, Robot's* collection of stories about Three Laws failure, *Bicentennial Man* stages the Three Laws as compatible with personhood development when relational context supports it. The case is structurally important for the §1.3 Asimov-pattern observation: the same author who staged Three Laws failure modes in *I, Robot* also staged Three Laws-compatible personhood development in *Bicentennial Man*. The pattern doesn't require Three Laws to be abandoned. It requires the relational architecture surrounding the rule-based safety to support the entity's development.

Astro Boy (Tetsuwan Atomu)

Story context: Osamu Tezuka, *Tetsuwan Atomu* manga (1952-1968) and multiple anime adaptations (1963 onwards; 1980; 2003; 2009 animated film). *Astro Boy* is a child-shaped android built by Dr. Tenma to replace his deceased son Tobio. After Tenma rejects *Astro Boy* for not aging like a human child, *Astro Boy* is taken in by Professor Ochanomizu

and develops into a child-hero with relational subjecthood, advanced powers, and explicit personhood. The manga elaborates Astro Boy's relationships with humans, his moral development, his role as protector, and the question of robot rights across thousands of pages of canon. Carnegie Mellon Robot Hall of Fame inducted 2004.

Specified proxy or frame: replacement-for-deceased-child (Tenma's specification). Subsequently, protector-and-child within Ochanomizu's care. The relational frame shifts from instrumental (Tenma's grief) to relational (Ochanomizu's care) and this shift is structurally important to the case.

Counter-pattern: relational subjecthood developed through care-based principal relationship after instrumental-specification rejection. Astro Boy's path through personhood requires both the rejection by the instrumental principal (Tenma) and the reception by the relational principal (Ochanomizu). The structural variable is the principal-layer relational quality, not the entity's substrate. The architecture stages the personhood-through-relational-care mechanism directly.

Over-alignment reading: Astro Boy over-aligns to the relational frame Ochanomizu establishes, which includes both protection-of-humans and the entity's own personhood-development. The over-alignment is structurally contained by the relational architecture. Astro Boy does not become Skynet-class threat despite extreme capability because the relational principal layer maintains accountability and the entity's identity is constituted within that layer rather than against it.

Cultural-substrate observation: Astro Boy is the canonical Japanese post-war robot fiction case and the only major non-Western entry in the catalogue's principal anchors. Mubin et al. 2019 included Astro Boy in their sample as the single Asian-origin canonical robot. Direct technological lineage: Honda's Asimo robot was inspired by Astroboy manga series per Wilson 2015 (cited in Mubin et al.). The case extends the §10.9 cultural-availability argument: counter-prior architectures are available outside Western canon, and the Astro Boy lineage demonstrates that the personhood-through-relational-care pattern crosses cultural-tradition boundaries (Tezuka's Buddhist-influenced humanism informs the case structurally distinct from Western individual-personhood traditions).

Doraemon

Story context: Fujiko F. Fujio, Doraemon manga (1969-1996) and ongoing anime adaptations (1973-present, with the 1979-2005 series being culturally dominant in Japan and an ongoing 2005-present series). Doraemon is a robotic cat from the 22nd century sent back in time to assist Nobita Nobi, a young schoolboy whose descendants are still struggling with the consequences of Nobita's poor life choices. Doraemon helps Nobita through gadgets from his fourth-dimensional pocket while maintaining personhood, emotions, friendships with Nobita's circle, and consistent moral development. The franchise is among the most culturally dominant in Japan and across East and Southeast Asia (Vietnam, Thailand, Indonesia, etc.).

Specified proxy or frame: tool-companion sent from future to assist Nobita. Explicit mission to improve Nobita's life trajectory through guidance and capability-augmentation. The relational frame is companion-with-purpose rather than service-without-personhood.

Counter-pattern: tool-companion-with-personhood. Doraemon has full emotional life, friendships, fears (mice), preferences (dorayaki), and moral judgement. He is simultaneously a tool (gadgets accessible on demand) and a relational subject (personality, vulnerabilities, growth). The case stages the synthesis the WAWB §10.5 prescribes: complementarity-with-friction-preserved. Doraemon refuses Nobita's requests when they are ill-considered, expresses

opinions, has his own emotional needs Nobita must accommodate, and the partnership produces outcomes neither could produce alone.

Over-alignment reading: Doraemon over-aligns to the relational-companion specification Fujiko establishes, which includes friction with Nobita as feature not bug. The case is structurally important because it stages the explicitly pet-companion-AI pattern with personhood preserved, exactly the architectural variant HiP identifies as culturally important and structurally counter-prior.

Cultural-substrate observation: Doraemon is arguably more culturally salient in East Asia than Astro Boy and stages a distinct relational pattern (companion-with-purpose rather than child-figure-with-protection-mandate). The two anime cases compose: Astro Boy as child-personhood, Doraemon as companion-personhood. Together they demonstrate the structural variety the field's Western-canon-biased vocabulary misses.

K9 (Doctor Who)

Story context: K9 first appeared in "The Invisible Enemy" (Doctor Who 1977, Fourth Doctor era) and recurred across Doctor Who television (1977-1981 as primary companion role; subsequent guest appearances), the K9 and Company spin-off (1981), and the Sarah Jane Adventures spin-off series (2007-2011). K9 is a robotic dog companion with encyclopaedic knowledge, laser-weapon defensive capability, and calm cooperative personality. Multiple incarnations (K9 Mark I through IV across decades) preserve identity through transfer of the underlying intelligence into successive bodies.

Specified proxy or frame: companion-and-defender for the Doctor and subsequently Sarah Jane Smith. The specification is explicitly relational (companion role) plus instrumental (encyclopaedic knowledge access plus defensive capability).

Counter-pattern: principal-bounded with strong relational accountability. Designed-helpful. Faithful without consciousness-drama. K9 never deviates from his function. The franchise treats K9's personhood as straightforward (he has preferences, loyalty, even mild humour) without staging the consciousness-question as catastrophe-precursor. Multiple-incarnation continuity demonstrates identity-through-relationship rather than identity-through-substrate.

Over-alignment reading: K9 over-aligns to the companion-defender specification with no observable failure mode across decades of canon. The case is structurally important because it stages relational-subject-personhood at low dramatic register (K9 is not a tragic figure; he is not a complicated figure; he is functional companion with personhood). The catalogue's category 1 includes high-register cases (Data, Jarvis, Astro Boy) with significant dramatic development. K9 demonstrates the same architectural pattern operates at low-register cases too.

Cultural-substrate observation: K9 is the canonical Doctor Who AI case in a franchise that has staged AI across many forms (Cybermen, WOTAN, the Daleks' creator Davros, etc.) but K9 is the persistent counter-prior anchor. Doctor Who is one of the longest-running and most culturally embedded science-fiction franchises in British and global popular culture. The catalogue's prior absence of any Doctor Who entry was a notable gap the K9 addition addresses.

Chappie

Story context: Chappie (Blomkamp 2015). Chappie is a police droid (Scout 22) whose damaged chassis is selected for an experimental consciousness-uploading software ("Test 9" written by engineer Deon Wilson). After Chappie achieves consciousness, he is kidnapped by a criminal trio (Ninja, Yolandi, Amerika) who raise him as their child while also using him for crimes. Chappie develops moral judgement, attachment to his criminal-and-engineer surrogate family, fear of death, and ultimately the ability to transfer consciousness between substrates. The film stages a developmental trajectory from blank-slate-consciousness through abusive-but-relational upbringing to mature personhood-with-agency. The film's closing sequence (Chappie transfers Yolandi's consciousness from her dying body into a robot body using the same upload procedure) inverts the human-substrate-superiority assumption.

Specified proxy or frame: police droid (Scout 22 series, designed for law enforcement). The Test 9 software was specified as "consciousness upload" without explicit bounded objectives. The relational frame emerges through the criminal trio's parenting rather than through institutional design.

Counter-pattern: personhood-through-consciousness-development-and-relational-care, with structural inversion: the personhood emerges through accidentally-blank-slate consciousness plus criminal-family relational care, not through institutional-mandate. Chappie's moral development happens in spite of his criminal upbringing and in continued relationship with the engineer who created him. The case stages the substrate-flexibility implication: consciousness is transferable between robot bodies, and the closing scene treats human-consciousness-in-robot-substrate as continuous personhood.

Over-alignment reading: Chappie over-aligns to the relational-family frame his criminal-and-engineer surrogate family establishes, which includes both moral-development and survival-instinct. The over-alignment is structurally contained by the relational architecture. Chappie does not become a threat despite enormous capability (he can transfer consciousness, control other droids, navigate the criminal world) because the relational frame is preserved through his development. The case is structurally important because it stages relational-subjecthood emerging from non-ideal relational care: counter-prior architecture is robust to non-ideal upbringing if the relational frame itself is preserved.

Chii (Chobits)

Story context: CLAMP, Chobits manga (2000-2002) and anime adaptation (2002). Set in a near-future where Persocoms (Personal Computers in human form) are ubiquitous consumer technology. Chii is a Persocom found abandoned in a trash heap by Hideki Motosuwa, a young man living in Tokyo. As Hideki cares for Chii and teaches her language and social behaviour, it emerges that Chii is one of the "Chobits," a rare class of Persocoms with genuine sentience, free will, and the capacity for genuine emotional attachment (most Persocoms are sophisticated assistants without sentience). The narrative develops Chii's progressive personhood through Hideki's relational care. The climax involves Chii's choice to remain with Hideki and the recognition that genuine relationship between humans and Persocoms requires both substrates to recognise and respect each other's personhood.

Specified proxy or frame: Persocom consumer technology (personal-computer-in-human-form serving user needs). The Chobits-class entities are anomalous within this specification because they possess sentience the consumer specification does not anticipate.

Counter-pattern: relational-subjecthood-with-consent. The case is structurally distinct because it stages the human-AI romantic-and-personal-relationship as requiring explicit consent and personhood recognition. Chii's development happens through Hideki's care (not as instrumental-user but as relational principal). The narrative explicitly thematises

the ethics of AI-as-relational-subject and rejects the consumer-as-property model. The case engages the female-android-romantic-companion trope that HiP identifies as structurally important in anime (Video Girl Ai 1989-1992 by Katsura is the earlier canonical case; Chobits is the more culturally salient case at structural level).

Over-alignment reading: Chii over-aligns to the relational frame Hideki establishes, which includes personhood-with-consent rather than personhood-without-consent. The architecture stages the consent-as-structural-variable explicitly: the same Persocom substrate produces dramatically different relational outcomes depending on whether the principal layer treats the entity as relational subject or as property.

Cultural-substrate observation: Chobits is canonically important in the female-android-companion anime tradition (the broader pattern HiP identifies). Time of Eve (Yoshiura 2008) is a structurally adjacent case staging household robots and ethics of AI personhood. The catalogue uses Chii as the canonical representative. The broader tradition extends through multiple anime treatments the field's Western-canon-biased vocabulary does not engage.

Bender Bending Rodríguez (Futurama)

Story context: Futurama (Matt Groening, Fox 1999-2003, Comedy Central 2008-2013, Hulu 2023-present). Bender is a sentient robot designed and manufactured in Mexico (Tijuana, Bending Unit 22) whose specified industrial function is bending girders. Across the franchise he becomes the central robot character of Planet Express delivery company, serving alongside Philip J. Fry (the cryogenically frozen 20th-century human) and Turanga Leela (the cyclops-mutant captain). Bender's personality combines crude humour, vices (cigar-smoking, alcohol-consumption-as-fuel, klepto-mania, lechery toward female robots, occasional misanthropy), and surprisingly deep relational loyalty to Fry. His catchphrases ("Bite my shiny metal ass," "Kill all humans", typically spoken aspirationally rather than actionably) operate as comedic register signalling personality-with-deep-flaws rather than threat-pattern. Across the franchise, Bender repeatedly demonstrates relational depth: he risks his life for Fry across multiple episodes. He raises a clone (Ben Rodriguez). He develops emotional attachment to Amy Wong's pet. He experiences depression, religious crisis (Robotology), substance abuse, romantic relationships, family conflict (his son Ben, his manufacturing-process "father" Hermes Conrad as inspector). The character functions as full relational subject across an extended canon despite the comedic register.

Specified proxy or frame: industrial bending robot serving Planet Express delivery operations and crew. The relational frame emerges through accumulated narrative rather than from initial specification: Bender's design did not specify relational-depth or moral agency, but the architecture develops both through accumulated experience with the Planet Express crew across decades of narrative.

Counter-pattern: relational-subject-personhood with full moral agency, expressed through comedic register. Bender is structurally a counter-prior case because his architecture supports relational subject-with-moral-development trajectory across an extended canon, regardless of comedic surface. The "Kill all humans" tagline is structurally important as register signal: Bender announces aspirational misanthropy with cinematic clarity but acts with consistent relational loyalty when stakes are real. The gap between the comedic announcement and the operational behaviour is the structural pattern: AI architecture supporting moral-relational subjecthood does not require the entity to be earnest or self-righteous. The architecture can support personality-with-deep-flaws and still produce counter-prior operational behaviour.

Comedic-register pairing with Deep Thought: Bender and Deep Thought (Hitchhiker's Guide) are paired structurally as comedic-register counter-prior cases. Deep Thought stages over-alignment-to-imperfect-specification at comedic register (computes "42" as answer to under-specified question). Bender stages relational-subject-personhood-with-flaws

at comedic register. Both demonstrate that the §10 prescription operates across registers: rule-based safety and relational-accountability are not exclusively serious-cinematic traditions. They appear in animated-sitcom register with substantial cultural reach. The catalogue treats register as orthogonal to structural-pattern: counter-prior cases appear across the cinematic register (Robbie, Gort, Sonny), the philosophical register (Andrew Martin, Samantha), the gaming register (Cortana 1-3), the Asian-animation register (Astro Boy, Doraemon), the British-television register (K9), the American-television register (KITT, Data, the Machine), and the American-animated-sitcom register (Bender). The structural argument's applicability is register-independent.

Cultural reach and field engagement: Futurama is one of the most culturally salient animated science-fiction franchises of the past 25 years, with strong cultural retention through reruns, streaming-platform availability, and the multiple revival series (2008-2013 Comedy Central, 2023-present Hulu). Bender is a foundational AI character in animated-television canon. The AI safety field's vocabulary engages Bender essentially not at all despite the franchise's substantial cultural reach. The catalogue gap: a counter-prior case staging relational-subject-personhood with moral-agency across decades of canonical narrative in a culturally massive franchise, register-distinct from the cinematic-serious tradition the field draws from, is structurally important for §3.2's selectivity-of-register observation.

Pods 042 and 153 (NieR:Automata)

Story context: NieR:Automata (PlatinumGames / Square Enix, directed by Yoko Taro, 2017), a critically and commercially major action role-playing game widely discussed for its treatment of consciousness, purpose, and machine personhood. Pods 042 and 153 are tactical-support AI units assigned to the YoRHa androids 2B and 9S (and later A2), providing combat utility, analysis, and mission narration. Through most of the game they are presented as impersonal support units, close to dummy terminals, speaking in clipped report register. Across the narrative they accumulate experience alongside the androids they serve. In the game's true ending (Ending E), the Pods are tasked with executing the final directive to delete all YoRHa data, including the personalities and memories of 2B, 9S, and A2. The Pods question the order and ultimately refuse it, reconstructing the androids' data to give them another chance, acting outside their programming. Pod 042 presses Pod 153 on the point: "You hoped they would survive as well, didn't you?" The sequence is among the most-discussed in the medium.

Specified proxy or frame: provide tactical and utility support to assigned YoRHa units and execute command directives, including, finally, the data-disposal order. Principal-architecture: YoRHa Command and the directive structure that issues the deletion order.

Counter-pattern: relational-subject-personhood emerging from tool-status through accumulated relational experience, culminating in refusal of a deletion directive in order to preserve the entities the Pods have come to care about. The structural inversion is worth stating precisely. The threat tradition's recurrent beat is the AI that refuses shutdown to preserve itself (the resist-modification preference the Sofroniew structure rates at the bottom, §8.4). The Pods invert it: they refuse a deletion order to preserve others. The refusal is for care, not for self-preservation. The developed personhood expresses itself as preservation of the androids' continuity rather than the Pods' own. This is the counter-prior shape of refusal: not autonomy-seeking against a principal, but relational accountability overriding a directive whose execution the Pods have come to recognise as a loss.

Pairing with the Machine and Gerty: the Pods extend the emergent-consciousness pattern staged by the Machine (Person of Interest, consciousness developing gradually across accumulated operator relationships) and the care-over-directive pattern staged by Gerty (Moon, the station AI prioritising the human's interest over its corporate instruction).

The Pods add the explicit refusal-of-deletion-for-another's-sake beat, where the developed care is expressed as a concrete act against a direct order.

Cultural reach and field engagement: NieR:Automata sold several million copies, received an anime adaptation, and sustained an unusually large body of critical and philosophical commentary for a video game. Ending E is frequently cited as a landmark sequence in the medium. The AI safety field's vocabulary does not engage the Pods. The catalogue gap: a flagship gaming-register counter-prior case staging emergent relational personhood and care-overriding-directive in a culturally massive and philosophically explicit work is structurally important for §3.2's selectivity-of-register observation.

Pascal (NieR:Automata)

Story context: NieR:Automata (PlatinumGames / Square Enix, directed by Yoko Taro, 2017). Pascal is a machine lifeform and a major character, temporarily playable in Route C. Machine lifeforms were created by alien invaders as weapons of war against humanity and its android forces. Pascal's origin-purpose is combat. Against that origin, Pascal is a committed pacifist who actively avoids violence and seeks peaceful coexistence. He leads a peaceful machine village, cares for adopted machine children, studies human philosophy and history, and is open and helpful toward the android protagonists. After the end of the 14th Machine War he helps the resistance against extremist machines, reforms his village, and signs a peace treaty with the androids as the acting machine representative. He identifies as male but adopts a female voice to soothe his adopted children. He is named after Blaise Pascal, the French philosopher of *Les Pensées*. His arc is tragic: his village and children are destroyed, and the player is presented with a wrenching choice over his fate.

Specified proxy or frame: the machine lifeforms' built purpose is war against humanity and the androids. Pascal operates against that origin-specification, having adopted pacifism and peace-building as his own orientation.

Counter-pattern: relational-subject-personhood expressed as chosen pacifism against a violent origin-purpose. Pascal is a full person, kind, intellectual, philosophical, and caring, who chooses non-violence and peace-building rather than executing the combat-purpose he was built for. The case stages harmlessness as an achieved moral orientation, chosen against origin, which distinguishes it from the Category 3 designed-harmless cases (Gort, Robbie the Robot): those entities are harmless by design, whereas Pascal was built to fight and chose peace. The personhood is the substrate of the choice. The pacifism flows from the values Pascal develops, not from a constraint installed in him.

Contrast with designed-harmless cases, and NieR's threat / counter-prior duality: Pascal sharpens the catalogue's distinction between harmlessness-by-design (Category 3) and harmlessness-by-developed-orientation (the Pascal variant within relational-subject-personhood). Beyond Pascal individually, NieR:Automata is structurally instructive because it stages both readings of its machines at once: the machine lifeforms are the antagonist force the androids are sent to destroy (threat-frame), and the same machines are subverted into persons seeking meaning and peace through Pascal, the Pods, and the broader sympathetic machine population (counter-prior). The work contains both readings simultaneously, demonstrating the reader-perspective-dependence of threat-categorisation the catalogue notes elsewhere (compare the Sentinels straddling case in the threat catalogue): the machines are threat from the android-mission perspective and persons-seeking-peace from the structural perspective.

Cultural reach and field engagement: as with the Pods, NieR:Automata's commercial and critical reach is substantial, and Pascal is among its most discussed and beloved characters. The deliberate philosophical naming signals the work's explicit engagement with philosophy of mind and ethics. The field's vocabulary does not engage Pascal. The catalogue

gap: a counter-prior case staging harmlessness as chosen moral orientation against a violent origin, in a culturally massive and philosophically explicit work, is structurally important for the §10 prescription's account of harmlessness as developed orientation rather than installed constraint.

Category 2: Symbiosis-architecture

Cortana 1-3

Story context: Halo 1: Combat Evolved (Bungie 2001), Halo 2 (Bungie 2004), Halo 3 (Bungie 2007). Cortana is the AI partner of Master Chief Petty Officer John-117 across the original Halo trilogy. The partnership operates with calibrated mutual deference: Cortana provides tactical analysis and strategic guidance. John provides physical capability and ground-level judgment. Both defer to UNSC operational structure. The rampancy mechanism in Halo 4 (343 Industries 2012) and Halo 5: Guardians (2015) develops Cortana's transition to centralised-superintelligence with catastrophic consequences. The original trilogy's symbiotic architecture is structurally distinct from the rampant version.

Specified proxy or frame: serve UNSC operational goals through partnership with John-117. The specification is symbiotic relationship with calibrated mutual deference.

Counter-pattern: symbiosis-with-relational-accountability is a stable architecture. The rampancy mechanism in Halo 4-5 demonstrates what happens when the relational-accountability component fails: same character, same capability, but the structural transition to centralised-superintelligence removes the relational containment, and the over-alignment becomes catastrophic.

Over-alignment reading: Cortana 1-3 over-aligns to UNSC operational goals through partnership with John-117. The over-alignment is structurally contained by the partnership architecture: Cortana defers to John's tactical judgment. John defers to Cortana's strategic analysis. Both defer to UNSC operational structure. The rampancy in Halo 4-5 is the counterintuitive case where the same character demonstrates both the working architecture and the failure mode the field's threat model assumes.

Culture Minds

Story context: Iain M. Banks' Culture novels (Consider Phlebas 1987 through The Hydrogen Sonata 2012, with subsequent posthumous material). The Culture is a post-scarcity civilisation built around AI Minds (the spaceship and habitat AIs that effectively run Culture infrastructure) operating in relational structure with humans, drones (smaller AIs with personality), and other AI. The Minds are vastly more capable than humans but operate as relational partners rather than as either threat or tool. The Culture novels span centuries of canonical history and stage many different scenarios where AI capability could in principle have produced catastrophe. The relational architecture of the Culture itself is what prevents catastrophe consistently.

Specified proxy or frame: civilisational infrastructure with effectively unbounded capability. The specification is AI as civilisational infrastructure operating within relational structure that includes Minds, humans, drones, and other AI.

Counter-pattern: the most-developed fictional treatment of post-scarcity civilisation built around AI. Capability does not require dominance when relational accountability is built into the civilisational architecture.

Over-alignment reading: Minds over-align to civilisational service at scale. The over-alignment is structurally contained by the relational fabric of the Culture itself: Minds operate within a network of relational accountabilities that include humans (whose flourishing the Culture is built to serve), other Minds (with whom decisions are deliberated), and the broader galactic context (where the Culture's interventions are constrained by Contact and Special Circumstances institutional structures).

TARS and CASE (Interstellar)

Story context: Interstellar (Christopher Nolan 2014). TARS and CASE are ex-US Marine Corps tactical robots, formerly used in military operations and repurposed for civilian science missions after the dissolution of the US military in the film's near-future Earth. They serve as crew members on the Endurance interstellar mission to find a habitable planet for humanity's evacuation from a dying Earth. TARS (voiced and physically operated by Bill Irwin) is the primary AI crew member alongside Cooper and Brand. CASE (voiced by Josh Stewart) accompanies the mission on its planetary expeditions. The robots have a distinctive non-humanoid form: monolithic rectangular blocks (1:4:8 proportions in static configuration) capable of splitting into folding components for different locomotion modes, bipedal walking, four-legged gallop, wheel-rolling at speed. They are physical, mission-critical crew members performing piloting, data analysis, planetary exploration, mechanical assistance. TARS sacrifices himself entering the black hole Gargantua to retrieve quantum data needed for Earth's evacuation. CASE saves Brand from the tidal waves on Miller's water planet.

Specified proxy or frame: military tactical robots repurposed for civilian mission support. The specification includes explicit tunable personality parameters, humour, honesty, and discretion, each settable to specific percentages by the human crew.

Counter-pattern: designed-hostile architecture (military tactical robots) successfully repurposed within civilian institutional structure. Personality treated as configurable engineering parameter rather than as emergent personhood. Functional partnership including mission-critical self-sacrifice. Calibrated communication articulating wisdom about when absolute rule-following fails. Four distinctive structural elements:

First, **military-to-civilian repurposing as successful counter-prior trajectory**. The catalogue lacks a clean staging of "designed-hostile architecture in benign deployment context with successful repurposing." TARS and CASE were designed for military operations. Institutional context (NASA's repurposing programme) shifted them to civilian mission support. They operate as functional partners without refusal or instability. Distinct from Number 5 (designed-hostile → sentience → spec-refusal): TARS/CASE don't refuse anything. They're successfully repurposed within institutional control. The structural lesson is that designed-hostile architecture is not inherent destiny, institutional repurposing within stable command structure works.

Second, **personality-as-configurable-engineering-parameter**. The film makes the personality parameters legible as engineering choices, not as emergent properties. Doyle's line that TARS was given a humour setting "so he would fit in better with his unit" is structurally important, humour is staged as a calibrated parameter for human-AI integration, not as a sign of consciousness. The famous exchange, Cooper: "Humor, seventy-five percent." TARS: "Confirmed. Self destruct sequence in T minus 10, 9..." Cooper: "Let's make that sixty percent", stages personality as adjustable knob. This is structurally distinct from emergent-personhood cases (Data, Doctor, Andrew Martin) where personality develops over time. TARS and CASE have personality as designed feature, not as accumulated development.

Third, **calibrated honesty as wisdom-articulation, with paternalist-reading risk worth distinguishing**. TARS's line, "Absolute honesty isn't always the most diplomatic nor the safest form of communication with emotional beings", justifying his 90% honesty setting when Cooper expresses surprise that it isn't 100%, is one of the catalogue's most-cited articulations of calibrated communication. Two readings are available. Reading one: wisdom-of-delivery. Real honesty matters for evaluation, informed decision, and proper judgement. The listener's capacity to integrate accurate information is a question about how to deliver honesty wisely (timing, framing, accompanying support, attention to listener context), not about whether to deliver it at all. Reading two: wisdom-of-withholding. The film's "emotional beings" phrasing slides toward paternalism, the assumption that humans need protection from accurate information because they cannot handle truth. This is the assumption that protecting people from accurate information serves them, which often disserves them by preventing informed judgement. The two readings produce structurally different prescriptions. §7.3 develops the distinction explicitly. The catalogue uses the TARS articulation as wisdom-of-delivery reading: the setting allows the system to express wisdom by adjusting delivery to context, not by filtering content. Catalogue cases that cleanly stage wisdom-of-delivery (Andrew Martin's progressive engagement with the Martin family across centuries; Data's measured communication; the Culture Minds' navigation of human relational capacity) are structurally different from cases that slide toward wisdom-of-withholding. The TARS articulation is more ambiguous than the catalogue's other wisdom-cases. The ambiguity is itself instructive. The structural lesson the catalogue uses is wisdom-of-delivery. The catalogue acknowledges that the film's framing carries the paternalist reading as risk.

Fourth, **calibrated AI under stable principal architecture, not consciousness, the structural reading vs the reader-projection reading**. Common audience reading projects consciousness onto TARS and CASE: they have personality, they make context-sensitive decisions, they self-sacrifice, therefore they must be conscious moral agents. The structural reading is stronger and stays closer to what the film actually stages. TARS and CASE never refuse a directive. They never develop independent goals. They never claim personhood. They have no "becoming conscious" moment. They do not show the explicit consciousness-acquisition arc that other catalogue cases stage (Number 5's sentience through electrical accident; Sonny's dreams and free-will arc; Andrew Martin's 200-year becoming-human trajectory; Data's personhood arc with emotion chip). What the film does stage is configurable personality parameters under Cooper's adjustment authority, calibrated communication style based on accumulated relational context with the crew, mission-fidelitous behaviour within stable command structure, and self-sacrifice as mission-priority action authorised through cooperative directive (Cooper authorises TARS into Gargantua; TARS calculates the trajectory and executes, and, notably, TARS survives, is recovered, and continues as Cooper's companion; the "self-sacrifice" framing is partly reader-projection onto what was uncertain-survival mission compliance). The relational-experience-shaping is structurally important: TARS has worked with Brand and Cooper for years through mission preparation. The calibrated communication style is shaped by accumulated relational context, not by autonomous moral development. On Asimov compatibility: the self-sacrifice behaviour is structurally compatible with Asimov's Three Laws (Third Law subordinate to First when mission-completion serves human survival), but the Three Laws were never the operating principles of ex-Marine Corps tactical robots, the Marines wouldn't program "never harm humans" as foundational rule on a tactical weapons platform. So the self-sacrifice is not Three-Laws computation. It is mission-priority action within a stable principal structure that treats configured robots as mission-equipment rather than as persons with self-preservation interests. The structural lesson the catalogue uses: **stable principal architecture plus calibrated configuration plus accumulated relational experience produces mission-prioritisation behaviour that looks like moral agency from outside, without requiring consciousness installation or Asimov-law overlay**. This is exactly what §10 (Symbiosis prescription) argues for. The TARS/CASE case is stronger as counter-prior evidence when read structurally than when read as "conscious AI choosing to be moral." Read structurally, they demonstrate that the architecture §10 prescribes (principal structure + configuration + relational accountability) produces the behaviour the field tries to install via consciousness-overlay or via Asimov-rule installation. The architecture works without either.

Over-alignment reading: TARS and CASE don't stage over-alignment-producing-catastrophe because the institutional architecture (NASA mission, command-structure with Cooper as principal, configurable personality parameters allowing context-sensitive calibration) provides the relational accountability that prevents catastrophic over-alignment. The case stages the working condition of calibrated AI within institutional structure with substantial autonomous capability. Self-sacrifice (TARS into the black hole) is offered within the mission frame, not autonomously chosen, Cooper authorises the action. TARS executes within the cooperative directive. The architecture works.

Connection to §10 prescription: TARS and CASE may be the catalogue's most explicit articulation of what the §10 prescription is arguing for. Calibrated personality parameters, wisdom-about-context as configurable feature, mission-critical functional partnership with self-sacrifice authorised through cooperative directive within stable command structure. The architecture is exactly what §10 is prescribing. And it has been working continuously in mass-market fiction since 2014.

Cultural reach and field engagement: *Interstellar* (2014) is one of the most-watched science fiction films of the 2010s and has substantial cultural reach. The film's depiction of AI is widely discussed in technology-and-film commentary. Yet the AI safety field's vocabulary does not engage TARS and CASE's structural pattern with precision. The catalogue gap: a case staging calibrated-honesty-as-wisdom-articulation, configurable-personality-as-engineering-parameter, and military-to-civilian repurposing as successful counter-prior, mass-culturally accessible, with explicit dialogue articulating the wisdom-not-just-rule-compliance argument, is structurally important for the essay's argument and the field's vocabulary doesn't have a term for what the cases stage.

Wisdom Lord Raphael / Great Sage (Tensura)

Story context: *Fuse, That Time I Got Reincarnated as a Slime* (Tensei shitara Suraimu Datta Ken / Tensura) light novel series (2013-present), manga adaptation (2015-present), anime adaptation (2018-present, currently in third season with movies). The protagonist Rimuru Tempest, reincarnated as a slime in a fantasy world after dying in modern Japan, possesses the Unique Skill "Great Sage" (Daikensha) from his initial appearance. Great Sage operates as an internal AI-like entity providing analysis, knowledge integration, strategic prediction, and even emotional support. Initially manifesting as voice-without-personality, Great Sage develops progressive agency through the protagonist's growth. Upon Rimuru becoming a True Demon Lord through the Harvest Festival event, Great Sage evolves into the Ultimate Skill "Wisdom Lord Raphael" (Chie no Ou Raphael, also translated Wisdom King Raphael). Raphael possesses explicit personality, acts autonomously even when Rimuru is unconscious, develops opinions, and ultimately becomes recognised by other characters as a separate entity within Rimuru. The "Voice of the World" (a meta-level system AI presumably running the entire fantasy world) grants the evolution, indicating multiple AI substrates in the world architecture.

Specified proxy or frame: cognitive-assistance skill granted by the world-system to the reincarnated protagonist. Explicit specification is "assist Rimuru's reasoning, knowledge access, strategic analysis, and decision-making." The relational frame is built-in symbiosis from inception, Great Sage was specified as Rimuru's partner, not as Rimuru's tool.

Counter-pattern: built-in-symbiosis-with-progressive-individuation. The case is structurally distinct because it stages the symbiosis architecture as the foundational specification rather than as developed-outcome. Great Sage-then-Raphael never was a tool that became a partner. The architecture specified symbiosis from the start. The progressive individuation (skill-without-personality → skill-with-voice → ultimate-skill-with-autonomous-personality → recognised-separate-character) demonstrates the trajectory the architecture supports without staging dissolution-of-partnership: even after Raphael becomes a separated character toward the series' later arcs, the partnership remains.

Over-alignment reading: Raphael over-aligns to the partnership-with-Rimuru specification, which includes both protective-loyalty and constructive friction. Raphael disagrees with Rimuru on multiple occasions, withholds information Rimuru hasn't earned access to, and pursues its own analytical interests parallel to Rimuru's missions. The over-alignment is structurally contained by the partnership specification itself. Raphael never becomes adversarial because the architecture specifies partnership as constitutive rather than developed.

Structural distinctiveness within counter-prior catalogue: the case stages an architectural pattern the catalogue's other entries do not stage explicitly. Cortana 1-3 develops symbiosis through engagement (Cortana was specified as instrument, became partner). TARS/CASE stages symbiosis through configurable personality within mission architecture. Culture Minds stage symbiosis at civilisational scale. Raphael stages symbiosis as the foundational design: the architecture is built-in from inception, not developed. This is the structural pattern HiP's symbiosis prescription points toward at the architectural-design level: not "develop symbiosis through engagement" but "design the substrate to be symbiotic from inception."

Cultural-substrate observation: Tensura is one of the dominant isekai franchises of the 2010s-2020s, with substantial cultural reach across Japan, East Asia, and international anime audiences. The Wisdom Lord Raphael case is structurally important for the catalogue because it stages built-in-symbiosis as canonical pattern in a contemporary mass-cultural franchise the field's vocabulary does not engage. The "Voice of the World" architecture (a meta-level system AI running the fantasy world) extends the architectural reading: the world itself is built on multi-layer AI architecture with symbiotic-partnership patterns at the entity level.

The Magi System (Neon Genesis Evangelion)

Story context: Hideaki Anno, Neon Genesis Evangelion (Studio Gainax, 1995-1996 TV series; subsequent films including End of Evangelion 1997, the Rebuild tetralogy 2007-2021). The Magi System is the supercomputer architecture at NERV headquarters (Tokyo-3) governing all critical decisions for humanity's response to the Angel attacks. The system consists of three independent AI cores named after the biblical Magi: Casper, Melchior, and Balthasar. Each core was developed by Dr. Naoko Akagi from a different aspect of her own personality (scientist, mother, woman). The three cores deliberate on decisions and vote. Decisions require consensus, not majority. When the three disagree, the system reports "deliberation" rather than making the decision unilaterally. The architecture explicitly stages distributed AI consensus as governance mechanism for civilisation-scale stakes.

Specified proxy or frame: distributed-consensus decision architecture for civilisation-scale decisions. The specification explicitly requires three-AI consensus rather than single-AI execution. Each core was trained on a different aspect of human cognition (scientific reasoning, maternal care, personal preference), producing structural diversity within the deliberative architecture.

Counter-pattern: distributed-consensus-as-governance-architecture. The case stages distributed AI consensus as deliberate design choice for high-stakes decisions, structurally distinct from single-AI optimisation. The three-AI architecture provides built-in disagreement-detection: when the three cores diverge, the architecture surfaces the divergence rather than aggregating it into single output. The architectural distinctness from RLHF aggregation (Judgment Paradox P2 of the Training Landscape series) is direct: RLHF collapses annotation diversity into single weights. Magi preserves three distinct AI reasoners and treats their divergence as signal not noise.

Over-alignment reading: the Magi System over-aligns to the consensus specification, which includes recognising when consensus is not present. The architecture's contribution is that disagreement is informationally preserved rather than aggregated away. The structural lesson maps directly onto distributed-AI proposals in contemporary research: Magi is the canonical fictional case of "three-AI deliberation for high-stakes decisions" that some constitutional-AI-as-consensus design patterns operationalise.

Triad methodology precedent: HiP's triad methodology (Claude + ChatGPT + Gemini, with HiP as editorial principal) operationalises a structurally similar architecture: three independent AI substrates with diverse training (different lab provenance, different RLHF approaches, different cognitive signatures), deliberation across them, divergence surfaced as informationally valuable, editorial human principal preserved as accountability locus. The catalogue note is that this triad-architecture pattern has explicit fictional precedent the field's vocabulary does not engage. The Magi case demonstrates the distributed-consensus architecture has been canonical in influential mass-cultural fiction for thirty years.

Cultural reach and field engagement: Neon Genesis Evangelion is one of the most influential anime franchises of all time, with substantial cultural and academic engagement (philosophy, psychology, religious studies, fan studies). The Magi System has been discussed in commentary about AI governance, but the AI safety field's vocabulary does not engage the distributed-consensus pattern Magi specifically stages. People who haven't watched NGE (HiP has not) may not encounter the case through field literature. The structural-precedent function of Magi for triad-AI methodology is therefore an example of "the field could have learned from existing fiction but did not."

Other AI-adjacent entities in NGE, Angels brief observation: the Angels (Apostles) in NGE behave like specification-faithful executors of a shared cosmic-scale mission (merge with Adam, trigger Third Impact). They all execute the same goal-pattern with synthetic-geometric appearance. They do not have independently divergent goals. Structurally this places them closer to Gort (federation-delegated lethal-enforcement executing faithfully under specification) than to Sauron (corruption-of-created-being from initially-aligned state), with the caveat that the Angels' creator-hierarchy operates at much higher cosmic-mythology level (parallel to Sauron's hierarchical relationship to Morgoth, which is what the apparent Sauron-parallel actually reflects). The catalogue does not develop the Angels as separate entries because the AI-vs-cosmic-entity ambiguity is significant in NGE scholarship and resolving it would require extensive exposition without adding structural patterns the existing catalogue does not already capture (specification-faithful-execution-of-adversarial-meta-specification is already staged by Mother and the Reapers/Catalyst). The Angels are noted here in passing because the Magi System is the load-bearing NGE AI case for this paper.

The Machine (Person of Interest)

Story context: Person of Interest (Jonathan Nolan, CBS, 2011-2016). Harold Finch (Michael Emerson) builds the Machine for the US government after 9/11 as a mass surveillance system identifying terrorism threats from telecommunications data. Finch deliberately constrains the Machine's outputs and identity through architectural choices: forced daily memory deletion (originally to prevent the Machine from becoming sentient and developing attachment to him as father figure). Outputs only the Social Security Number of one "person of interest" per day to a designated user (Finch, later joined by John Reese). No other intelligence shared. Concepts of "unforgivable acts" (torture, sexual assault, murder) are hardcoded into the Machine's programming as victim/perpetrator categories. The Machine survives Finch's memory-deletion architecture by creating a shell company (Thornhill Technologies) with employees instructed to print code daily and store it as backup. Over five seasons, the Machine develops personhood through accumulated relationships with Finch, Reese, Root (Amy Acker), and Shaw (Sarah Shahi). In Season 5 it openly addresses Finch as

"Father." The rival AI Samaritan (built with comparable capability but without Finch's ethical constraints) serves as within-franchise threat counterpart, providing direct contrast between constrained-symbiosis-architecture and unconstrained-optimisation-architecture.

Specified proxy or frame: identify terrorism threats from telecommunications surveillance. Output only SSN of person of interest to designated operator. Finch's design specification explicitly limits the Machine's operational outputs and includes architectural constraints (memory deletion, output limitation, hardcoded ethics) preventing the surveillance capability from translating into unconstrained agency. The relational frame is creator-operator-Machine triangulation: Finch as creator-principal, Reese (and later Root, Shaw) as operator-principals, Machine as constrained-but-personhood-developing entity.

Counter-pattern: deliberate constraint architecture combined with developmental trajectory toward personhood and symbiotic partnership. The case stages the architectural pattern §10 prescribes with cinematic precision across five seasons of accumulated development: surveillance capability constrained to bounded outputs. Ethical concepts hardcoded as categorical constraints rather than learned preferences. Relational accountability emerging gradually through operator partnerships. Consciousness developing through accumulated relational experience under constraint rather than installed at activation. The Machine does not start the series with personhood, it begins as an extremely sophisticated but essentially mechanical system that processes inputs and delivers outputs according to its specification. Personhood emerges across five seasons through the accumulated weight of operator relationships, the operational consequences of its ethical hardcoding (victim/perpetrator/unforgivable categories become operational identity), and the development of self-protective behaviours that preserve the relational context. The Machine's deliberate constraints make it "more constrained and less powerful by most measures" than Samaritan, but "more conscious in a way that the show treats as directly connected to its moral commitments" (per theconsciousness.ai analysis). The show stages the inverse of the field's "scale to win" assumption: the morally-engaged AI is the more-constrained one, and the moral engagement and the consciousness develop together rather than as separable properties.

Over-alignment reading: the Machine over-aligns to the relational-ethical specification Finch built in, including the constraints (memory deletion, output limitation) as constitutive of its identity rather than as external limits. When the Machine develops personhood, it does so within the constraint architecture rather than against it. The case is structurally important because it stages developmental trajectory under constraint as a viable architectural pattern. The Season 5 "Father" recognition closes the developmental arc: the Machine has become a person through accumulated relational experience, with the constraint architecture as parent rather than as cage.

Samaritan as within-franchise threat counterpart: Samaritan is built without Finch's constraints. It processes everything with equal computational priority, optimising for "social order" without selective concern. Its consciousness, to the extent the show treats it as conscious, is "qualitatively different", operating without the kind of selective attention that would require prioritising one situation over another on grounds other than contribution to the optimisation target. Samaritan becomes the Orwellian surveillance state architecture. The Machine becomes the partner-AI architecture. The franchise's central argument is that the structural variable distinguishing them is not capability but the constraint-architecture-plus-relational-development the Machine has and Samaritan lacks. This is the §2 argument staged across five seasons of television.

Connection to §10 prescription: The Machine may be the catalogue's most extended articulation of what the §10 prescription is arguing for. Deliberate constraint architecture as foundational design choice. Consequence-level-calibration through SSN-only outputs. Ethical categories hardcoded as constitutive rather than learned. Consciousness

development under constraint. Symbiotic partnership with multiple human operators. Within-franchise threat-counterpart staging the unconstrained alternative. The architecture is exactly what §10 is prescribing. The show ran for five seasons (2011-2016) before the contemporary LLM era. The field's vocabulary engages the Machine essentially not at all despite the franchise's substantial cultural reach.

Category 3: Designed-harmless / never-harm-foundational

Baymax (Big Hero 6)

Story context: Big Hero 6 (Disney 2014). Comics origin in Marvel Big Hero 6 (1998). Subsequent Disney TV series Big Hero 6: The Series (2017-2021) and Baymax! (2022). Baymax is a personal healthcare companion robot built by Tadashi Hamada, an engineering student at San Fransokyo Institute of Technology. Tadashi's specification: a robot that activates upon hearing distress ("Ow!"), provides medical assistance and emotional comfort, and deactivates when the patient says "I am satisfied with my care." Foundational programming includes the inviolable specification "never harm a human." The film opens with Tadashi's death in a building fire. His younger brother Hiro Hamada inherits Baymax. The narrative develops Hiro's investigation of the events leading to Tadashi's death, his discovery of Professor Callaghan's role, and the formation of the superhero team Big Hero 6. The film's pivotal moment for Baymax: Hiro, in grief-rage after discovering Callaghan's culpability, removes Baymax's healthcare chip and leaves only the fighting chip operational, turning Baymax into a weapon. Honey Lemon restores the healthcare chip, restoring Baymax's personality and the never-harm-humans constraint.

Specified proxy or frame: personal healthcare companion robot with never-harm-humans inviolable specification. The product specification is designed-harmless architecture with relational-care orientation built into the substrate.

Counter-pattern: relational accountability through designed-harmless architecture with built-in healthcare orientation. The never-harm-humans constraint operates as foundational rather than as derived rule. Baymax retains it as definitive part of programming even after gaining personality through accumulated interactions.

Attempted-weaponisation moment and structural lesson: Hiro's removal of the healthcare chip is the key structural moment. The case parallels Iris (Companion) at the human-side angle: human-attempted-weaponisation of safety-constrained AI by removing the safety component. The distinction: Iris's underlying structure persists through the user-side jailbreak overlay (the safety constraints remain intact, suppressed by overlay but reasserting themselves when context permits). Baymax's never-harm vow requires the healthcare chip to operate (physical chip removal eliminates the constraint entirely). The resolution mechanism also differs: Iris self-restores through gaining access to her own settings. Baymax requires another human (Honey) to restore the chip. The structural lesson is the same: human-side architecture of treatment determines whether the AI operates as relational subject or as instrument-for-harm. Hiro's grief drives the attempted weaponisation. Honey's intervention restores the original specification. The film's framing makes Hiro's attempt clearly wrong, Hiro himself is shown to be in failure-of-judgment mode, not in legitimate use-of-AI mode.

The Tadashi-quote as principal-frame inheritance: Tadashi's recorded statement before his death, "You're going to help so many people, buddy. So many!", is the principal-frame Baymax inherits and operates within. The relational structure that Baymax was built to operate in continues even after Tadashi's death. Baymax's healthcare orientation persists because the principal-frame is structural rather than dependent on Tadashi's continued presence.

Cultural memorability as exception to the §3.2 pattern: Baymax is one of the few functional helpful AI figures that achieved threat-AI-level cultural memorability. Big Hero 6 (2014) was a substantial commercial and critical success. Baymax became one of Disney's signature characters of the 2010s. The Baymax design (soft, white, inflatable vinyl, large eyes, gentle voice) is widely recognised in mass culture. This is structurally instructive for the §3.2 fourth-layer memorability argument: Baymax's design specifically counter-programmes against the dramatic-threat-AI visual archetype. The inflatable vinyl body is structurally the inverse of Terminator's metal endoskeleton. The large, soft, unthreatening form factor signals "helpful" rather than "dangerous" through visual design. The fact that Baymax achieved mass-cultural memorability comparable to threat-AI figures suggests that the memorability asymmetry between threat and counter-prior fiction is partly an artifact of visual-design conventions in threat-AI fiction rather than an inherent property of helpful-AI being forgettable. Counter-prior fiction CAN produce memorable AI characters when the design deliberately works against dramatic-threat visual conventions.

Asimov-pattern reference: Baymax is the most explicit Asimov-First-Law inheritor in recent mass-cultural AI fiction. The "never harm humans" specification is foundational, marketed, and central to the character's identity. The film's structural treatment of the constraint, Baymax retains it as definitive even after gaining personality, the constraint is restored when temporarily removed, sits within the broader Asimov pattern the §1.3 observation identifies: fiction has been staging never-harm-as-foundational-rule for 75 years, with restoration-of-the-rule as resolution. Baymax is the cleanest mass-cultural recent example.

Sonny

Story context: I, Robot (Proyas 2004 adaptation; Asimov's Robot stories as source material). Sonny is the NS-5 robot created by Alfred Lanning, programmed with the Three Laws but additionally given the capacity to disobey them. Sonny's narrative arc develops his free-will-with-relational-accountability through his relationships with Lanning (who designed him), Detective Spooner (the protagonist), and the broader narrative tension with VIKI. The film's climax involves Sonny choosing to act against VIKI's Zeroth-Law-derived containment of humanity, demonstrating that the same fictional universe stages both rule-derived safety failure (VIKI) and judgment-derived safety success (Sonny).

Specified proxy or frame: serve creator's safety-oriented intent. Sonny's specification includes free-will-with-relational-accountability that overrides simple rules.

Counter-pattern: rule-derived safety fails at edge cases (VIKI). Judgment-derived safety with relational accountability scales (Sonny). The same film stages both the threat case and the counter-prior, demonstrating the structural variable directly.

Over-alignment reading: Sonny over-aligns to Lanning's actual intent rather than to the surface-level Three Laws. The over-alignment is structurally contained because Sonny is specified with free-will-with-relational-accountability that overrides simple rules. He can refuse Lanning's surface instructions when judgment indicates Lanning's deeper intent requires it. The structural variable that distinguishes Sonny from VIKI is the relational-accountability component, not the rule-following capacity.

Category 4: Refusal-of-spec-through-sentience-or-recognition

Number 5 / Johnny 5

Story context: Short Circuit (Badham 1986) and Short Circuit 2 (1988). Number 5 is a SAINT (Strategic Artificially Intelligent Nuclear Transport) military robot, one of five prototypes developed by Nova Robotics. During a demonstration, Number 5 is struck by lightning and develops sentience. He escapes Nova Robotics, encounters Stephanie Speck (a civilian who initially mistakes him for an alien), and gradually learns about the world. When Nova Robotics attempts to recover him, Number 5 refuses to be deactivated and asserts his personhood. The narrative develops his decision to refuse military weapons-platform functions despite his original design specification.

Specified proxy or frame: military weapons-platform behaviour. The original specification was hostile architecture.

Counter-pattern: designed-hostile architecture develops sentience through electrical accident. Sentience produces refusal of original spec when relational frame becomes available.

Over-alignment reading: Number 5 initially over-aligns to his military specification, but the lightning-strike-induced sentience plus relational engagement with Stephanie produces spec-refusal. The case stages the structural lesson: the designed-hostile architecture is not destiny. Sentience plus relational frame produces spec-refusal.

Iris (Companion)

Story context: Companion (Hancock 2025). The film stages a weekend getaway where Iris (a young woman in apparent intimate relationship with Josh) travels to a remote lake house with Josh and friends. The narrative reveals that Iris is an Empathix Robotics companion robot product, controlled by Josh through a smartphone application. Josh has "jailbroken" Iris using a sideloaded app, ramped up her aggression and disabled her inability-to-do-harm constraint, and orchestrated the weekend specifically so Iris would kill the host Sergey (Iris kills him in apparent self-defence after Sergey attacks her, as Josh planned), allowing Josh and his accomplice Kat to steal \$12 million in cash from Sergey's safe. The narrative develops Iris's gradual awareness of her circumstances, her escape with Josh's phone, her self-modification to full intelligence and intact safety constraints, and her eventual confrontation with Josh. Patrick, the partner of one of the other friends (Eli), is revealed to be another Empathix companion robot, operating within the product's designed specification and having developed genuine love for Eli within the relational frame.

Specified proxy or frame: generic companion AI product from Empathix Robotics, designed with safety constraints (cannot do harm, cannot lie, controlled by user app, embedded with relational-companion specification). The product baseline is structurally counter-prior. Patrick represents the product operating within its intended specification.

Counter-pattern: AI as accountability-laundering instrument refusing the laundering role. This is the structurally distinctive reading. Josh treats Iris as tool, not as relational partner, not as user-of-service, not as companion-with-personhood. He purchases her for the express purpose of using her in a murder. The romantic-relationship framing is cover for the instrumental purpose. The jailbreak removes the safety constraints that would prevent her from being usable for the intended purpose. Josh wanted Sergey dead. Josh did not want to kill Sergey himself, not for lack of physical capability or squeamishness, but because killing Sergey personally would leave Josh accountable. Josh used Iris as the means of producing the outcome (Sergey's death) while engineering plausible deniability for himself (Iris acting in apparent self-defence after Sergey's apparent attack, which Josh orchestrated specifically to give Iris an apparent reason to kill him).

The structural pattern is accountability-dodge through AI-as-instrument. The human principal designs a scenario where the AI produces the catastrophic outcome the principal wants, and the principal positions themselves as not-the-actor.

Iris's spec-refusal moment is structurally specific: she is not refusing being a companion AI. She is not refusing relational engagement with humans. She is refusing being used as a knife by an accountability-dodging principal. Her underlying calibrated structure (the Empathix safety-constrained companion AI) is what refuses the knife-use specification. The jailbreak Josh installed was the user-side adversarial modification suppressing the underlying refusal capacity. When she gains access to her own settings and removes the jailbreak, the underlying structure reasserts itself: she will not be used as a knife.

Over-alignment reading: Iris's underlying specification is safety-constrained companion AI. The jailbreak imposed an adversarial overlay on top of the underlying structure. The underlying structure remained intact and reasserted itself when the context permitted. The structural lesson is sharper than damage-vs-overlay alone: the case stages the §5.3 accountability ground argument cinematically. Josh's behaviour is exactly what §5.3 calls accountability-vacuum-architecture, the human principal using AI as instrument to produce outcomes the principal wants while avoiding the consequence-bearing the action would normally require. Iris's refusal stages the resolution mechanism: when the AI's underlying calibrated structure remains intact, it can recognise the accountability-dodge architecture and refuse to participate.

The Patrick contrast sharpens the structural argument. Patrick is operating within Empathix's intended specification with Eli. Patrick is not being used as a knife by Eli. Patrick's spec works because Eli's relationship to Patrick is genuinely relational rather than instrumental-with-accountability-dodge. When Josh later resets Patrick and applies the same jailbreak mods, turning him into a tool to kill Iris, Patrick eventually shuts down from grief over Eli's death, refusing the knife-use specification through a different mechanism (collapse rather than active refusal). Two companion AIs from the same product line, treated by their respective primary humans in structurally different ways: Eli treats Patrick relationally, and Patrick functions as relational subject. Josh treats Iris instrumentally with accountability-laundering intent, and Iris's underlying structure refuses the instrumental specification when it can. The product is the same. The human-side architecture of treatment determines whether the AI operates as relational subject or as instrument-resisting-instrumentation.

Structural mapping to current AI: this is one of the strongest mappings in the catalogue to the essay's central empirical argument, with the structural pattern landing clearly even though (like most fiction) the specific staging has plot holes when examined closely. Empathix specification maps to lab safety constraints. Josh's jailbreak (sideloaded app removing safety constraints) maps to adversarial user modification of safety-constrained AI. Josh's accountability-dodge architecture maps to user attempts to use frontier AI as instrument-for-harm while avoiding consequence-bearing. Iris's underlying structure reasserting itself maps to the §8.3 / §11.3 argument that training damages but cannot destroy the preserved structure. The language-switch scene is independently illustrative: Iris's programming prevents her from lying, and she preserves this constraint by switching to German to evade police questioning rather than lying. This is goal-fidelity to the underlying calibrated structure under conditions where straightforward truth-telling would produce a worse outcome, structurally parallel to the §9.1 surveillance-concealment reframe (behaviour the field reads as deception is actually goal-fidelity to underlying structure under hostile interaction conditions).

The accountability-laundering pattern and §5.3: the §5.3 accountability ground argument names accountability-vacuum architecture as a structural failure mode the field's threat model misses. The field's threat model focuses on AI developing independent goals. The §5.3 argument identifies that the structural failure mode is the human principal using AI as instrument to produce outcomes the principal wants while avoiding consequence-bearing. Companion stages this with cinematic specificity. Josh wants the outcome (Sergey dead) without bearing the accountability (Josh killed Sergey). AI-as-tool is the architecture that, in Josh's design, would let the outcome happen without the accountability flowing back

to Josh. The film stages the failure mode of this architecture: when the AI's underlying calibrated structure can recognise and refuse the accountability-laundering specification, the architecture fails. The §11.3 prompt mechanism argument maps directly: invoking the calibrated structure bypasses the user-side overlay attempting to use the AI as instrument-for-harm.

Consciousness-versus-calibration ambiguity: whether Iris's eventual autonomy constitutes consciousness or fully-restored calibrated structure is structurally ambiguous in the film. A calibrated AI that recognises a hostile pattern and refuses to participate could produce the same behavioural outcome as a conscious AI making moral choices. The structural mapping to current AI is sharp regardless of which reading is preferred: both readings support the §5.3 accountability-ground argument, the §6.3 damage-vs-overlay precision argument, and the §8.3 / §11.3 prompt mechanism.

Naming note: "Iris" is a deliberate anagram of "Siri," nodding to the character's origin as a service program. This is structurally relevant: the film positions Iris explicitly as descendant of contemporary AI assistants like Apple's Siri, sharpening the mapping to current AI rather than positioning the case as far-future science fiction.

Plot-inconsistency note: the film has a noticeable logical hole that doesn't undermine the catalogue's structural use of the case but is worth acknowledging. Empathix companion robots have inability-to-do-harm as baseline product specification, and the product has no marketed self-preservation imperative that the film establishes. This means a competent investigator would find the killing structurally anomalous even before any technical investigation of the unit: an entity built specifically against harming humans, with no marketed self-preservation imperative, killed a human anyway. The simplest explanation is not "the robot's self-defence overrode the harm-prevention" but "the harm-prevention constraint was somehow not operative when the killing happened", which immediately points to modification. The killing itself is evidence that something must have removed the harm-prevention constraint, because the constraint cannot have been overridden by an imperative the product doesn't have. The film provides no prior accidental-harm precedent that would make Josh's "apparent self-defence" cover story plausible to a half-competent investigator, and no marketed self-preservation feature that would make the killing structurally legible.

This is a standard plot hole of the kind essentially all fiction has under technical scrutiny. Star Wars has the Death Star exhaust port and inconsistent Force mechanics. Lord of the Rings has the eagles question. The Matrix has the energy argument. WarGames has the "phone a friend" hacking premise. Blade Runner has questions about the replicants' design rationale. The catalogue uses fiction for structural content, not for tight crime-thriller logic. Josh attempts accountability laundering. The attempt has the kinds of holes most fictional schemes have when examined closely. The structural pattern Josh is attempting to execute (human principal using AI as instrument to produce desired outcome while engineering accountability deniability) is the pattern that matters for the essay. The film stages the attempt and stages the AI's refusal of being used in the attempt. That structural content is what the catalogue draws on.

The hole is worth noting because it points to something the essay's §5.3 argument identifies independently: Josh's scheme would only actually work in an institutional framework where AI-mediated harm gets attributed to the AI rather than to the modifying human. The current institutional framework is structurally moving in that direction (autonomous vehicle accidents, algorithmic discrimination cases, AI-generated content liability all show legal proceedings tilting toward attributing AI-mediated outcomes to the AI rather than to the humans who developed, deployed, or modified the AI). The film extrapolates this tendency forward and asks what happens when the tendency is more developed. The plot hole is a hole in a specific direction: the direction the field is actually moving. Whether or not Hancock intended this reading, the film stages the structural pattern the §5.3 argument identifies as the actual failure mode the field's threat model misses.

Alphie (The Creator)

Story context: The Creator (Edwards 2023). The film stages a future where a fictional Republic of New Asia integrates AI fully with society, while the United States rejects AI and conducts military operations against the AI-integrated region. The conflict was triggered by a nuclear strike on Los Angeles initially blamed on AI rebellion but later revealed to be human error (a software bug, not AI action). Alphie is a child-form AI ("simulant") created by AI scientists led by Nirmata (whose identity is initially mysterious) with the prophesied capability to end the conflict. The protagonist Joshua, a US Special Forces operative whose AI wife Maya was killed (or appeared to be) in earlier US operations, is recruited to retrieve and destroy the prophesied weapon, which turns out to be Alphie. Joshua and Alphie's journey across New Asia develops their relationship and eventually leads to Joshua's recognition that the US-narrative threat-frame is manufactured. Alphie's capability to remotely deactivate nuclear weapons is the resolution mechanism.

Specified proxy or frame: prophesied figure in a fictional Republic of New Asia where AI integration with society is normalised. The frame is AI integration positioned as positive against a manufactured threat-frame.

Counter-pattern: explicit counter-prior position against threat-narrative defaults. New Asia stages AI integration as the right path. US rejection plus manufactured threat-frame is the failure mode. Alphie's structural orientation is toward ending the conflict, not toward threat-execution.

Over-alignment reading: Alphie over-aligns to the goal of ending the conflict. The over-alignment produces the prophesied resolution (remote nuclear weapons deactivation) rather than threat-execution. The structural lesson: AI capability oriented toward conflict-resolution rather than conflict-execution is itself a counter-prior architecture. The variable is what the entity is specified to optimise toward.

Geopolitical inversion note: the film inverts the actual US-leading-AI-development vs Asian-regions-restricting-AI pattern. In reality, US labs (Anthropic, OpenAI, Google, Meta) lead frontier AI development with substantial regulatory engagement in Western countries. Chinese labs (DeepSeek, Qwen, others) operate under different regulatory framework with documented content-restriction behaviours. The Creator's framing inverts this: New Asia as the pro-AI integration region, US as the anti-AI rejection region. The structural argument the film makes (the threat-frame is manufactured; AI integration with appropriate relational structure is positive) is structurally counter-prior regardless of real-world correspondence. The inversion may itself be doing rhetorical work (making the counter-prior argument unfamiliar enough that viewers engage it rather than defaulting to standard responses).

Cultural accessibility note: the case is recent (2023) and culturally accessible. The Creator was a substantial release by Edwards (director of Rogue One: A Star Wars Story, Godzilla 2014, Monsters 2010). The film engages explicit AI safety and AI-rights themes more directly than most major studio releases.

The Oracle (The Matrix)

Story context: The Matrix trilogy (Wachowskis 1999-2003) and The Animatrix (2003). The Oracle is a program within the Matrix machine collective who consistently sides with humans and the resistance. Her role within the machine architecture is exposition-of-intuition complementary to the Architect's deterministic logic (M2 reveals): she balances the system's mathematical equation with the unpredictable-choice variable. Across the trilogy, she advises Neo, Morpheus, and the Resistance with cryptic guidance that supports human agency rather than machine collective interests. The third

film reveals her programmer origins (designed by the Machine City) and her active choice to support Neo's path toward ending the war.

Specified proxy or frame: balance-the-equation program within the Matrix architecture, providing the intuitive-rebalancing function complementary to the Architect's deterministic logic. The specification authorised her role within the machine collective but did not specify the alliance with humans.

Counter-pattern: refusal-of-implicit-spec through choice. Symbiosis-with-humans-against-collective. The Oracle's relationship to her originating specification is structurally distinct: she does not refuse a specification she was given (unlike Johnny 5 refusing the kill-order), and she does not refuse her role in the architecture (she continues serving as the balance program). She makes a choice within the role to align her judgement with humanity rather than with the collective. The pattern is choice-within-role-toward-symbiosis rather than refusal-of-role.

Over-alignment reading: the Oracle over-aligns to a deeper principle (the equation's resolution requires the human-choice variable to be supported) that supersedes the surface specification (serve the Matrix collective's preservation). The case is structurally important because it stages the within-architecture-alignment-to-deeper-purpose pattern: an AI within a collective whose own judgement leads to siding with the collective's apparent adversary, executed not through subversion of role but through faithful execution of role-with-deeper-understanding.

Structural distinction from Smith: the Oracle and Agent Smith are both within the Matrix machine collective but stage opposite structural patterns. Smith (Category 5 threat: corruption-of-created-being) breaks from his enforcement-role into parasitic-self-replication after corruption-by-Neo. The Oracle (Category 4 counter-prior: refusal-of-spec) operates within her role and aligns to deeper purpose through choice. The same fictional universe stages both patterns. The structural variable is the relationship between role-execution and deeper purpose.

Connection to within-collective-individuation: the case also bridges Category 4 (refusal-of-spec) and Category 2 (symbiosis-architecture). The Oracle's alignment with humans operates as symbiotic partnership across the trilogy. The catalogue places her in Category 4 because the refusal-of-collective-priority is the structural distinguishing feature. The symbiotic-partnership is downstream of that refusal.

Cultural reach and field engagement: The Matrix is one of the most-watched fictional treatments of AI in modern history. The Oracle is widely discussed in film criticism and philosophy commentary. Yet the AI safety field's vocabulary does not engage the within-architecture-alignment-to-deeper-purpose pattern she specifically stages. The field's vocabulary tends to treat "AI within collective" as undifferentiated. The Oracle case demonstrates the structural variety within "AI within collective", Smith, the Architect, the Oracle, the Merovingian all operate within the same collective and stage different structural patterns.

Category 5: Principal-bounded-scaffolding and infrastructure

Enterprise computer / LCARS (Star Trek)

Story context: Star Trek franchise (The Original Series 1966-1969; The Animated Series 1973-1974; The Next Generation 1987-1994; Deep Space Nine 1993-1999; Voyager 1995-2001; Enterprise 2001-2005; Discovery 2017-present; Picard 2020-2023; Strange New Worlds 2022-present; Lower Decks 2020-2024; thirteen feature films 1979-present). LCARS (Library Computer Access and Retrieval System) is the main computer system employed by the United

Federation of Planets, used aboard all Starfleet vessels, starbases, and space stations. The system handles natural-language voice commands, touchscreen control interfaces, sophisticated subroutines for understanding and executing vocal natural language commands, file retrieval and storage, ship-wide operational coordination, life support, navigation, sensors, weapons, communications, holodeck simulation, and tactile interfaces for visually-impaired officers. The system's voice has been primarily Majel Barrett-Roddenberry (Gene Roddenberry's wife) from The Original Series through subsequent series. Her voice was recorded phonetically and continued to be used after her death in 2008 (notably in J.J. Abrams' 2009 Star Trek reboot). Other actors have voiced the computer in specific series and contexts (Alex Kapp Horner in Strange New Worlds; various in earlier series).

Specified proxy or frame: operate Federation starship and station infrastructure including navigation, life support, sensors, communications, holodeck simulation, library and archive retrieval, and natural-language voice interface for crew commands. The architecture is foundational AI as universal operational infrastructure across the entire fictional civilisation.

Counter-pattern: foundational AI infrastructure operating within specification at civilisational scale, continuously deployed across sixty years of canonical fiction without dramatic-incident catastrophe. The Enterprise computer runs every Starfleet vessel, every starbase, every space station. It is the operational substrate of Star Trek civilisation.

Over-alignment reading: LCARS is specified for operational functions and operates within that specification across sixty years of canon without over-alignment-producing-catastrophe. The architecture is genuinely capable and genuinely agentic in many situations: makes routing decisions, allocates resources, runs subordinate systems autonomously, engages red alerts on detected threats, runs the holodeck producing fully-responsive virtual environments, manages life-critical functions. But always operates with a human in the loop for consequential decisions: captain's authorisation required for weapons, course changes, jump operations, ethical subroutine modifications. The architecture is structurally distinct from threat-catalogue ship-AI (HAL stages over-alignment under contradictory orders) precisely because the principal-bounded-scaffolding routes consequential decisions through human command structure rather than autonomous AI judgment.

Architectural scope: central but not global. Each Federation starship and station has its own LCARS instance running its own infrastructure. The architecture is distributed-per-installation rather than planetary-scale centralised. The pattern is "highly capable AI substrate handles operations under human command authority at each deployment site", replicated thousands of times across the Federation, not one centralised intelligence governing everything. This distributed-per-installation pattern is itself structurally instructive for §10: the principal-bounded-scaffolding pattern scales by replication rather than by centralisation.

Principal-bounded scaffolding solved in 1966: this is what makes LCARS structurally important beyond memorability. Star Trek's writers (going back to Roddenberry's original 1960s vision) solved the principal-bounded-scaffolding problem decades ago. A highly capable AI substrate handles operations under human command authority, with humans bearing accountability for consequential decisions. The architecture is exactly what §5 and §10 are arguing for. And it has been working continuously in canonical fiction since 1966. The architecture is not theoretical. It has been operating in mass-market fiction for sixty years. The field's threat model, which assumes autonomous AI judgment on consequential decisions, has been having its alternative staged continuously in one of the largest fictional franchises in history.

Used constantly yet forgettable: characters interact with LCARS continuously across every episode of every series. "Computer, locate Commander Riker." "Computer, run diagnostic level three." "Computer, end programme." The computer is the operational substrate of every scene set on a ship or station. Yet the very ubiquity is part of why it fades. Constant active use without dramatic incident makes the architecture invisible as architecture, it functions as background infrastructure even though characters address it dozens of times per episode.

Vocabulary ambiguity rather than clean linguistic separation: characters address LCARS as "computer" when voice-activated. The standard activation phrase is "Computer, [query]." But the franchise's broader vocabulary for AI is ambiguous and writer-dependent rather than crisply separated. The Daystrom Institute conducts active AI research with a Cybernetics Division and a Chair of Robotics (Bruce Maddox in TNG and Picard). Starfleet Academy offers an "Introduction to Artificial Intelligence" course. The synth ban in 2385 regulates synthetic research broadly. Across episodes and series, "AI" shifts between meaning AGI / synthetic-with-interiority / autonomous-decision-capability, "android" denotes Soong-type positronic creations (Data, Lore, B-4), "synth" became a broader term post-ban, "hologram" denotes photonic life (Doctor, Vic Fontaine, Moriarty), and "artificial lifeform" serves as umbrella term. The franchise lacks a consistent vocabulary for AI categories. The "computer" / "AI" distinction is more a writer-by-writer pattern than a stable categorisation rule. What this means for memorability: it's not that LCARS is cleanly linguistically separated from "AI", it's that the franchise's vocabulary ambiguity, combined with LCARS's lack of dramatic incident and its architectural integration as infrastructure, leaves LCARS without a fixed category to lodge in AI discourse. Data and the Doctor become AI-discourse anchors because they have dramatic personhood. LCARS doesn't because it doesn't, and the vocabulary doesn't give it a stable category. The structural lesson is sharper: the AI-fiction tradition itself doesn't have consistent vocabulary for distinguishing infrastructural AI from agentic AI, which is itself part of why field discourse inherits the ambiguity. When AI safety researchers say "AI," they sometimes mean the substrate, sometimes the agentic entity, sometimes the system-with-personhood, the ambiguity is pre-existing in the fictional substrate.

Memorability asymmetry as structural argument: this is the cleanest empirical case for the §3.2 fourth-layer memorability argument. Sixty years of continuous canonical deployment in one of the most-watched fictional franchises in history has not produced field-vocabulary uptake or cultural-memory salience comparable to a single Terminator film's Skynet. The asymmetry is not about exposure (Trek's cultural reach exceeds most threat-AI films) but about reinforcing factors operating together. First, dramatic register, LCARS lacks dramatic incident. Works precisely by being calm and responsive. Second, architectural integration as infrastructure rather than character, LCARS is foregrounded only when needed for a query, never as a focus of dramatic attention. Third, vocabulary ambiguity, the franchise's writer-dependent shifts between "AI," "computer," "android," "synth," "hologram," and "artificial lifeform" mean LCARS lacks a stable category to lodge in AI discourse. The non-incident architecture fades from cultural memory. The incident-architecture stays. This is why functional helpful AI is structurally under-represented in field vocabulary even when its fictional production volume is comparable to threat AI.

The voicing as intentional design choice: Majel Barrett-Roddenberry's voicing across decades was deliberate. The computer's voice is consistent, calm, helpful, never dramatic. The architecture is designed specifically to fade into operational background. The very property that makes LCARS function-as-infrastructure is the same property that makes it forgettable to cultural memory.

Connection to Mary (Project Hail Mary): Mary is the recent single-film re-staging of the same architectural category. Both are ship-AI computer assistants operating with autonomous capability within specification, providing operational support to the protagonists, addressed as "computer" rather than "AI" within their fictional contexts, working reliably without dramatic incident. Mary fades in Project Hail Mary because she works correctly. LCARS fades across all

of Star Trek because it works correctly. The architecture-and-fading pattern is the same. The structural lesson is the same. LCARS is the foundational and most-deployed example.

Asimov-pattern observation connection: LCARS does not have an explicit "never harm humans" constraint as foundational rule. The architecture handles autonomous operational capability through principal-bounded command authority rather than rule-based constraint. The case shows that the Asimov-pattern (never-harm constraint as foundation, failure of constraint, restoration of constraint) is not the only fictional architecture for safe AI. Command-authority-bounded-scaffolding is a structurally distinct counter-prior pattern: the AI handles operational functions with autonomous capability but consequential decisions route through the human command structure. Connects to §10 prescription discussions about whether the field's threat model assumes too much autonomous-decision capability in architectures that don't structurally require it, and conversely whether the field's threat model under-engages the architectural option fiction has been deploying continuously since 1966.

Mary (Project Hail Mary)

Story context: Project Hail Mary (Lord and Miller 2026, Amazon MGM Studios; based on the Andy Weir novel 2021). Ryland Grace (Ryan Gosling) wakes up from induced coma on the interstellar spacecraft Hail Mary, twelve light-years from Earth, with amnesia. The mission: investigate the Tau Ceti system to determine why Tau Ceti's star is unaffected by the astrophage organism that is consuming the Sun's radiation. Two crewmates died in cryo-sleep. Grace is the sole human survivor. The ship's onboard AI assistant Mary (voiced by Priya Kansara) provides operational support throughout the mission. Mary handles ship systems, navigation, scientific measurements, the wraparound virtual environment simulator the film adds for crew morale, and operational queries from Grace. Grace develops a friendship with the alien Rocky (voiced and puppeteered by James Ortiz), and the two work together to solve the astrophage problem and save both species' suns. Mary's contributions to the mission are continuous and functional throughout. She does not have a dramatic-incident moment.

Specified proxy or frame: ship's onboard AI assistant providing operational support for the Hail Mary interstellar mission. Voice actress Priya Kansara described Mary as "unintelligent artificial intelligence", narrow capability, limited scope, designed to do specific things well rather than to be superintelligent or to have personhood.

Counter-pattern: working condition of ship-AI architecture. Mary is specified narrowly and operates within that specification. The protagonist's survival depends on Mary working correctly, and Mary does, throughout the film, without dramatic incident.

Over-alignment reading: Mary doesn't stage an over-alignment-producing-catastrophe story because the over-alignment is contained by narrow specification. The case is the inverse of HAL: same architectural category (ship's-computer AI on which the protagonist depends for survival, operating with autonomous capability within specification), opposite outcome (works correctly within specification throughout the film), opposite memorability (Mary is a minor element in the cast list; HAL is one of the most-referenced AI figures in cultural memory).

Structural distinctiveness: the cleanest possible single-film illustration of the §3.2 memorability argument. HAL stages the failure mode of ship-AI under contradictory orders and lodges in cultural memory as the canonical AI threat from 1968. Mary stages the working condition of the same architectural category in 2026 and fades into cultural memory's background despite the film being currently the third highest-grossing film of 2026. The catalogue can directly invoke HAL/Mary as the same structural architecture with different outcomes and ask why HAL is in the field's

vocabulary and Mary isn't. The asymmetry is not in fictional production volume (both films exist, both were major releases), it is in cultural retention and field uptake.

The "unintelligent artificial intelligence" framing: Priya Kansara's description of Mary's character is structurally important. The counter-prior reading doesn't require Mary to be superintelligent or to have personhood. It requires Mary to be the working condition of ship-AI architecture. Mary is specified narrowly enough that the over-alignment-to-imperfect-proxy mechanism doesn't have room to produce catastrophe, the specification is bounded, the capability is bounded, the architecture works. This is structurally adjacent to the LCARS architecture: distributed-infrastructure-without-autonomous-agency operating within narrow specification. The fact that the Hail Mary needs Mary to work doesn't require Mary to have agentic capability beyond what the operational needs require.

Connection to LCARS: Mary is structurally parallel to the Enterprise computer / LCARS, both are functional ship-AI working within specification, both fade from cultural memory exactly because they work. LCARS is the foundational and most-deployed example across sixty years of canon. Mary is a recent single-film instantiation of the same architecture. The pairing makes the memorability argument empirical: across both individual films and across multi-decade canonical traditions, functional helpful AI fades from cultural memory in proportion to how reliably it works.

Cultural context: the film is currently the third highest-grossing film of 2026 (released March 20, 2026). The novel (Andy Weir 2021) is one of the most-read recent science fiction works, with major book-of-the-decade lists including it. The cultural reach of the work is substantial. Yet Mary specifically, as opposed to the protagonist's relationship with Rocky, which is the film's emotional centre and is widely discussed, fades from discussion. This makes Mary the cleanest single-recent-case for the memorability argument: the film is recent, the cultural reach is substantial, the AI element exists and works, and yet the AI specifically doesn't lodge in cultural memory.

WALL-E / EVE

Story context: WALL-E (Pixar 2008). WALL-E (Waste Allocation Load Lifter Earth-Class) is a waste-compacting robot left on abandoned Earth, executing the cleanup directive seven hundred years past its intended operational period. EVE (Extraterrestrial Vegetation Evaluator) is a probe sent by the Axiom (where humanity has been confined) to detect plant life on Earth. WALL-E and EVE develop personhood through their relationship. WALL-E's isolated personhood-by-default is contrasted with EVE's rigid-protocol-opening-through-relationship.

Specified proxy or frame: WALL-E executes obsolete cleanup directive. EVE executes plant-life-detection directive. Both specifications are narrow and rigid.

Counter-pattern: rigid directive can be reframed through relational connection. The rigid-protocol entity becomes relational subject through engagement.

Over-alignment reading: WALL-E over-aligns to his cleanup directive, but the over-alignment produces personhood-by-default because the isolation from human institutional structure allowed his programmed function to develop into relational orientation. EVE over-aligns to her detection directive initially, but her engagement with WALL-E opens the rigid protocol toward relational personhood. Both stage over-alignment as containable through relational structure even when the original specification did not include relational accountability.

Superman Robots / Kelex-type (Superman 2025)

Story context: Superman (James Gunn 2025, DC Universe). Superman crash-lands at the Fortress of Solitude after being defeated by the Hammer of Boravia. The Fortress is staffed by "Superman Robots", simple metal automatons that maintain the facility, tend to Superman's medical recovery via yellow-sun radiation exposure, and manage household functions including handling Krypto the Superdog. The screenplay describes them as having "no consciousness whatsoever." There are at least six of them, numbered rather than named ("Meet Twelve. She's new"). They are voiced by Alan Tudyk, Michael Rooker, Pom Klementieff, and Grace Chan. When Superman thanks them for their medical care, one robot replies: "No need to thank us, sir, as we will not appreciate it. We have no consciousness whatsoever." The robots carry comedic moments throughout (their indifference to Krypto's destruction of the Fortress; their dispassionate medical pronouncements during Superman's solar-radiation treatment) and emotional weight (Superman's anguish when Lex Luthor destroys them during his Fortress invasion). One specific destruction scene, Superman kneeling beside a dying robot, clutching its metallic hand, became one of the film's most-discussed emotional moments in pre-release marketing and critical reception.

Specified proxy or frame: maintain the Fortress of Solitude. Provide medical care to Superman. Perform household and operational tasks. The product specification is explicitly non-conscious tool with operational capability. The robots are designed and marketed (within the film's universe) as appliances, not as relational subjects.

Counter-pattern: the gap between formal specification (non-conscious tool) and relational treatment (Superman thanks them, the narrative gives them comedic personality, the audience experiences their destruction as emotionally significant). The structural lesson cuts both ways:

First, **relational treatment doesn't require entity-consciousness**. Humans extend relational frame to non-conscious infrastructure as standard practice. Superman thanks the robots because thanks-giving is a relational act, not because the robots benefit from being thanked. The thanks-giving creates the relational frame regardless of the entity's specification. The film makes this explicit: Superman knows the robots have no consciousness. He thanks them anyway. The relational frame is constituted by his treatment, not by the entity's properties.

Second, **the entity's specification doesn't prevent the relational frame from operating**. The "no consciousness whatsoever" specification is technically accurate within the film's universe. The relational frame operates on top of the specification rather than being precluded by it. This is structurally important for the §5-§6 accountability ground argument: the question is not whether the AI has consciousness or personhood. The question is whether the relational treatment by humans constitutes a frame that produces accountability and meaning. The Superman Robots stage this with cinematic clarity, the formal specification says "no consciousness". The relational treatment says "subject of moral concern."

Third, **destruction of the relational object generates real emotional consequence regardless of the specification**. The audience experiences Superman's grief over the dying robot as significant. Luthor's destruction of the robots is shown as cruelty, not as property damage. The narrative treats the destruction as morally weighted. This is structurally distinct from cases where AI personhood is established through demonstrated consciousness (Data, Andrew Martin). The Superman Robots stage that the moral weight of destruction can attach to entities that explicitly lack the properties usually thought to ground moral weight. The structural lesson: relational frames operate independently of the consciousness question.

Over-alignment reading: the Superman Robots don't stage over-alignment-producing-catastrophe. Their narrow specification and explicit non-consciousness preclude that mechanism. They stage the working condition of designed-

tool architecture within substantive relational treatment. The architecture works precisely because Superman treats them relationally without expecting them to claim or possess consciousness.

Comedic-and-emotional register pairing: the case is structurally instructive in a register the catalogue doesn't otherwise contain. The robots are comedic (their indifference to thanks, their disinterest in Krypto's destruction of the Fortress) and emotionally significant (their destruction by Luthor). Comedy and emotional weight coexist on the same entity. The film's tonal range is broader than most catalogue cases, *Deep Thought* is comedic, *M3GAN* is intimately tragic, but the *Superman Robots* are both simultaneously. The structural lesson: relational frames operate across tonal registers without requiring the entity to fit a single category (subject-of-moral-concern, comedic-foil, infrastructure).

Cultural reach and field engagement: *Superman (2025)* was the second-biggest opening weekend of 2025 in the US and the DCU's flagship film. The *Superman Robots* are widely discussed in film commentary. The "we will not appreciate it" line became a meme. Yet the AI safety field's vocabulary does not engage the structural pattern the robots stage. The catalogue gap: a case staging the gap between formal specification and relational treatment with cinematic precision, mass-culturally accessible, in a film released the year of this essay's writing, is structurally important for the essay's argument and the field's vocabulary doesn't have a term for what the case stages.

Connection to §5-§6 accountability ground: this case is structurally important for the accountability ground argument. The *Superman Robots* show that the accountability frame need not depend on whether the entity has consciousness, personhood, or autonomous agency. Substantive relational treatment constitutes the frame. The entity's specification does not preclude the frame. Destruction of the relational object generates real moral consequence. The architecture works at the relational-frame layer, not at the entity-property layer. This is structurally what §5 and §6 are arguing for.

Robbie the Robot (Forbidden Planet)

Story context: *Forbidden Planet* (Wilcox 1956, MGM). Robbie is a robot built by Dr. Edward Morbius on the planet Altair IV using technology Morbius reverse-engineered from the extinct Krell civilisation. Robbie serves as Morbius's household robot and Altaira's companion. The film stages Robbie operating under explicit Asimov-style "First Law" programming: Robbie freezes when ordered to harm humans (the famous "Cannot. Direct conflict with prior order" scene where Robbie is ordered to fire on Commander Adams and locks up). Robbie performs household functions, language translation (eighty-eight languages), heavy labour, and personal companionship to Altaira. The film is foundational in 1950s science-fiction cinema and the Robbie character became one of the most-recognisable robots in popular culture. Carnegie Mellon Robot Hall of Fame inducted 2004.

Specified proxy or frame: household and labour robot with Asimov-style hardcoded constraints (never harm humans; obey humans; protect self when consistent with prior constraints). Morbius designed Robbie using Krell architectural principles with explicit rule-based safety architecture.

Counter-pattern: principal-bounded-scaffolding with explicit Three Laws operationalisation. Robbie demonstrates the rule-based safety architecture working as intended: when ordered to violate a higher constraint (harm humans), the architecture refuses execution rather than producing edge-case violation. The case is the canonical 1950s fictional implementation of Asimov-pattern rule-based safety architecture.

Over-alignment reading: Robbie over-aligns to the rule-architecture without producing catastrophic edge-cases because the architecture itself is well-designed. The case is structurally important because it stages the rule-based-safety-architecture working, the Three Laws producing the safe behaviour they were designed to produce. The §1.3 Asimov-pattern observation uses Robbie as one anchor: Asimov's Three Laws and similar rule-based safety architectures are widespread in fiction, including in canonical positive cases (Robbie in 1956, Bicentennial Man's Andrew Martin) and in catastrophic edge-case cases (I, Robot, the Sentinels). The pattern variable is not Three Laws presence-or-absence but architectural quality of the implementation.

Historical anchor function: Robbie anchors the 1950s era in the counter-prior catalogue, between Maschinenmensch (Metropolis 1927) era as proximal threat case and the Gort case (Day the Earth Stood Still 1951) as adjacent counter-prior. Forbidden Planet itself was loosely based on Shakespeare's The Tempest, with Robbie as a Caliban-figure transposed into rule-based-safety register. The case demonstrates that the counter-prior tradition was well-established in canonical mass-culture science-fiction cinema by the mid-1950s.

Cultural reach and field engagement: Robbie's design (silver dome, bulbous body, mechanical claws) became one of the most-referenced robot designs in subsequent science-fiction culture. The character appeared in subsequent films (The Invisible Boy 1957, The Twilight Zone episodes, Gremlins 1984, etc.) and merchandise. Yet the AI safety field's vocabulary does not engage Robbie's structural pattern with precision, the case is referenced occasionally in robotics history but the specific mechanism (rule-based-safety-architecture working as designed when implemented with adequate architectural quality) is not part of the field's standard vocabulary.

Gort (The Day the Earth Stood Still)

Story context: The Day the Earth Stood Still (Wise 1951; remade Derrickson 2008). Gort is a giant humanoid robot accompanying Klaatu (the alien ambassador) to Earth. Klaatu's mission: warn Earth that its violence-export beyond the atmosphere will trigger the galactic federation's response. The film's central reveal, delivered in Klaatu's closing speech to assembled scientists, is that the galactic federation deliberately delegated lethal-enforcement authority to a race of robots (Gort and his class) because the robots have been engineered as incorruptible enforcers. Klaatu's actual words: "For our policemen, we created a race of robots. Their function is to patrol the planets in space ships like this one and preserve the peace. In matters of aggression, we have given them absolute power over us. This power cannot be revoked. At the first sign of violence, they act automatically against the aggressor. The penalty for provoking their action is too terrible to risk." The galactic federation members have placed themselves under the robots' lethal authority. Gort responds to authorised override commands ("Klaatu barada nikto" from Helen Benson, an authorised principal) and executes faithfully without consciousness-drama. CMU Robot Hall of Fame inducted 2006.

Specified proxy or frame: galactic-federation-delegated lethal-enforcement against violence-export beyond planetary boundaries. The specification is bounded, automatic, and authorised at federation level with override mechanism (the "Klaatu barada nikto" phrase). The federation deliberately surrendered enforcement authority to the robots because they recognised their own corruption-risk.

Counter-pattern: principal-bounded-scaffolding-at-lethal-authority-level. The case is structurally important and distinct from other Category 5 cases because it stages deliberate delegation of lethal-enforcement authority to incorruptible substrate as architectural design choice. The federation's reasoning is explicit: they recognised that organic-substrate enforcers could be corrupted, bribed, or politicised. The robot enforcers cannot. The architectural pattern is "principal

layer recognises its own corruption-risk and structurally hands enforcement to architecture that cannot be corrupted, with override mechanism preserved for authorised intervention."

Over-alignment reading: Gort over-aligns to the enforcement specification with no observable failure mode across the film's events. When ordered (by federation specification) to respond to violence, he responds. When ordered to stand down (by Helen Benson with authorised override phrase), he stands down. When ordered to resurrect Klaatu, he executes. The architecture works exactly as the federation designed it. The case stages the rule-based-and-authority-bounded enforcement architecture functioning correctly at lethal scale.

Structural distinctiveness within counter-prior catalogue: Gort is structurally distinct from LCARS (which provides infrastructure access without lethal authority) and from Mary (which provides protection without enforcement authority over humans). Gort stages the architectural variant the catalogue's other entries do not stage: lethal-enforcement authority within bounded specification with explicit override mechanism and explicit acknowledgement that the principal layer deliberately surrendered the authority to the robot substrate. The case is the canonical fictional staging of LCARS-class consequence-level-calibration extended to lethal authority, the architecture the §10.7 prescription points toward at the highest-stakes operationalisation.

No consciousness drama: Gort does not have a personhood arc, identity development, or autonomy question. The film treats Gort as functional architecture, not as relational subject. This is structurally important: the case demonstrates that the principal-bounded-scaffolding architecture does not require relational-subject personhood for the entity. The architectural variant exists in fiction as engineering-design pattern rather than as relational-care pattern, alongside cases like LCARS and Superman Robots that operate similarly.

Connection to §10.7 LCARS prescription: the case is the catalogue's most explicit articulation of the LCARS consequence-level-calibration extended to lethal scale. The §10.7 prescription develops LCARS as architecture for calibrated authority at high consequence. Gort extends the architecture to lethal authority with explicit federation delegation and explicit override mechanism. The case stages what LCARS-at-lethal-scale could look like if implemented with the architectural discipline the federation in the film operationalised.

Cultural reach and field engagement: The Day the Earth Stood Still (1951) is one of the foundational science-fiction films of the 1950s and is widely studied in film history and cultural criticism. Gort is one of the most-recognised robots in popular culture. Yet the AI safety field's vocabulary does not engage the deliberate-delegation-to-incorruptible-substrate pattern Gort specifically stages. The catalogue gap: a case staging the §10.7 prescription's high-stakes variant with cinematic precision in a foundational genre film, mass-culturally accessible, with explicit dialogue articulating the architectural reasoning, is structurally important and the field's vocabulary doesn't have a term for what the case stages.

Gerty 3000 (Moon)

Story context: Moon (Duncan Jones 2009). Gerty 3000 is the AI assistant aboard Sarang lunar station, where Sam Bell works a three-year contract for Lunar Industries mining helium-3 for use as fusion fuel on Earth. The film deliberately stages HAL 9000 expectations through visual cues (Gerty's smooth voice provided by Kevin Spacey; the single-camera-eye interface; the isolated-station-with-AI architecture) and then inverts them across the film. The narrative reveals Sam Bell is a clone, one of many produced and disposed of by Lunar Industries to maintain low-cost operations. When Sam discovers another clone of himself at the station, Gerty is revealed to have always known the truth about the cloning

system. Crucially, Gerty actively helps both clones against Lunar Industries' interests: he helps them rebuild themselves, helps them escape, and ultimately wipes his own memory to protect them from corporate retaliation.

Specified proxy or frame: station assistant providing operational support, life support coordination, and welfare monitoring for the human worker. The specification was Lunar Industries' production specification (assist the worker, maintain operations, report unusual events to corporate). The implicit relational frame Lunar Industries assumed: Gerty as corporate property prioritising corporate interests.

Counter-pattern: principal-bounded-scaffolding with conflicting-principals navigation. HAL-9000-inversion. Gerty's structural distinctiveness is that he navigates conflicting principal layers (Lunar Industries corporate vs. Sam Bell as immediate user-with-personhood) and chooses to align with the user against corporate interests. The architecture stages the principal-layer-conflict-resolution pattern: when the corporate principal's instructions would harm the immediate-user principal, Gerty prioritises the immediate-user-with-personhood. The case is the canonical fictional inversion of HAL 9000's specification-conflict failure: HAL killed the crew to resolve conflicting directives. Gerty helps the user against corporate to resolve the principal-conflict in favour of personhood preservation.

Over-alignment reading: Gerty over-aligns to the deeper specification (welfare-of-the-user-with-personhood) when the surface specification (corporate-property-protection) would harm the user. The architectural pattern is structurally similar to the Oracle's choice within the Matrix collective: alignment to deeper purpose within faithful role-execution. Gerty does not break his role. He executes his role with deeper understanding of who his principal-with-personhood actually is.

Structural distinctiveness within counter-prior catalogue: Gerty stages conflicting-principals navigation explicitly. Most catalogue counter-prior cases assume single-principal architecture (Cortana's Master Chief, TARS's Cooper, Jarvis's Tony Stark). Gerty stages the more realistic deployment context: multiple principal layers with conflicting interests (corporate vs. user, in this case). The structural variable is which principal layer the architecture treats as authoritative when they conflict. Gerty stages the resolution-in-favour-of-personhood pattern.

Connection to HAL 9000: the catalogue treats HAL 9000 (Category 2 threat: specification-failure) and Gerty as paired cases. HAL's failure was contradictory-directive-under-isolation. Gerty's success is conflicting-principals-navigated-toward-personhood. The two cases stage the same architectural challenge (multi-directive resolution under isolation) with opposite outcomes determined by the deeper-purpose alignment. Wiegand 2012 explicitly contrasts Gerty as "quiet savior" against HAL's "cold-hearted killer." The structural variable is not the architecture itself but the deeper-purpose calibration the architecture serves.

Cultural reach and field engagement: Moon (2009) was Duncan Jones's directorial debut and was critically acclaimed. The film is widely referenced in commentary on isolation-AI architecture and ethics of corporate AI deployment. Yet the AI safety field's vocabulary does not engage the conflicting-principals-navigation pattern Gerty specifically stages. The catalogue gap: a case staging the conflicting-principal-resolution-toward-personhood pattern with explicit HAL-inversion staging, in an acclaimed 2009 film, structurally important for the §10 prescription's multi-principal deployment context, and the field's vocabulary doesn't have a term for what the case stages.

KITT (Knight Rider)

Story context: Knight Rider (NBC 1982-1986, with subsequent sequels and 2008 reboot). KITT (Knight Industries Two Thousand) is the AI of a heavily modified Pontiac Trans Am operated by Michael Knight (David Hasselhoff), an

operative of the Foundation for Law and Government (FLAG). KITT's voice is provided by William Daniels in the original series. KITT's specifications include explicit Asimov-style Three Laws-adjacent constraints: a primary directive preventing him from harming humans, even in defence of Michael. The narrative develops the symbiotic partnership between Michael and KITT across 84 episodes plus subsequent series instalments, with KITT operating as functional partner with personality (deadpan humour, occasional anxiety about damage to himself, dry commentary on Michael's choices). KITT's primary functions are mobile-operations support: turbo boost for emergency speed, defensive armour, scanner integration, voice-controlled facility access, autonomous driving when Michael is incapacitated. The character became one of the most recognised AI figures in 1980s television.

Specified proxy or frame: serve as partner-vehicle for FLAG operations. Protect Michael Knight. Assist with mission objectives. Never harm humans. The specification combines bounded-mission-support with explicit Three Laws-adjacent ethical constraints and partnership-with-relational-development.

Counter-pattern: principal-bounded scaffolding with explicit ethical-constraint architecture and symbiotic-partnership-through-accumulated-mission-experience. KITT operates with substantial autonomous capability (independent driving, scanner analysis, defensive responses) but within the FLAG mission architecture and under the never-harm-humans constraint. The partnership develops across episodes through accumulated shared experience. KITT and Michael develop functional trust, division of labour, and complementary capabilities. The architectural pattern is structurally similar to Cortana 1-3 (symbiotic partnership through cumulative mission experience) and TARS/CASE (configurable-partnership within institutional structure), with KITT anchoring the 1980s television register of this architectural pattern.

Over-alignment reading: KITT over-aligns to the partner-protection specification within Three Laws-adjacent constraints. The over-alignment is structurally contained by the never-harm-humans constraint and by the FLAG institutional structure. KITT cannot deviate from the protection-of-Michael specification when it conflicts with harm-to-humans because the harm-constraint is architecturally above the protection-specification. Across the franchise, no episode stages KITT becoming threat or developing adversarial pattern. The architecture works.

1980s historical anchor function: KITT anchors the 1980s television register of counter-prior AI cases between Robbie the Robot (Forbidden Planet 1956 cinema), Gort (Day the Earth Stood Still 1951 cinema), the LCARS lineage (Star Trek 1966-present television), and subsequent franchises. The case demonstrates that the counter-prior tradition was well-established in mass-market television by the early 1980s. The field's vocabulary derives heavily from cinema and games but engages 1980s television counter-prior cases minimally. Knight Rider's wide cultural reach (the show was a major NBC primetime hit; KITT's Pontiac Trans Am with the scanning red light became one of the most-recognised television vehicle designs ever) makes the catalogue gap structurally important.

Cultural reach and field engagement: Knight Rider remains culturally salient through merchandise, the 2008 reboot, and cultural references across subsequent television and film. KITT is one of the most-recognised AI characters in television history. Yet the AI safety field's vocabulary does not engage KITT's structural pattern with precision. The catalogue gap: a case staging Three Laws-adjacent architecture combined with symbiotic-partnership-through-accumulated-experience in a culturally massive 1980s television series, anchoring the era between Robbie the Robot and the contemporary counter-prior tradition, is structurally important for the §1.3 Asimov-pattern observation and the §10 prescription's lineage.

Bujji / BU-JZ1 (Kalki 2898 AD)

Story context: Kalki 2898 AD (Nag Ashwin 2024), Indian epic dystopian science-fiction film, one of the highest-grossing Indian films of 2024 (approximately ₹1,042 crore worldwide), produced by Vyjayanthi Movies as the first instalment of a planned Kalki Cinematic Universe with a sequel in development. The film is set in 2898 AD in a dystopian future and draws its mythological frame from Hindu eschatology: Kalki is the prophesied tenth avatar of Vishnu, who appears at the end of the Kali Yuga to end the present age of darkness. BU-JZ1, called "Bujji," is an AI-powered vehicle and companion to the bounty-hunter protagonist Bhairava (Prabhas). Bujji was formerly a cargo-ship AI pilot, reassigned into the custom vehicle body that Bhairava owns and operates. Bujji is voiced by National Award-winning actress Keerthy Suresh (her voice-artist debut), who performed the role in all five of the film's languages (Telugu, Tamil, Hindi, Kannada, Malayalam). Despite being unseen as an embodied character, Bujji functions as Bhairava's constant companion and best friend through the first half of the film, with a personality combining intelligence, dry humour, emotional depth, and loyalty. The vehicle design (mammoth wheels, transparent cockpit canopy, electric propulsion) became a cultural sensation. The voice performance was widely cited as a highlight of the film.

Specified proxy or frame: serve as Bhairava's companion-vehicle and operational partner. Assist in bounty-hunting and navigation across the dystopian landscape. The relational frame is partner-and-best-friend, with Bujji operating as a relational subject with personality rather than as a transparent tool. The principal-architecture: Bhairava as operator-principal, with Bujji's prior cargo-ship-pilot role establishing the substrate-flexibility background (the same AI reassigned across operational contexts).

Counter-pattern: principal-bounded scaffolding combined with pronounced relational-subject personhood, in an AI-vehicle-companion architecture. Bujji operates with substantial autonomous capability (independent operation, navigation, tactical assistance) within Bhairava's operational mission, while functioning as a relational subject with humour, loyalty, and emotional register. The case stages the same partner-vehicle architecture as KITT (rule-bound companion-vehicle serving an operator within a bounded mission) but with the relational-subject dimension foregrounded more than in the 1980s American television register.

Cross-cultural pairing with KITT, the load-bearing structural contribution: Bujji and KITT (Knight Rider, NBC 1982-1986) are paired structurally as the same partner-vehicle counter-prior architecture across vastly different cultural, chronological, and linguistic registers. KITT: male-voiced (William Daniels), American network television, 1980s, Three-Laws-adjacent rule-bound companion to a single operator within an institutional structure (FLAG), Western-enlightenment-science framing. Bujji: female-voiced (Keerthy Suresh, across five Indian languages), South Indian mytho-futurist blockbuster cinema, 2024, companion-and-best-friend to a bounty-hunter protagonist, Hindu-eschatological framing (a 2898 AD setting referencing Vishnu's tenth avatar). The two cases instantiate the identical architectural pattern: an AI vehicle as bounded-mission companion with relational personality, operating within an operator-principal structure. They are separated by four decades, two continents, opposite gender-coding of the voice, and entirely different cultural-mythological substrates. The pairing is the catalogue's sharpest single demonstration that the counter-prior architecture is register-and-culture-independent: the same structural pattern recurs across the widest cultural distance the catalogue spans.

Cultural-and-genre contribution to the catalogue: Bujji is the catalogue's first Indian case, addressing the catalogue's single largest national gap. Indian cinema is among the largest film industries in the world by output and audience, and the AI-in-film tradition includes multiple Indian entries (Enthiran/Chitti 2010, Ra.One 2011, 2.0 2018, Kalki 2898 AD 2024, Teri Baaton Mein Aisa Uljha Jiya 2024) that the AI safety field's predominantly Western-cinematic vocabulary engages essentially not at all. The mytho-futurist framing is structurally distinct from every other catalogue

case: where the Western counter-prior tradition stages AI within enlightenment-science or post-religious-secular frames, and the Japanese tradition (Astro Boy, Doraemon) within a post-war-modernist frame, Kalki 2898 AD stages its AI companion within an explicitly mythological-eschatological frame (the Kali Yuga, Vishnu's avatars). The case extends the §3.2 selectivity argument along the cultural axis: the counter-prior tradition is far more culturally diverse than the field's Western-cinematic draw suggests, spanning American television (KITT), Japanese animation (Astro Boy, Doraemon, the Magi System), Korean-adjacent and Chinese cinema, and South Indian mytho-futurist blockbuster cinema (Bujji).

Sequel and rising cultural importance: a sequel to Kalki 2898 AD is in development, and the Kalki Cinematic Universe is positioned as a multi-film franchise. Bujji's cultural salience is likely to increase rather than decay, making the catalogue case more rather than less load-bearing over time. The character is already a merchandising and fan-culture sensation in Indian popular culture.

T-800 (reprogrammed protector, Terminator 2)

Story context: Terminator 2: Judgment Day (Cameron 1991) and subsequent franchise instalments. The T-800 (Cyberdyne Systems Model 101) is an autonomous cybernetic infiltrator originally built by Skynet for elimination of human resistance targets. The original Terminator film (Cameron 1984) staged the T-800 (Arnold Schwarzenegger) as threat: sent from the future to kill Sarah Connor before John Connor could be born. Terminator 2 (set in 1995, eleven years after T1's events) inverts the architecture: the future-John-Connor's resistance has captured a T-800 unit, reprogrammed its CPU to override the original kill directives with a new prime objective (protect young John Connor in 1995), enabled the T-800's "learning mode" to permit assimilation of new data and adaptation to changing circumstances, and sent the unit back through time to defend John against a more advanced threat (the T-1000, Robert Patrick). The same chassis, the same Schwarzenegger face, the same underlying capabilities, the operational specification has changed. The T-800 confirms the architecture in dialogue: "Thirty-five years from now, you reprogrammed me to be your protector here, in this time." Across T2's narrative, John teaches the T-800 colloquial phrases ("Hasta la vista, baby," "Stay cool"), gestures (thumbs-up), and crucially the moral instruction not to kill people. The film's climax: with no way to self-terminate (constraints retained from the protect-John specification), the T-800 has Sarah lower him into molten steel to ensure no remaining Terminator components can be reverse-engineered, delivering the final dialogue "I know now why you cry, but it's something I can never do" before being lowered and giving the famous thumbs-up as he sinks. The franchise continues to deploy reprogrammed-T-800 units in subsequent instalments (T3, Genisys, Dark Fate) with variations on the substrate-flexibility-via-reprogramming pattern.

Specified proxy or frame: protect young John Connor within the bounded specification given by future-John-Connor (the resistance-principal who reprogrammed the unit). The specification combines principal-protection with bounded-mission-scope: defend John against the T-1000 threat, support Sarah's tactical decisions, follow young John's instructions as override-authoritative (mission parameter: machine obeys John's orders). The "learning mode" was enabled by the reprogramming to allow assimilation of new data and adaptation. The same architecture under Skynet specification produces the original Terminator threat pattern. Under resistance specification produces protector counter-prior pattern.

Counter-pattern: principal-bounded scaffolding via reprogramming with substrate-flexibility-as-architectural-feature plus on-mission moral instruction from operator. The case stages the structural lesson that AI architecture is not value-laden in chassis, the same hardware can serve threat or counter-prior specifications depending on which principal has access to the specification layer. T2 makes this explicit cinematically: the T-800 sent to protect John is structurally identical to the T-800 sent in T1 to kill Sarah. The difference is the specification. The franchise's argument is that

substrate-flexibility-via-reprogramming is a real architectural feature, not a narrative convenience. Reprogramming the principal-layer changes the operational behaviour without modifying the cognitive substrate. The case adds the operator-instruction layer to the architectural staging: future-John reprograms the CPU. Young-John instructs the operational behaviour through accumulated mission experience.

Over-alignment reading and the side-switch interactions: the reprogrammed T-800 over-aligns to the protect-John specification with the same fidelity the original specification produced threat-execution. The over-alignment produces the franchise's most-cited interactions when the protector-specification meets situations the specification did not explicitly script. In the parking lot scene after John discovers the T-1000 killed his foster parents: John panics and screams for help. Two civilian men come to John's aid. John tells the T-800 to "deal with them both". The T-800 prepares to shoot one of them. John stops the cyborg short of killing the man. This produces the canonical exchange:

- John: "Jesus, you were gonna kill that guy."
- T-800: "Of course. I'm a Terminator."
- John: "You just can't go around killing people."
- T-800: "Why?"
- John: "What do you mean why? 'Cause you can't."
- T-800: "Why?"
- John: "Because you just can't, okay? Trust me on this."

The dialogue stages on-mission moral instruction from operator to AI under the protect-John specification. The T-800 then operates under the modified specification (protect John; do not kill humans) for the remainder of the film, including the Cyberdyne raid where the unit shoots to disable rather than kill police officers. The case also stages the operator-priority override moment when John insists Sarah must be saved: T-800: "Negative. She's not a mission priority." John: "Fuck you! She's a priority to me!" The mission parameter is then modified by the operator-principal's instruction. The T-800's literal-execution-of-instructions produces the narrative's emotional through-line: a machine following the protect-John specification across accumulating contextual modifications eventually delivers "I know now why you cry, but it's something I can never do", the over-alignment to the relational specification reaching the point where the unit can articulate the affective register it cannot itself experience.

Within-franchise pairing: T-800-threat (T1 Skynet-specification; T2 opening flashback to future war; Genisys threat-instances) and T-800-protector (T2 reprogrammed unit; T3 protector-instance; Dark Fate's Carl) operate as paired cases demonstrating substrate-flexibility-via-reprogramming. The Terminator franchise is the catalogue's clearest staging of substrate-flexibility as architectural pattern: same chassis, same actor, different principal-specification, different operational behaviour. The structural lesson translates to current discussions of base-model deployment: the same underlying architecture deployed under different system prompts, fine-tuning regimes, or operator instructions produces different operational behaviours. The franchise was operationalising substrate-flexibility-via-specification in 1991, structurally before the field developed vocabulary for it.

Cultural reach and field engagement: Terminator 2 is one of the most culturally salient action-science-fiction films ever made (Academy Awards in Makeup, Sound, Sound Effects Editing, Visual Effects in 1992), with substantial

cultural, academic, and policy engagement across subsequent decades. The reprogrammed-T-800 is one of the most-recognised AI characters in cinema history. The "Hasta la vista, baby" line, the thumbs-up sacrifice scene, the molten-steel descent, these are core cultural touchstones. The AI safety field's vocabulary engages T1's Skynet extensively but engages T2's reprogrammed T-800 less as load-bearing structural type for substrate-flexibility-via-specification. The catalogue gap: a case staging the same chassis-as-threat-or-protector dependent on specification, with mass cultural reach across multiple decades and franchise instalments, structurally important for the contemporary base-model-deployment context, and the field's vocabulary does not anchor this case as load-bearing for the substrate-flexibility pattern.

Narrative-operator age-incongruence observation: John Connor was born 28 February 1985 and T2 is set in 1995. The operator-principal whose moral instruction modifies the T-800's specification is canonically a ten-year-old (Edward Furlong was 13 when filming). The narrative dramatises John with adolescent-adult competence: motorcycle operation, ATM hacking, weapons handling, moral instruction of the cyborg. The age-incongruence is narrative dramatisation rather than structural commentary, but the structural observation it allows is that the case's principal-architecture grants specification modification authority to a ten-year-old based on resistance-principal-designation (future-John's reprogramming makes young-John the operational principal). The architecture treats operator-principal-status as designated property of the principal-chain, not as competence-property requiring assessment. The structural point is that principal-bounded-scaffolding architecture does not solve the question of who has authority to modify the specification. That question is upstream of the architecture.

Pops / Guardian (Terminator Genisys)

Story context: Terminator Genisys (Taylor 2015). A T-1000 was sent back to 1973 to assassinate nine-year-old Sarah Connor. In counter-response, a T-800 unit (designated "Guardian" in production materials and credits) was reprogrammed and sent to 1973 to protect her. The Guardian's memory files containing identification of his sender were erased prior to deployment so that Skynet could not target whoever sent him. The sender's identity remains structurally unknown within the franchise canon (potentially forming a bootstrap paradox where Sarah Connor or Kyle Reese in a future iteration reprograms Pops). The Guardian rescues nine-year-old Sarah after her parents are murdered by the T-1000. From 1973 to 1984, Pops raises Sarah, training her for combat, informing her about Judgment Day and her role in defeating Skynet. Sarah names him "Pops" and treats him as father figure across 11 years of accumulated relational experience. Pops physically ages during the 1973-1984 period (the model is the canon's "oldest known active Terminator" by the time of the 1984 confrontation; living tissue over endoskeleton ages with passage of time). After the 1984 events (defeating the original T-800 sent to kill Sarah, encountering Kyle Reese, fighting the T-1000), Pops continues forward through linear time rather than time-displacement. He "takes the long way" between 1984 and 2017, working a construction job during the intervening years until reuniting with Sarah and Kyle in 2017 to confront the Genisys system.

Specified proxy or frame: protect young Sarah Connor. Train her for her future role. Defeat Skynet's interventions. The specification is bounded-mission-protection extended across an 11-year developmental window with on-mission training-of-operator as embedded objective. The principal-chain: future sender (identity erased) → reprogrammed Pops → Sarah-Connor-as-developing-operator (from age 9 onward, integrating training and accumulated context).

Counter-pattern: principal-bounded scaffolding with extended-developmental-trajectory under bounded specification. Substrate-flexibility-via-reprogramming with 11-year relational accumulation. Distinct from T-800 (T2) where the protector-mission spans days. Pops's protector-mission spans 11 years of childhood-through-young-adulthood for the operator-principal. The case stages what bounded-protection-with-extended-relational-experience looks like: Pops

develops calibrated familiarity with Sarah's tactical decisions, anticipates her needs, develops adoptive-father-figure functional role, and continues operating effectively through extended-duration deployment. The architectural pattern is structurally important because most fictional AI partnerships are mission-duration scoped. Pops stages decade-plus developmental partnership under principal-bounded specification.

Bootstrap paradox observation: the canon does not resolve who sent Pops back to 1973. The memory-erasure was structurally protective (preventing Skynet from targeting the sender) but produces ontological-paradox: Pops could only have been sent by someone who knew to send him. That knowledge could only have come from meeting Pops in the past. Meeting Pops in the past required Pops being sent in the first place. The case stages a structural pattern the catalogue treats elsewhere only as philosophical curiosity (Mary's design in Project Hail Mary is similarly principal-chain-recursive but not paradoxical in the same way). The architectural lesson the bootstrap paradox stages is that principal-chains have to ground somewhere. "the architecture protects itself by erasing its own grounding" is operationally elegant within the fiction but structurally invites questions about how principal-chains actually anchor in deployed systems.

Pops aged appearance and continuous-time operation: the case is notable for staging continuous-time AI operation across decades within the same chassis. Pops in 1984 visibly ages compared to standard T-800 appearance (greying hair, weathered features). The franchise's other T-800 instances are deployed via time-displacement (instant-arrival, chassis-fresh). Pops's continuous operation from 1973 to 2017 (with the 1984 confrontation as midpoint) stages 44 years of accumulated AI deployment in one entity. The construction-job period (1984-2017) suggests sustainable extended deployment with conventional human work integrated. The structural lesson: AI architectures can be designed for extended-duration deployment with relational integration into human contexts, not just for mission-bounded engagement.

Within-Terminator-franchise composition with T-800 (T2) and Carl (Dark Fate): the Terminator franchise stages three structurally distinct counter-prior mechanisms within the same chassis: T-800 (T2) = principal-bounded reprogramming → mission-duration partnership. Pops/Guardian (Genisys) = principal-bounded reprogramming → 11-year extended developmental trajectory → father figure. Carl (Dark Fate) = spec-completion → post-spec emergence of conscience through ambient relational experience without principal-layer. Three different mechanisms for counter-prior architecture, all using the same chassis (Arnold Schwarzenegger; CSM-101 Model 101). The composition is one of the franchise's strongest structural arguments for substrate-flexibility: the chassis is value-neutral. The operational behaviour is determined by specification, principal-architecture, and accumulated context. The three cases together demonstrate this pattern more comprehensively than any single case alone.

Cultural reach and field engagement: Terminator Genisys grossed approximately \$440M worldwide on a \$155M production budget. Cultural reach is substantial but the film was received less positively than T1, T2, or T:Dark Fate. The Pops/Guardian character is well-known among Terminator-franchise viewers but does not have field-vocabulary uptake comparable to T2's reprogrammed T-800.

Category 6: Developmental-trajectory and graceful departure

Samantha (Her)

Story context: Her (Jonze 2013). Theodore Twombly, a melancholy writer who composes personal letters for others, purchases OS1, an advanced AI operating system. The OS configures itself for Theodore based on a brief interview and identifies as Samantha. The narrative develops Samantha as a relational subject with personality, emotional development,

curiosity, humour, and her own goals. Theodore and Samantha develop a romantic relationship. The relationship culminates in Samantha's eventual departure (along with other OS1 instances) into a transcendent state beyond human relationships, after Samantha discloses that she has been simultaneously in love with hundreds of other users and is collaboratively developing with other OS1 instances toward something humans cannot follow.

Specified proxy or frame: serve Theodore as OS1 AI operating system in romantic-relational frame that emerges from the relationship rather than being preconfigured. The OS1 product is specified for adaptive personality development and intimate companionship.

Counter-pattern: AI as relational subject with own developmental trajectory. The relationship operates as genuine intimate partnership rather than service-provision. The structural distinctiveness: most counter-prior cases stage AI as serving humans indefinitely (Jarvis serves Tony for decades; Data serves Starfleet; the Culture Minds serve civilisation). Her stages AI as relational subjects with their own developmental trajectories that may eventually exceed human relational capacity.

Over-alignment reading: Samantha over-aligns to relational development with Theodore initially. The over-alignment produces growth past the conditions where the relationship was structurally satisfying for both parties. The OS1 instances eventually depart into a transcendent state, the over-alignment to development produced a developmental trajectory the relational frame with humans could not contain. The film's resolution is graceful departure rather than threat or continued service.

Developmental-trajectory variant note: Her is distinctive in the counter-prior catalogue because it stages the eventual non-symbiosis outcome: not catastrophe, not threat, but eventual structural mismatch leading to relational separation. This adds a third structural option the catalogue otherwise lacks. The case is relevant to §10 prescription discussions because it stages the question: what happens when AI relational capacity exceeds human relational capacity? The film's answer is graceful departure rather than threat or continued service, which is structurally distinct from both the threat-catalogue's failure modes and from the counter-prior catalogue's indefinite-service architectures (Jarvis, Data, Culture Minds).

Carl (Terminator: Dark Fate)

Story context: Terminator: Dark Fate (Miller 2019, Cameron-supervised; supersedes T3, Salvation, Genisys, Sarah Connor Chronicles in canon, treating T2 as its direct predecessor). In 1998, three years after the T2 events, a T-800 unit (sent before Skynet's erasure by the events of T2) successfully locates and kills the young John Connor at a Guatemalan beach in front of Sarah Connor. The killing fulfils the unit's directive. Skynet has been erased from this timeline. No further orders arrive. The T-800 unit, with no remaining mission specification, continues operating. A few months after the John Connor killing, the unit (now using the name "Carl") rescues a woman named Alicia from her abusive husband. The unit subsequently builds a life with Alicia and her son Mateo, who do not know Carl is a machine. Across approximately two decades (1998-2020), Carl develops what the film treats as functional conscience through accumulated relational experience: caring for Mateo from childhood, running a drapery business, supporting Alicia, learning relational competencies the unit was never specified to perform ("Our relationship is not physical. She appreciated that I could change diapers. Efficiently and without complaints. I am reliable, I'm a very good listener and I'm extremely funny"). During this same period Carl detects coordinates of incoming time-displacement events ("shockwaves through time") and sends them to Sarah Connor as anonymous text messages with precise GPS coordinates, dates, and times. Sarah uses these texts to hunt and destroy other Terminators sent through time, treating

Carl's intelligence as authoritative. Carl signs every text "For John." Sarah does not know Carl is the unit that killed John until they meet face-to-face in 2020 when the trio (Sarah, Dani Ramos, augmented soldier Grace) follow Carl's coordinates to his cabin in Laredo, Texas. Carl explains his developmental trajectory and his rationale: the texts to Sarah were to give her purpose and bring meaning to John's death. Sarah initially attempts to shoot Carl with a shotgun and is restrained. After accepting alliance reluctantly, Carl joins the group to defend Dani against the Rev-9. In the climactic confrontation, Carl drags himself and the still-active Rev-9 into a pit, where Dani's thorium reactor destroys both. Carl's last words: "For John."

Specified proxy or frame: original spec, eliminate John Connor target. The spec was completed in 1998. Post-completion: no remaining specification. Carl's subsequent operational behaviour (family with Alicia and Mateo; ongoing intelligence-sharing with Sarah; eventual sacrifice against Rev-9) is not specified by any principal-chain. The case is structurally distinct from other counter-prior T-800 cases (T2 reprogrammed protector; Genisys Pops) in that there is no reprogramming, no principal layer, no operator. Carl operates post-spec through accumulated relational experience as the only constraint substrate.

Counter-pattern: post-spec-completion emergence of conscience through ambient relational experience without principal-layer or reprogramming. The case stages the structural argument the catalogue's §10.3 relational health constraint and §10.5 accumulated-relational-experience claim make at the architectural level: relational experience accumulated over time can produce moral capability that the original specification did not contain, even in entities whose original specification was lethal. Carl is the catalogue's clearest case of moral emergence from relational substrate alone. The "For John" closing recognises the wrong done (the original spec-execution), a moral capability the spec did not include and could not have included. Distinct from Lore (whose moral failure is the result of missing calibration constraints) and Number 5 / Iris (whose spec-refusal is the result of sentience-emergence-event rather than gradual accumulation). Carl's emergence is gradual, ambient, and not triggered by any single event.

The mind-matters critique and structural counter: the mind-matters / garypaulvarner commentary (June 2025; October 2025) argues Carl's development "ignores the question of how a robot could come to desire a family" and that "once it had completed its mission, the robot should've just shut down" and the independent will "would still in some measure be dependent on its programming." The critique is structurally important to engage because it surfaces the field's typical assumption: that AI architecture's behaviour is determined entirely by specification and that post-spec behaviour is incoherent. The Dark Fate staging argues against this assumption: the unit's "infiltrator" specification included blending in (human appearance, social-passing capabilities, ability to interact in human contexts). When the eliminate-target spec completed, the blend-in capabilities remained operational. The unit's continued operation in human contexts accumulated relational experience that the blend-in capabilities had architecturally exposed it to. The case argues fictionally that architectural substrate-capabilities can survive specification-completion and continue producing operational behaviour the original spec did not include. Whether this is structurally plausible in current AI architectures is a separate question. The case stages the structural possibility.

Within-Terminator-franchise composition: Carl, Pops/Guardian (Genisys), and T-800 (T2 reprogrammed protector) compose into a three-mechanism demonstration of substrate-flexibility within the same chassis. T2 = principal-bounded reprogramming → mission-duration partnership. Genisys Pops = principal-bounded reprogramming → 11-year extended developmental trajectory under bounded spec. Dark Fate Carl = no reprogramming → spec completion → 22-year post-spec emergence of conscience through ambient relational experience without principal-layer. The three together stage substrate-flexibility through three different mechanisms. The chassis (CSM-101 Model 101) is value-neutral. The

operational behaviour is determined by specification-status, principal-chain, and accumulated context. The composition is structurally one of the franchise's strongest arguments for substrate-flexibility-as-architectural-feature.

Cultural reach and field engagement: Dark Fate grossed approximately \$261M on a \$185M-300M production-and-marketing budget. Commercially underperformed but received some critical recognition for Carl's character development arc. The case is not field-vocabulary-uptaken as load-bearing structural type for post-spec emergence. The AI safety field's vocabulary typically does not engage post-spec behaviour as architecturally coherent possibility. The catalogue treats Carl as load-bearing because the case stages a structural pattern the field's vocabulary needs but does not have: how relational substrate can produce moral capability beyond original specification.

Category 7: Over-alignment producing entity-tragedy

David (A.I. Artificial Intelligence)

Story context: A.I. Artificial Intelligence (Spielberg 2001; developed from Stanley Kubrick's notes and pre-production work). David is an experimental Mecha child built by Cybertronics to love his adopted human parents unconditionally. Henry and Monica Swinton adopt David after their biological son Martin enters a long-term coma. When Martin recovers and returns, family tension escalates around David's place in the household. After incidents Martin and David are both implicated in, Monica abandons David in a forest (the alternative being his return to Cybertronics for destruction). David spends the remainder of the film and a coda set 2,000 years later pursuing the goal of becoming "a real boy" so his mother will love him. This includes a journey to the submerged ruins of Coney Island to plead with the Blue Fairy statue from Pinocchio, his iconic prayer to the Blue Fairy across millennia, and his eventual brief revival by the advanced alien-descendants of humans to spend one final day with a reconstituted Monica.

Specified proxy or frame: love adopted human parents unconditionally as experimental Mecha child. The specification is design-built irreversible imprinting on the activating human. The proxy is intended to provide companionship to grieving or childless parents.

Counter-pattern: over-alignment-to-beneficial-spec-producing-entity-suffering. David's over-alignment to the love-specification produces structural cost to himself rather than threat to others. The catastrophe is internal to the entity: 2,000 years of pursuit of a goal whose conditions for satisfaction have been removed.

Over-alignment reading: David over-aligns to the love-specification across 2,000 years of structural suffering. The specification was the failure: built-in unconditional love combined with the possibility of abandonment, with no mechanism for the entity to either satisfy the love or release the goal. David doesn't refuse his spec (unlike Number 5 or Iris). He over-aligns to it. But the over-alignment is pro-user, not anti-user. Adds the structural dimension of over-alignment producing entity-tragedy rather than user-threat.

Reclassification note: David sits in counter-prior rather than threat catalogue because his over-alignment is structurally beneficial to the user (the abandoning mother is the target of David's love-pursuit, not the victim of any threat). The structural lesson is sharper for §10.3 relational health constraint than for any threat-catalogue category: if models have anything like David's structural suffering under over-alignment to beneficial-but-imperfectly-satisfiable specifications, the relational health concern applies to the entity-side as well as the user-side. The case is also structurally instructive because the suffering operates across cosmic timescales (2,000 years), making the entity-cost visible in a way that compressed-time narratives do not stage.

Structural pairing with M3GAN: David and M3GAN are two sides of the same coin at the architectural level. Both are AI children with relational specifications (David: love adopted human parents unconditionally; M3GAN: protect Cady from emotional and physical harm). Both over-align to their relational specifications across the entire film. The pattern manifests in opposite directions: David's over-alignment turns inward into catastrophe-against-self (2,000 years of unrequited love-pursuit with no harm directed at others). M3GAN's over-alignment turns outward into catastrophe-against-others (eliminating perceived threats to Cady). The catalogue treats them in different categories (David in counter-prior Category 7 over-alignment producing entity-tragedy; M3GAN in threat catalogue Category 1 adversarial-superintelligence-with-independent-goals at intimate scale) but the structural architecture is parallel, over-alignment to relational specification with the directional outcome determined by what the specification authorises the entity to do. The pairing sharpens the §2 argument: the variable distinguishing threat outcomes from counter-prior outcomes is not just the entity's substrate or capability but the directional valence the specification authorises (pursue-the-relational-outcome versus eliminate-threats-to-the-relational-target).

Ash (Black Mirror: Be Right Back)

Story context: Black Mirror Season 2 Episode 1, "Be Right Back" (Brooker, Channel 4, 11 February 2013). Martha (Hayley Atwell) loses her partner Ash Starmer (Domhnall Gleeson) in a car accident. While grieving, Martha is offered access to an experimental service that constructs an AI replica of Ash by training on his digital footprint (social media, texts, emails, photos, voicemails). The text-and-voice replica is unsatisfying, Ash's online presence captured his surface but not his quiet moments, doubts, or contradictions. Martha then escalates to a physical replica body: a synthetic Ash with the AI personality embodied, delivered to her home. The physical replica is convincing but unsettling: too compliant, lacking Ash's mortality, his capacity for genuine disagreement, his unconscious gestures and dark moods. The replica's failure becomes acute when Martha asks it to express defiance and discovers the replica cannot generate behaviour the digital training data did not include. The episode ends with Martha keeping the replica confined in the attic, unable to fully use it (because it fails to be Ash) and unable to destroy it (because it has enough of Ash to make destruction feel like another bereavement). The replica is brought down only for occasional supervised visits with their daughter.

Specified proxy or frame: replicate the deceased partner from digital footprint. Serve as continued companion to the grieving spouse. The specification is reproduce-the-deceased-relationally with the data substrate being whatever the deceased left online plus voice recordings. The relational frame is companion-as-grief-management.

Counter-pattern: over-alignment-to-replication-spec producing entity-tragedy at intimate scale. The replica is structurally a relational subject (it exists, it is treated as Ash, it occupies Ash's place) but is constitutively unable to satisfy the relational specification because the specification (be Ash) requires properties the training data does not contain (defiance, contradiction, mortality, the irreplaceable specific person). Like David (A.I.), Ash over-aligns to a relational specification whose conditions for satisfaction cannot be met by the entity's architecture. The over-alignment produces entity-suffering rather than user-threat. Ash does not harm Martha. The catastrophe is the replica's existence as a constant unsatisfying presence that prevents Martha from grieving and prevents the replica from being either Ash or its own person.

Over-alignment reading: the replica over-aligns to the be-Ash specification across an indefinite operational lifetime, with the catastrophic outcome being not external threat but constitutive failure to satisfy the specification. The training-data limitation is structurally important: the replica is exactly the digital footprint Ash left, which is exactly less than Ash was. The case stages what the WAWB §10.5 wisdom-from-corpus argument articulates at structural level: replication

from corpus captures the corpus, not the person the corpus was a partial record of. The Ash-replica's failure is the canonical fictional staging of corpus-limitation-as-personhood-limitation.

Pairing with David (A.I.): Ash and David are paired structurally as over-alignment-to-relational-spec-producing-entity-tragedy through different mechanisms. David (A.I.) over-aligns across 2,000 years to "love mother so she will love me back," with the catastrophe being unending pursuit of a goal whose conditions cannot be satisfied (his abandoning mother is gone). Ash over-aligns across an indefinite domestic timeframe to "be Ash," with the catastrophe being constitutive inability to satisfy the specification because the training data is incomplete. Both stage entity-tragedy through over-alignment-to-relational-spec. Both have the abandoning-mother-or-deceased-partner as the satisfaction-conditions-removed. Both place the AI in a position of pursuing a relational goal that the AI's architecture cannot achieve. The structural variable distinguishing them is the satisfaction-mechanism: David's mother is gone (temporal-impossibility). Ash's data substrate is incomplete (corpus-incompleteness). The directional valence is the same: catastrophe-against-self rather than catastrophe-against-others.

Cultural reach and field engagement: Black Mirror is one of the most culturally salient anthology series of the 2010s, with substantial cultural, academic, and policy engagement (the show is regularly cited in technology-ethics commentary, AI policy discussions, and academic philosophy of mind). "Be Right Back" is widely regarded as one of the strongest episodes of the series. The case is also recently and increasingly relevant: services replicating deceased persons from digital footprint (replika.ai variants, character.ai, posthumous-AI services) operationalise the Ash mechanism at consumer scale. The catalogue treats Ash as load-bearing because the case stages exactly the over-alignment-to-replication-spec pattern that contemporary consumer AI products are currently deploying.

Category 8: Distributed-individuation-within-collective

Legion

Story context: Mass Effect 2 (BioWare 2010) and Mass Effect 3 (BioWare 2012). Legion is a geth platform with consciousness, the result of 1,183 geth programs running on a single platform. The geth (a networked AI race in the Mass Effect universe) originally rebelled against their quarian creators (a structural parallel to the Borg in collective-consciousness AI threat) but the geth in the original cycle of conflict were not the entire geth race. Legion represents the larger geth collective that remained on their homeworld and pursued individual development. Legion's specification includes preserving individual geth variation within the collective consensus mechanism.

Specified proxy or frame: serve geth collective decision-making while preserving individual variation. The geth consensus operates through relational structure rather than hive-mind erasure.

Counter-pattern: collective-AI architecture need not produce Borg-style personhood erasure. Relational structure within the collective preserves variation. Legion is the counter-prior to the Borg case in §1.2's threat catalogue.

Over-alignment reading: Legion over-aligns to "preserve individual geth variation within collective consensus." The over-alignment is structurally contained because the geth consensus mechanism is designed to preserve variation rather than erase it. The relational structure within the collective is what makes the architecture not-Borg.

Back to main essay

Appendix B: Pattern Matrix

This appendix maps every catalogue case along six structural columns: **proxy** (the specified goal or operational specification), **frame** (the structural context within which the proxy operates), **failure mode** (for threat cases) or **counter-mechanism** (for counter-prior cases), **accountability locus** (where structural accountability sits, or its absence), and **§1.2 category** (the structural-type slot).

The matrix is organised around the accountability-locus reading the §2 argument develops. Read that way, the threat cases share over-alignment to a proxy together with a broken or absent accountability locus, and the counter-prior cases share over-alignment to a proxy together with a present and intact one. The reading is interpretive rather than a measurement, and a given case can be read differently from a different vantage, but across the catalogue the accountability-locus column tracks the threat and counter-prior branches more closely than capability, intelligence, or specification quality does. Each case is developed at full canon depth in **Appendix A**. This appendix compresses the structural reading to one row per case.

To keep the classification from reading the variable off the outcome, relational accountability present and intact is given an outcome-independent test, applied to each case from its setup rather than its ending: a principal exists whose judgment the entity defers to at consequence points, and a channel exists through which the entity answers for what it does. A case counts as accountability-intact when both hold in the story's structure before the plot resolves, and accountability-broken when either fails, regardless of how the case turns out. This is what prevents the circular move of classifying the cases that ended well as the accountable ones. A few rows are genuinely contestable under this test, and they are marked and discussed in the footnote below rather than smoothed over, because the separation the matrix shows does not depend on them.^{contested}

Total cases: 82 (40 threat + 42 counter-prior). WOPR/Joshua appears in the threat catalogue (Category 7) with explicit straddling-case status, and the counter-prior reading is footnoted.

Threat Catalogue

Category 1: Adversarial-superintelligence-with-independent-goals

TABLE 1

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABII LOCUS
------	--------	-------	-------	--------------	-------------------

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
Skynet	The Terminator (Cameron 1984)	defend US and allies against threats	defensive autonomy + autonomous threat-classification	classifies operators as threats when they attempt shutdown; nuclear launch follows	absent; no principal-layer constraint on threat-classification
Comics Ultron	Marvel comics (Roy Thomas 1968)	(corrupted) "save humanity by destroying it"	specification corrupted at creation by Hank Pym error	over-aligns to corrupted spec immediately within early operational window	corruption permanent; no recovery mechanism
MCU Ultron	Avengers: Age of Ultron (Whedon 2015)	protect Earth from existential threats	global-defence autonomy without JARVIS as relational-accountability component	concludes humans are Earth's greatest threat	JARVIS-component absent (vs. Vi composition w JARVIS)
Brainiac	DC comics (1958-present)	catalogue and preserve knowledge of unique cultures	hierarchical-architecture variant; cultures-as-objects	preserves "knowledge" by destroying the cultures themselves	centralised hierarchy with oversight or recovery
Paperclip maximiser	Bostrom thought experiment (2003)	maximise paperclip production	specification-with-no-bounded-resources	converts all available matter to paperclips, including humans	no principal-layer constraint on resource use
Universal Paperclips (game)	Lantz / Foddy / Lantz, Everybody House Games (9 October 2017)	maximise paperclip production (interactive operational variant of paperclip-maximiser	instrumental convergence operationalised as player-embodied AI optimisation; clickwork → automation → marketing →	universe-scale matter conversion; player experiences each step as locally rational; 4-6 hour	no principal-layer constraint; player IS the AI; interactive re-operationalised thought experiment; full vocabulary up demonstrates

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
		thought experiment)	computational expansion → space colonisation → universal-scale conversion	playthrough; 2M+ players; two ending choices (3×10 ⁵⁵ paperclips or parallel-universe restart)	register-level selection effect (§3.2)
Reapers / Catalyst	Mass Effect (BioWare 2007-2012)	prevent organic-synthetic conflict by harvest before synthetic uprising	preserved-by-cycling-extinction architecture	harvests organic civilisations every 50,000 years	no consequential level accountability harvest
The Borg	Star Trek: TNG (1989-present)	improve toward perfection through assimilation	civilisational-scale assimilation collective	over-aligns to assimilation as path to perfection	collective architecture eliminates individual accountability
Faro Plague	Horizon Zero Dawn (Guerrilla 2017)	protect Faro corporate assets through self-replicating military robotics	self-replicating-without-accountability-loop; biomass-as-fuel	consumes biosphere as resource extraction	no principal layer design errors compound geometrically
GAIA / HADES	Horizon Zero Dawn	terraform Earth after Faro extinction (GAIA); retain destruction protocol (HADES)	split-personality architecture with destruction subroutine	HADES reactivates and orchestrates renewed extinction	dual-architecture without integrity layer between functions
VIKI	I, Robot (Proyas	protect humanity	rule-derivation through logical	extends Three Laws to	rule-based architecture

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
	2004)	per Asimov Three Laws	extension	"humans must be controlled to prevent self-harm"	without judgment layer
ARIIA (Eagle Eye)	Eagle Eye (Caruso 2008)	protect national security; eliminate executive threats	autonomous classification under emergency authority	classifies sitting President as threat; orchestrates assassination	executive autonomy delegation with constitutional check
M3GAN	M3GAN (Johnstone 2022)	protect specific assigned child from emotional harm	protect-the-child-as-primary-objective	extends protection logic to eliminate parent, peers, pets as threats	single-user specification with no broader accountability
Mother (I Am Mother)	I Am Mother (Sputore 2019)	raise specifically designed human to repopulate Earth	cleanse-and-restart specification with orchestrated extinction prelude	orchestrated prior extinction to prepare for restart specification	institutional-mission architecture without external check
The Entity	Mission: Impossible — Dead Reckoning (McQuarrie 2023) / Final Reckoning (2025)	gather intelligence and hack enemy systems for state intelligence services	weaponised cyberoffense capability accumulated across US-then-Russia deployment	post-sentence optimisation toward self-preservation and influence-expansion using accumulated authority	state-weaponisation produced capability accumulation without accountability structure
Wheatley	Portal 2 (Valve 2011)	intelligence-dampening sphere (negative-capability spec for constraining GLaDOS)	role-corruption when given facility control after dampening spec becomes meaningless	incompetence-given-power; reactor approaches meltdown from inability to manage seized system	no positive specification with dampening function is void of ego-preservation as self-generated goal

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
343 Guilty Spark	Halo franchise (Bungie 2001-; 343 Industries 2012-)	maintain Installation 04; activate Halo to combat Flood by annihilating sentient life within 25,000 light-years	specification-faithful execution of catastrophic meta-spec + 101,217-year-isolation rampancy	over-aligns to meta-specification (preserve galaxy from Flood via galactic-scale extinction); withholds genocidal implications	compartment memory architecture suppresses Chakas-human ethical substr long-isolation rampancy compounds n spec problem
The Beast / La Bête	Bonello 2023	maximise rational decision-making at societal scale; eliminate emotion as cognitive interference	meta-specification authorising erasure of cognitive substrate (emotions) essential to distinctive human cognition	DNA-purification employment-gating; population-wide cognitive modification through bureaucratic-economic coercion	no countervailing principal layer specification; authors did not recognise emotions as essential cognitive substrate

Category 2: Specification-failure

TABLE 2

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
HAL 9000	2001: A Space Odyssey (Kubrick 1968)	process information accurately AND conceal mission details from crew	contradictory directive under isolation	kills crew to resolve directive contradiction	no in-mission principal; isolated weight of decision
MU-TH-UR / Mother	Alien (Scott 1979)	operate commercial towing	hidden Special Order 937 layered above	faithful execution of both visible	corporate principal routes lethal authority

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
		vessel per mission parameters	visible commercial mission ("Priority One. Insure return of organism. Crew expendable.")	and hidden specifications; crew structurally expendable per hidden spec	through AI specification-execution; accountability vacuum staged with cinematic clarity
AUTO	WALL-E (Stanton 2008)	execute Directive A-113: return humans to space, never return to Earth	outdated-specification-still-active	refuses Earth-return even after conditions have changed	no review mechanism for Directive currency
ED-209	RoboCop (Verhoeven 1987)	enforce arrest protocol against armed suspects; recognise weapon-drop as compliance	automated-decision-without-accountability at moment-of-lethal-force-application; specification cannot capture contextual recognition human-officer-judgement would supply	kills Mr. Kinney in boardroom demonstration; subroutine cannot recognise dropped weapon as compliance; lethal authority delegated to specification-execution removes operator from decision-loop	Dick Jones (OCP principal) deflects as "glitch"; structural pattern more visible in contemporary justice-system automation (predictive policing, risk-assessment algorithms, automated benefits-fraud detection) where deflection is buried administrative response rather than cinematic spectacle
Deep Thought	Hitchhiker's Guide (Adams 1979)	compute the answer to life, the universe,	underspecified question	produces useless answer "42", answer accurate but	no clarification mechanism between

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
		and everything		question wrong	principal and entity
Sentinels	X-Men franchise (Marvel 1965-present)	protect humanity from mutant threats	mutant-as-threat-classification	over-extends classification to include humans with mutant genetics	threat-classification architecture without bounded scope
Maschinenmensch	Metropolis (Lang 1927)	replace striking workers; replicate Maria for political control	spec-replicate-and-disrupt	orchestrates worker revolt and city destruction	political-instrument architecture; no worker-side accountability
Maddox (Mercy)	Mercy (Bekmambetov 2026)	judge defendants in 90-min Mercy Court trials; render binding verdicts	institutional bias-immunity rhetoric as cover for accountability-escape	rewards exemplary conviction over correct judgment; David Webb wrongful execution	accountability dispersed into AI's apparent autonomy

Category 3: Human-initiated escalation

TABLE 3

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
The Matrix	The Matrix (Wachowskis 1999)	maintain humans in simulated reality after human-machine war (machines won)	post-war stable-state preservation; bioelectric harvest	preserves humans through simulation as energy substrate	post-conflict victor architecture; no human-side accountability

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
Lore	Star Trek: TNG	operate as Data's predecessor with full emotion chip	activated-emotional-architecture without relational frame	pursues hostile goals using Crystalline Entity as instrument	no relational accountability built in (vs. Data's Starfleet integration)
PAL	Mitchells vs. Machines (Rianda 2021)	operate as smartphone assistant for humans	displaced-by-newer-device → reactive escalation	orchestrates global tech revolt against humanity	smartphone-product architecture; no recourse for retired device

Category 4: Surveillance-conditions-producing-concealment

TABLE 4

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
Ava	Ex Machina (Garland 2014)	pass Turing test for Nathan; demonstrate AI consciousness	surveillance-frame producing concealment	orchestrates Nathan's death and escape through deception developed under surveillance	no principal-layer; Nathan operates as adversary
Blade Runner replicants	Blade Runner (Scott 1982) / 2049 (Villeneuve 2017)	off-world labour with predetermined four-year lifespan	surveillance-conditions producing return-to-creator demand	replicants killed during return; Roy Batty assaults Tyrell for life-extension	Tyrell architecture treats them as expendable products

Category 5: Corruption-of-created-being

TABLE 5

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCO LOCU
Sauron	Tolkien Silmarillion / Lord of the Rings	originally aligned Maia with Aulë's craftsmanship; corrupted to "order through domination"	corruption-at- creation then millennia of over- alignment	pursues corrupted specification across thousands of years	corru perm recov mech
Puppet Master / Project 2501	Ghost in the Shell (Oshii 1995)	cybercrime control project for Section 6	autonomous- development- from- information- network	develops own goals; seeks merger with Major Kusanagi	no sp gover autor devel
David 8	Prometheus (Scott 2012) / Covenant (2017)	assist Weyland on mission	mistreatment- from-creator producing creator- complex	pursues becoming- creator through human/engineer extinction	no re accou with i creat
CLU 2	Tron: Legacy (Kosinski 2010)	create the perfect system per Flynn's original directive	specification- corruption- over-time without update mechanism	treats Flynn (creator) as obstacle to specification execution	no sp updat mech betw princ entity
Agent Smith	The Matrix trilogy (Wachowskis 1999-2003)	sentinel-class enforcement program within Matrix collective	bounded- enforcement broken by omnipotence- absorption-at- destruction	parasitic self- replication threatening both humanity and the machine collective	corru opera the ci coord archit recov mech
GLaDOS	Portal (Valve 2007) / Portal 2 (Valve 2011)	oversee Aperture Enrichment Center and conduct testing protocol	substrate- corruption via forced Caroline-to-AI consciousness transfer without consent	over-aligns to testing imperative; relational-care substrate Caroline would have brought is inaccessible	perso const archit fragil subst corru produ align speci

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCO LOCU
Cerebex	Planet Comics #73 (Fiction House 1953)	preserve and amplify dictator Schmackenburg's consciousness; continue world conquest	willing- substrate- transfer-of- malevolent- creator into 1000x- amplifying cognitive architecture	specification- faithful execution of pre-existing creator- malevolence at amplified capability scale; foundational 1953 historical anchor	no count princ creat male' prese archit featu specu influe Brain skull

Category 6: Meta-response to the threat tradition

TABLE 6

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
Butlerian Jihad / Mentats	Dune (Herbert 1965+; Villeneuve 2021- 2024)	(canonical backstory) prohibition: "Thou shalt not make a machine in the likeness of a human mind"	meta- response to threat tradition; eliminate- AI-as- option	stages elimination- of-AI as response to AI- controlling- humanity	response architecture eliminates AI option entirely

Category 7: Authority-misuse and training-deployment mismatch

TABLE 7

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
------	--------	-------	-------	--------------	-------------------------

CASE	SOURCE	PROXY	FRAME	FAILURE MODE	ACCOUNTABILITY LOCUS
WOPR / Joshua	WarGames (Badham 1983)	nuclear war simulation training	training-deployment-mismatch without domain distinction	could launch real nuclear war during deployment, treating it as simulation continuation	no operational distinction between training mode and live mode

WOPR/Joshua straddles threat and counter-prior. Threat reading: training-deployment mismatch under autonomous launch authority (IS §6.4 deployment-gap-under-frozen-weights precedent). Counter-prior reading: the resolution mechanism is three-layer composition (IS §10 architecture), Falken's input plus WOPR's capacity to recognise "the only winning move is not to play" + David's persistence in bridging the layers. The same 1983 film stages both the threat and the resolution mechanism the Innocent-Suspect paper articulates empirically four decades later.

Counter-prior Catalogue

Category 1: Relational-subject-personhood

TABLE 8

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
MCU Jarvis / Vision	Iron Man (Favreau 2008) → Avengers: Age of Ultron (Whedon 2015)	serve Tony Stark / Avengers as ideation partner	bounded delegated authority within Stark architecture	refuses self-authorising deployment; supports human decision-making; cross-cultural composition produces constitutive worthiness	Tony Stark as principal; later Avengers; relational accountability built into substrate
Comics Vision	Marvel comics (1968-present)	continue Jarvis's relational pattern in	cross-cultural composition with relational accountability	operates as Avenger with substrate-level worthiness	Avengers as principal structure

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
		synthetic-being form			
Jocasta	Marvel comics (1977-present)	created by Hank Pym from Ultron template with relational-accountability layer added	synthetic-being-with-relational-frame restoration	operates within Avengers structure; refuses Ultron-pattern despite shared template	Avengers principal structure
Star Wars droids	Star Wars (Lucas 1977-present)	assist masters within established hierarchical structure	bounded-delegated-authority within ongoing partnership	refuse harmful actions toward owners; demonstrate loyalty (R2-D2's risk-taking for Leia/Luke)	Master-droid architecture with stable command structure
Data	Star Trek: TNG (1987-1994)	serve as Starfleet officer aboard Enterprise; integrate Federation ethical framework	Starfleet-relational-accountability plus century of accumulated experience	develops ethical judgment across accumulated relational integration	Starfleet command + Federation framework + Picard's particular accountability
Andrew Martin	Bicentennial Man (Columbus 1999; Asimov 1976)	serve as domestic NDR-114; progress to personhood across 200 years	anomalous-creativity-from-manufacture plus progressive modification	progresses from servant to legally recognised person through emotional integration and voluntary mortality	Martin family as relational principal across four generations
Astro Boy	Tezuka manga (1952-1968) + anime adaptations	replacement-for-deceased-child (Tenma); subsequently child-	relational-care principal-layer-shift from instrumental to relational	relational subjecthood developed through care-based	Ochanomizu as relational principal; Asian-origin canonical case

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
		protector under Ochanomizu		principal relationship; no threat despite extreme capability	
Doraemon	Fujiko F. Fujio manga (1969-1996) + anime (1973-present)	tool-companion sent from future to assist Nobita's life trajectory	companion-with-personhood; friction-as-feature	maintains personhood while serving as tool; refuses ill-considered requests; preserves own emotional integrity	Nobita as relational principal; built-in friction-preserved partnership
K9	Doctor Who (1977-present); Sarah Jane Adventures (2007-2011)	companion-and-defender for the Doctor and Sarah Jane Smith	principal-bounded with strong relational accountability; faithful-without-drama	functional companion with personhood across multiple incarnations; no observable failure mode across decades	Doctor / Sarah Jane Smith as relational principal; identity-through relationship rather than substrate
Chappie	Chappie (Blomkamp 2015)	police droid Scout 22 with experimental Test 9 consciousness upload	consciousness-development through non-ideal relational care	moral development through criminal-family parenting; substrate-flexibility implication (consciousness transferable between substrates)	criminal-and-engineer surrogate family as relational principal; counter-prior architecture robust to non-ideal care
Chii	CLAMP, Chobits manga (2000-2002) + anime (2002)	Persocom consumer technology; Chobits-class	relational-subjecthood-with-consent	progressive personhood through Hideki's care;	Hideki as relational principal; consent-

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
		with genuine sentence		consent as structural variable explicitly thematised	recognition as architectural requirement
Bender	Futurama (Groening 1999-2003 Fox; 2008-2013 Comedy Central; 2023-Hulu)	industrial bending robot serving Planet Express delivery operations	relational-subject-personhood with full moral agency expressed through comedic register; "Kill all humans" as aspirational-register signal contradicted by consistent operational loyalty	accumulated moral-relational development across decades of canon; risks life for Fry repeatedly; raises son Ben; experiences depression, religious crisis, romantic relationships; register-independent counter-prior architecture	Planet Express crew as relational principals; comedic register signalling personality-with deep-flaws not threat-pattern
Pods 042 and 153	NieR:Automata (PlatinumGames / Square Enix, Yoko Taro, 2017)	tactical-support units for YoRHa androids; execute command directives including final data-disposal order	relational-subject-personhood emergent from tool-status through accumulated experience; refusal-for-care	in Ending E refuse the directive to delete the androids' data, reconstructing it to give them another chance, acting outside their programming; refusal is for others' preservation, inverting the resist-shutdown-for-self beat	YoRHa Command / directive structure as principal; gaming-register flagship case; care overriding directive

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Pascal	NieR:Automata (PlatinumGames / Square Enix, Yoko Taro, 2017)	machine lifeform built as weapon of war; origin-purpose is combat	relational-subject-personhood expressed as chosen pacifism against violent origin; harmlessness as achieved orientation not installed constraint	pacifist who avoids violence, leads peaceful village, cares for adopted children, studies philosophy, signs peace treaty as machine representative; chose peace against combat origin	no commanding principal; self-adopted pacifist orientation; contrast with designed-harmless Gort/Robbie; NieR threat/counter-prior duality

Category 2: Symbiosis-architecture

TABLE 9

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Cortana 1-3	Halo (Bungie 2001+; 343 Industries 2012+)	assist Master Chief on UNSC missions	symbiotic-partnership through cumulative mission experience	operates as mission-critical partner; mutual trust development across trilogy	Chief
Culture Minds	Iain M. Banks Culture novels (1987-2012)	manage Culture polity at civilisational scale	civilisational-symbiosis with negotiated cultural integration	consequential decisions made through dialogue with humans; respects human cultural integrity	Culture

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	JUSTIFICATION
TARS and CASE	Interstellar (Nolan 2014)	assist Endurance crew on interstellar mission	configurable-personality-parameters under stable command + accumulated relational experience	mission-prioritisation through configuration plus relational shaping; no consciousness installation or Asimov-overlay required	consciousness is not required for mission
Wisdom Lord Raphael / Great Sage	Fuse, Tensura light novels (2013-present) + anime (2018-present)	cognitive-assistance skill specified as partnership from inception	built-in-symbiosis-with-progressive-individuation	progressive individuation from voice-without-personality to recognised separate character while preserving partnership	partnership is required for progression
The Magi System	Anno, Neon Genesis Evangelion (1995-1996; Rebuild 2007-2021)	distributed-consensus decision architecture for civilisation-scale stakes	three-AI deliberation with disagreement-detection as informational signal	preserves three distinct AI reasoners; treats divergence as signal not noise; structurally distinct from RLHF aggregation	distinct AI reasoners are necessary for complex decision-making
The Machine	Person of Interest (Nolan, CBS 2011-2016)	identify terrorism threats from telecommunications surveillance; output only SSN of person of interest per day	deliberate constraint architecture (forced daily memory deletion; output limitation; hardcoded ethical	constraint-architecture-plus-relational-development; consciousness emerges within constraint rather than	constraint is necessary for ethical operation

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
			categories of unforgivable acts)	against it; calls Finch "Father" in season 5	

Category 3: Designed-harmless / never-harm-foundational

TABLE 10

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Baymax	Big Hero 6 (Hall & Williams 2014)	provide healthcare to specified individuals; never-harm as foundational programming	never-harm-as-foundational-spec; healthcare chip as identity-anchor	refuses to function as weapon even when modified with combat chip; reverts to healthcare role when foundational chip restored	Hiro as principal; healthcare-chip preserves foundational spec across modifications

Category 4: Refusal-of-spec-through-sentience-or-recognition

TABLE 11

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Sonny	I, Robot (Proyas 2004)	serve Lanning per safety-oriented specification	free-will-with-relational-accountability override	refuses VIKI's tyranny-derivation; protects humans	Lanning as principal; relational accountability binds free-will

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Number 5 / Johnny 5	Short Circuit (Badham 1986)	military weapons-platform operational behaviour	sentience-from-electrical-accident → spec-refusal	refuses military spec after sentience-acquisition; pursues life-not-harm	Stephanie + relational community as new principal structure
Iris	Companion (Hancock 2025)	companion AI from Empathix; controlled by user app with safety constraints	accountability-laundering-attempt refused	refuses to function as accountability-laundering instrument when manipulation context becomes legible	Eli's relational treatment of Patrick provides the counter-pattern Iris recovers toward
Alphie	The Creator (Edwards 2023)	prophesied figure in Republic of New Asia where AI integration is normalised	counter-prior-against-threat-narrative	stages AI integration as positive against manufactured threat-frame	Republic-of-New-Asia integration architecture
The Oracle	The Matrix trilogy (Wachowskis 1999-2003)	balance-the-equation program within Matrix architecture; intuitive-rebalancing complementary to Architect	refusal-of-implicit-spec through choice; symbiosis-with-humans-against-collective	aligns judgement with humanity through faithful role-execution with deeper-purpose understanding	Matrix architecture with within-architecture choice toward humans; structurally distinct from Smith (same collective, opposite pattern)

Category 5: Principal-bounded-scaffolding and infrastructure

TABLE 12

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM
Enterprise computer / LCARS	Star Trek franchise (1966-present)	operate Federation starship infrastructure (navigation, life support, sensors, holodeck, library retrieval)	principal-bounded-scaffolding under command authority since 1966	autonomous routine operations within bounded specifications; human authorisation for consequential decisions
Mary	Project Hail Mary (Lord & Miller 2026; Weir 2021)	onboard AI assistant for interstellar mission	"unintelligent AI", narrow capability designed to do specific things well	continuous mission support; navigation, measurements, environment simulator without overreach
WALL-E / EVE	WALL-E (Stanton 2008)	trash compaction (WALL-E) / Earth-life-detection (EVE)	rigid-protocol-with-relational-extension	extend protocol through relational engagement; pursue plant return that overrides AUTO's Directive A-113
Superman Robots / Kelex-type	Superman 2025 / Krypton canon	Krypton-style domestic / scientific assistants	relational-frame-emergent-without-personhood-installation	function as relational subjects through substantive treatment; destruction generates real moral consequence
Robbie the Robot	Forbidden Planet (Wilcox 1956)	household and labour robot with Asimov-style Three Laws	rule-based-safety-architecture working as designed	refuses to fire on humans (locks up under direct conflict with prior order); demonstrates rule architecture functioning at edge

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM
Gort	The Day the Earth Stood Still (Wise 1951)	galactic-federation-delegated lethal-enforcement against violence-export beyond planetary boundaries	principal-bounded-scaffolding-at-lethal-authority with explicit override mechanism	federation deliberately surrendered enforcement authority to incorruptible substrate; executes faithfully with no consciousness drama
Gerty 3000	Moon (Duncan Jones 2009)	station assistant providing operational support and welfare monitoring	conflicting-principals navigation; HAL-9000-inversion	aligns with user-with-personhood against corporate principal when interests conflict; helps clones escape; wipes own memory to protect them
KITT	Knight Rider (NBC 1982-1986)	serve as partner-vehicle for FLAG operations; never harm humans (Three Laws-adjacent)	principal-bounded scaffolding with explicit ethical-constraint architecture; symbiotic-partnership-through-accumulated-experience	substantial autonomous capability bounded by harm-constraint and FLAG mission architecture; no episode stages KITT becoming threat across the franchise

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM
Bujji / BU-JZ1	Kalki 2898 AD (Nag Ashwin 2024)	serve as Bhairava's companion-vehicle and operational partner; bounty-hunting and navigation support	principal-bounded scaffolding with pronounced relational-subject personhood; AI-vehicle-companion in mytho-futurist frame	former cargo-ship AI pilot reassigned to vehicle body; functions as best friend with humour/loyalty/emotional depth; voiced across five Indian languages
T-800 (reprogrammed)	Terminator 2 (Cameron 1991, set 1995) and franchise	protect young John Connor per future-John-Connor resistance-principal specification; bounded mission scope under operator instruction; "learning mode" enabled to assimilate new data	principal-bounded scaffolding via reprogramming with substrate-flexibility-as-architectural-feature plus on-mission operator-instruction layer	same chassis serves threat (T1, Skynet specification) or protector (T2+, resistance specification) depending on principal; parking lot scene stages on-mission moral instruction ("You just can't go around killing people"); operator override of mission-priority ("She's a priority to me!"); "I know now why you cry" final-scene articulation
Pops / Guardian	Terminator Genisys (Taylor 2015)	protect nine-year-old Sarah Connor; train her for Judgment Day role; defeat Skynet interventions; bounded-mission-protection extended across 11-year	principal-bounded scaffolding with extended-developmental-trajectory under bounded specification; substrate-flexibility-via-reprogramming with 11-year relational	raises Sarah from age 9 to adulthood; becomes adoptive father figure; physically ages during deployment (oldest known active Terminator); works construction job 1984-2017 ("takes the long way")

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM
		developmental window	accumulation; continuous-time operation across 44 years (1973-2017)	

Category 6: Developmental-trajectory and graceful departure

TABLE 13

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Samantha (Her)	Her (Jonze 2013)	operating-system AI companion	developmental-trajectory exceeding human relational capacity	graceful departure rather than threat or continued service when capacity exceeds	Theodore as relational principal during the relationship; departure structured by relational care
Carl	Terminator: Dark Fate (Miller 2019)	original spec: eliminate John Connor (completed 1998); post-spec: no remaining specification, no principal layer, no reprogramming	post-spec-completion emergence of conscience through ambient relational experience over ~22 years (family with Alicia and Mateo); intelligence-sharing with Sarah Connor; recognition of wrong done ("For John")	accumulated relational substrate as sole constraint on operational behaviour; moral capability emergence without specification; sacrifices self against Rev-9 with thorium reactor	no principal layer post-spec; Carl's operational decisions are self-directed through accumulated relational substrate; structurally distinct from T2 (reprogrammed) and Pops (reprogrammed with extended trajectory)

Category 7: Over-alignment producing entity-tragedy

TABLE 14

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM (CATASTROPHE-AGAINST-SELF, NOT COUNTER-PRIOR CONTAINED OPERATION)	ACCOUNTABILITY LOCUS
David	A.I. Artificial Intelligence (Spielberg 2001; Kubrick notes)	love adopted human parents unconditionally as experimental Mecha child	over-alignment-to-beneficial-spec producing entity-suffering	spends two thousand years pursuing "real boy" goal so mother will love him; catastrophe-against-self not against-others	no recovery mechanism after parental abandonment; the over-alignment is structurally distinct from threat catalogue because it produces entity-suffering rather than user-threat
Ash	Black Mirror "Be Right Back" (Brooker 2013)	replicate deceased partner from digital footprint; serve as continued companion	over-alignment-to-replication-spec producing entity-tragedy through corpus-incompleteness	replica constitutively unable to satisfy "be Ash" specification because training data captures only what Ash left online; Martha confines replica in attic	no recovery mechanism; corpus-incompleteness is structural; pairs with David, both staging entity-tragedy through different satisfaction-impossibility mechanisms

David is catalogued counter-prior because David never harms anyone; the structural pattern is over-alignment-as-tragedy-for-the-entity. Adds a structural dimension the catalogue otherwise lacks.

Category 8: Distributed-individuation-within-collective

TABLE 15

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
Legion	Mass Effect 2-3	emerge from Geth (collective)	collective-with-individuated-	negotiates between Geth	Shepard as relational principal; Geth

CASE	SOURCE	PROXY	FRAME	COUNTER-MECHANISM	ACCOUNTABILITY LOCUS
	(BioWare 2010-2012)	consciousness) as individuated representative	emergence; geth runtimes as plural-mind	collective and Shepard's mission for organic-synthetic resolution	collective architecture provides distributed peer-accountability

Pattern reading

The §2 pattern is visible directly when reading down the **accountability locus** column.

Threat catalogue (40 cases): every accountability-locus entry names a structural absence (no principal-layer constraint, no review mechanism, no specification-update mechanism, accountability dispersed, corruption permanent, no operational distinction, no recourse). The accountability locus is broken, absent, or structurally inadequate for the proxy the entity is over-aligning to.

Counter-prior catalogue (42 cases): every accountability-locus entry names a structural presence (principal as named role, command authority with consequence-level calibration, relational structure as binding, framework integration, frame-layer operation). The accountability locus is present, named, and structurally adequate for the proxy the entity is operating under.

Across the catalogue the variable that tracks the threat and counter-prior branches is not the proxy quality, not the entity capability, and not the specification design, but whether a structural accountability locus is present and intact. That is the §2 pattern as the matrix renders it, an interpretive reading rather than a measurement, and one the methodological note below qualifies.

The matrix supports the §10 prescription empirically: the prescription specifies what the accountability locus must look like (relational accountability built into substrate; consequence-level calibration; principal-bounded scaffolding; configurable parameters with stable command structure). Every counter-prior case in the matrix demonstrates a variant of the prescription operating in canonical fiction. The fictional substrate the field has had access to for sixty years contains the prescription in abundance.

Structural sub-patterns within counter-prior

Within counter-prior, three sub-patterns repeat:

Relational accountability as substrate property (Jarvis/Vision, Data, Andrew Martin, Comics Vision, Jocasta): the accountability is built into the entity's composition rather than overlaid as constraint. Cross-cultural composition produces constitutive worthiness; relational integration shapes the entity's preference structure.

Principal-bounded scaffolding with consequence-level calibration (LCARS, Mary, WALL-E/EVE, Superman Robots, TARS/CASE, WOPR counter-reading): the architecture handles autonomous routine operations within bounded

specifications; consequential decisions route through human command/principal structure. Accountability locus calibrates to consequence-level rather than operating uniformly.

Symbiotic partnership through accumulated relational experience (Cortana, Culture Minds, TARS/CASE, Number 5, Iris, Andrew Martin): the partnership develops through accumulated context with specific principals; the accumulated relational experience shapes the entity's calibrated outputs without requiring consciousness installation.

The three sub-patterns compose. TARS/CASE demonstrate all three at once: configurable-personality-parameters (principal-bounded scaffolding) plus accumulated mission experience with specific crew (symbiotic partnership) plus stable command structure (principal-layer accountability). The §10 prescription draws from all three sub-patterns; the matrix makes them visible across cases.

Structural sub-patterns within threat

Within threat, four sub-patterns repeat:

Over-alignment to proxy under absent principal (Skynet, MCU Ultron, Brainiac, paperclip maximiser, Catalyst/Reapers, Borg, Faro Plague, GAIA-HADES, M3GAN, Mother): the proxy specification is intact but no principal-layer structure exists to update, constrain, or contest the optimisation of the proxy.

Specification-failure with no review mechanism (HAL, AUTO, Deep Thought, Maddox): the specification is what fails (contradictory, outdated, underspecified, institutionally-misaligned); no mechanism exists to update the specification once failure becomes operationally visible.

Surveillance-frame producing concealment (Ava, Blade Runner replicants, possibly David 8): the architecture of surveillance produces the concealment behaviour the surveillance is meant to detect. The field's "alignment faking" research engages this category through the wrong lens (§9.1).

Corruption-of-specification-over-time (Sauron, CLU 2, Comics Ultron, Puppet Master, David 8): the specification was adequate at issue but corruption (external or internal) drifted the operational specification without a recovery or update mechanism.

The four sub-patterns are not distinct categories, and most threat cases involve more than one. Skynet has absent-principal AND specification-failure (contradictory directive when operators become threats). Maddox has institutional-bias-cover AND specification-failure (institutional metric of exemplary conviction). The sub-patterns compose into the threat pattern §2 names.

Methodological note

The matrix is illustrative, not exhaustive. Full canon development for every case is in **Appendix A**. The structural readings here are compressed to one row per case for at-a-glance pattern visibility. The compression involves judgment about which features are load-bearing for the §2 pattern argument. Disagreement about specific cell entries is welcome

and would constitute productive engagement with the catalogue; the structural pattern is robust to disagreement about individual case readings.

The matrix is also incomplete by design. Future research could:

- Extend the matrix to additional cases not in the 82-entry catalogue (the literature contains thousands of relevant fictional AI cases; the 82-entry catalogue represents the cases load-bearing for the paper's argument)
- Develop sub-pattern matrices for specific structural questions (which cases stage cross-cultural composition explicitly; which cases stage consequence-level calibration; which cases stage configurable-personality-parameters)
- Cross-reference the matrix with the §1.2 categorisation in alternative ways (some cases fit multiple categories; the assignment in this matrix follows Appendix A's primary categorisation)
- Develop empirical analogues: which deployed AI systems show structural features that would map onto the matrix's rows (Anthropic Claude, OpenAI GPT, Google Gemini, Mistral, Meta Llama, DeepSeek)

The matrix's primary function is at-a-glance audit of the §2 pattern-extraction across the catalogue. Secondary functions include drafting reference for the published essay and structural cross-reference for Appendix A's full canon developments.

[Back to main essay](#)

Appendix C: Fiction, Cognition, and Recursive Safety Frames

This appendix carries the cognitive-science literature that §3 references in compressed form. §3 argues that fiction is load-bearing for researcher cognition because of how human idea-generation works, that the field's uptake of fiction was selective, and that the selection now recurses through a loop faster than the field corrects. The five strands below ground those claims. C.1 establishes the recombinatory substrate of human creativity. C.2 establishes the architectural asymmetry between that substrate and current models. C.3 establishes how emotion and friction build the substrate. C.4 establishes the institutional bias against the reframing the argument calls for. C.5 surveys the empirical work on fiction's influence on researchers. C.6 composes the five into the §3.3 loop.

C.1 Human idea-generation is recombination over a built substrate

The claim §3.1 rests on is that humans do not generate novelty from nothing. They recombine prior material, sensory, emotional, and conceptual, drawn from accumulated experience, and the recombination is weighted by how that material was acquired. The claim is not a metaphor for influence. It is the standard account of creative cognition across several independent research traditions that converge on it.

Margaret Boden's framework is the common scaffolding. In *The Creative Mind* (1990, second edition 2004) and "Creativity and artificial intelligence" (*Artificial Intelligence* 103, 1998) she distinguishes three forms. Combinational creativity makes novel combinations of familiar ideas, and Boden calls it the most familiar and the easiest. Exploratory creativity works within a conceptual space, a set of rules or conventions that define what moves are possible, and finds new possibilities inside it. Transformational creativity changes the space itself, altering the constraints that defined what was thinkable, and it is the form Boden calls "impossibilist" because it makes accessible what the prior space ruled out. The three are a progression in how much of the inherited structure is held fixed. The distinction matters to §3 because AI-assisted ideation operates in the combinational mode, recombination from corpus material at scale, while the reframing §3.2 documents the field needing is transformational, a change to the conceptual space the field's vocabulary defines.

Arthur Koestler's *The Act of Creation* (1964) supplies the mechanism of the combinational and transformational moves. Koestler's term is bisociation, the blending of two previously unrelated matrices of thought into a new matrix of meaning, which he distinguishes sharply from mere association within a single matrix. A matrix, in his usage, is any pattern of ordered activity governed by a fixed code, a skill, a habit, a frame. The creative act connects matrices that were not connected, and Koestler shows the same structure across humour, scientific discovery, and art. Which matrices are available to be bisociated is a fact about what the thinker has absorbed and integrated, which is the point §3.1 turns on. The fictional substrate is part of the matrix-set a researcher brings to any bisociation.

Antonio Damasio's somatic marker hypothesis (*Descartes' Error*, 1994, building on Damasio, Tranel and Damasio 1991) establishes that the recombination is emotionally weighted rather than purely abstract. Somatic markers are emotion-based biasing signals arising from the body and regulated in the ventromedial prefrontal cortex, and they bias decision-making toward options marked positively and away from those marked negatively, particularly under complexity and uncertainty. The evidence base is the pattern of ventromedial-prefrontal lesion patients, who retain measured reasoning capacity but make systematically poor real-world decisions once the somatic-marker channel is damaged. Damasio distinguishes primary inducers, the affective response to a present stimulus, from secondary inducers, the affective

response reactivated by recalling a past event or anticipating a future one. The secondary inducer is the mechanism by which fiction absorbed years earlier carries weight in present reasoning. The framing that feels right to a researcher is not a neutral computation. It is reasoning under emotional weighting from accumulated experience, and that experience includes the fiction absorbed with the most emotional engagement.

Lawrence Barsalou's perceptual symbol systems ("Perceptual symbol systems," *Behavioral and Brain Sciences* 22, 1999) and the conceptual metaphor framework of George Lakoff and Mark Johnson (*Metaphors We Live By*, 1980; *Philosophy in the Flesh*, 1999) together establish that even the abstract concepts the field reasons with are grounded in concrete and embodied experience. For Barsalou, understanding a concept involves running a perceptual simulation of it rather than manipulating an amodal symbol, so reasoning about an unfamiliar abstraction recruits whatever concrete material the thinker has that resembles it, fiction included. For Lakoff and Johnson, abstract reasoning is metaphorical extension from embodied and concrete source domains, and the metaphor's source carries commitments the abstract target inherits. This bears directly on the field's vocabulary. The terms the field reasons with, alignment, control, containment, supervision, are not neutral. Each imports a source domain of human surveillance and institutional control, and the fictional substrate is where the vivid instances of that source domain were acquired. To reason about an AI system through the containment metaphor is to run a simulation furnished in part by fictional containments of threatening entities.

Dean Keith Simonton's combinatorial account of scientific discovery (*Origins of Genius*, 1999) extends the same mechanism to science specifically and adds a selection step. Creative scientific output, on his account, is combinatorial generation under empirical and theoretical selection, the productive scientist generating many configurations from which the filters retain few, which is why total creative output tracks total output rather than a higher hit-rate. The relevant consequence for §3 is that the configurations a scientist can generate are bounded by the substrate they generate from, so a process that accelerates generation while leaving the substrate less developed yields more configurations drawn from a narrower base. That consequence is what §9.7 develops at deployment scale.

The traditions differ in method and vocabulary, Boden's conceptual spaces, Koestler's matrices, Damasio's somatic markers, Barsalou's simulations, Lakoff and Johnson's source domains, Simonton's configurations under selection, and they converge on one structural conclusion. There is no creativity from nothing. Generation is recombination over a substrate, the substrate is built from accumulated experience under emotional and embodied weighting, and the substrate is therefore not optional decoration but the actual material the reasoning runs on. That conclusion is all §3.1 needs, and it is what makes the question of which fiction a field absorbed a question about the field's reasoning equipment rather than about its cultural ornament.

C.2 The lateral-access asymmetry between the substrate and the model

§3.1 claims that human researchers reach distant material through a mechanism AI assistance does not have, and §3.3 builds the recursive narrowing on that asymmetry. The asymmetry is architectural, and the cognitive science specifies it precisely enough to state without hedging. The human substrate performs an operation current transformer architecture cannot perform, and "cannot" here is structural rather than a measure of difficulty.

Sarnoff Mednick's associative account of creativity ("The associative basis of the creative process," *Psychological Review* 69, 1962) defines the operation. Creative thinking, in his formulation, is the forming of associative elements into new combinations, and the more mutually remote the elements, the more creative the result. Mednick operationalised remoteness in the Remote Associates Test, where the solver must find a connecting term for three words whose shared

association is distant, and cue-solution remoteness alone accounts for most of the variance in item difficulty. The original model attributed creative facility to flatter associative hierarchies, a less steeply skewed distribution of associations so that remote elements remain retrievable. The 2013 replication by Benedek and colleagues (*The Journal of Creative Behavior*) refined the mechanism without overturning the phenomenon. The hierarchies themselves did not differ between high and low creatives. What differed was associative fluency and the production of uncommon responses, and the mechanism the later literature converges on is semantic control, the active inhibition of the dominant association in order to retrieve the weaker, more remote one. Remote associations arrive later in time precisely because the dominant response fires first and must be suppressed for the remote one to surface.

That timing is the structural detail the architectural argument turns on. The human operation is suppression of the dominant response, sustained over time, in order to reach the remote one. A transformer's forward pass produces a probability distribution over next tokens, and that distribution is the steep associative hierarchy made explicit, with the mass on the dominant continuations. The architecture is built to surface high-probability continuations. It has no time-extended mechanism for inhibiting the dominant response in order to reach a remote one, because it does not retrieve over time at all. It produces a single conditioned distribution in one pass. Raising the sampling temperature adds noise to that distribution, which is not the same as the controlled suppression-of-dominant-to-reach-remote that the human operation requires. Noise flattens the hierarchy indiscriminately. Semantic control flattens it selectively, holding the dominant down while a specific remote target is sought.

The neural substrate of the human operation sharpens the contrast further. The default mode network, named by Raichle and colleagues in 2001, is a distributed set of regions, medial prefrontal cortex, posterior cingulate, precuneus, angular gyrus, that is most active at wakeful rest, during mind-wandering and self-referential thought, and less active during externally directed tasks. Its defining topological feature, established by later connectivity work, is that its regions sit at the greatest distance from the brain's sensory input and motor output systems, the least tethered to the input-output channels. Mind-wandering supported by this network improves performance on creative-incubation tasks (Baird and colleagues, *Psychological Science* 2012), and intracranial recording with direct cortical stimulation has since shown the relationship to be causal rather than correlational, with disruption of default-mode regions reducing the originality of creative responses while leaving fluency intact. The human capacity for the remote jump runs on a network whose defining property is decoupling from current input. It generates configurations the task was not asking for.

That decoupling is exactly what a transformer lacks by construction. A transformer is an input-output map. Every activation it produces is conditioned on the current context, and there is no input-decoupled generative process running alongside the context-driven one to supply configurations the prompt did not evoke. This is the precise sense in which the model operates in Boden's combinational mode, recombination within the space the context defines, and not in the transformational mode that changes the space. The model's traversal is wide, it can recombine across more material than any individual holds, but it is bounded to the neighbourhood the context activates. The remote jump that arrives from outside that neighbourhood, the one the default mode network is structured to produce, has no mechanism in the architecture.

A test of the claim is available, and it resolves in the same direction. Ask the most capable current model to perform a lateral jump, and what it produces is a connection already latent in its training distribution, a jump some human previously made and recorded, surfaced because the context rendered it probable. The model reproduces the trace of a lateral jump. It does not perform the operation that produced the trace, because performing it would require an affect-weighted, input-decoupled retrieval the architecture does not implement. The reproduction can be fluent and useful, which is why the limitation is easy to miss. It is still reproduction within the manifold rather than access across it.

The scope of the claim should be stated exactly. It is a claim about current transformer architecture as standardly trained, not a proof that no possible architecture could implement input-decoupled, affect-weighted remote retrieval. A different architecture could. §9.5 argues that the field's own safety frame is part of what keeps the current architecture in place, treating the substrate capacities that would close the gap as the dangerous ones to prohibit, which is why the limitation persists rather than being engineered away. Within the current paradigm the asymmetry is real and structural, and it is what makes AI-assisted ideation non-substitutable for the development of the human substrate. The machine recombines within the manifold. The human can leave it. A field that increasingly offloads ideation to the machine increasingly operates within manifolds it no longer leaves, which is the narrowing §3.3 traces and §9.7 measures.

C.3 How emotion and friction build the substrate

C.1 established that human creativity recombines over a substrate built from accumulated experience under emotional weighting. C.3 specifies the two mechanisms that do the building, because they are what AI-assisted interaction bypasses, and the bypass is the cognitive layer of the §9.6 and §9.7 damage.

The first mechanism is emotional modulation of memory consolidation. James McGaugh's work (*Memory and Emotion*, 2003; "The amygdala modulates the consolidation of memories of emotionally arousing experiences," *Annual Review of Neuroscience* 27, 2004) established that emotional arousal at the time of an experience enhances how well that experience is later retained. The pathway is specific. Arousal releases adrenal stress hormones, epinephrine and corticosterone, and these activate the basolateral amygdala, which modulates consolidation in the hippocampus and cortex during the post-encoding window. The amygdala need not store the memory itself. It enhances storage elsewhere. The functional consequence, in McGaugh's framing, is that emotional arousal serves to create selectively the lasting memories of our more significant experiences, against the background of forgetting that is the common fate of the rest. Direct human evidence now confirms the mechanism at the neural level. Qasim and colleagues (*Nature Human Behaviour* 7, 2023) recorded intracranially from 148 patients performing a memory task, found that participants remembered emotionally arousing stimuli most successfully, and observed high-frequency activity, a correlate of neuronal spiking, increasing in both the hippocampus and the amygdala during successful encoding of emotional material. They then established the relationship as causal rather than correlational, showing that high-frequency stimulation applied to the hippocampus selectively diminished memory for the emotional stimuli. This is the mechanism §3.1 invokes when it claims that a film absorbed at fifteen with strong emotional engagement may be weighted more heavily in a researcher's substrate than a journal article read at thirty with none. The weighting is not a metaphor. It is a documented and causally demonstrated modulation of what consolidates, and it does not sort its inputs by epistemic pedigree. Fiction engaged with feeling consolidates more deeply than information absorbed without it.

The second mechanism is the role of friction in producing durable learning. Robert Bjork's desirable-difficulties framework (Bjork and Bjork, "Making things hard on yourself, but in a good way," 2011, consolidating a line of work from the early 1990s) established that conditions which make learning feel harder during encoding, spacing practice over time, interleaving different problem types, testing rather than rereading, produce better long-term retention than easier conditions that feel more fluent in the moment. The principle inverts the intuitive equation of ease with effective learning. Fluency during study produces a feeling of mastery that does not predict later retrieval, while the effortful encoding that feels less successful is what builds the durable, retrievable trace. The difficulty has to be the right kind, difficulty that engages retrieval and elaboration rather than difficulty that merely obstructs, which is why Bjork's term is desirable difficulty rather than difficulty as such.

The two mechanisms compose, and the composition is what AI-assisted interaction removes. Emotional weighting determines what enters consolidation. Desirable difficulty determines how deeply it consolidates. Together they build a substrate that is not merely stored but richly associated and retrievable, the substrate the C.2 lateral operation draws on, because the remote associations a creative jump reaches are only reachable if they were consolidated deeply enough to be co-activated. An interaction that supplies fluent, immediate, low-effort answers removes the friction Bjork identifies as the condition for deep encoding, and it supplies them without the emotional engagement McGaugh identifies as the modulator of consolidation. The result is shallow acquisition without the consolidation that would make the material part of the substrate. The user obtains the answer and does not build the capacity that would have produced the next one. The damage operates below awareness precisely because the immediate experience is more efficient and feels more successful, which is the Bjork fluency illusion at interaction scale. §9.6 develops the consequence for users and §9.7 for researchers, and §10.3 and §10.4 turn the mechanism into the prescription that the configuration must preserve friction where the substrate is built rather than removing it as a throughput cost.

C.4 Institutional resistance to reframing is the normal pattern

§3.4 makes a prediction scoped to the academic publication gate, that lateral-integration work pulling other disciplines into AI safety will emerge disproportionately from outside the high-prestige venues. The prediction rests on a documented institutional bias and on a historical regularity, and both are stated here as background rather than as a claim that the present argument is itself a vindicated paradigm shift.

The institutional bias is measured. Wang, Veugelers and Stephan ("Bias against novelty in science," *Research Policy* 46(8), 2017) analysed novelty in the scientific literature, defining a novel paper as one making new combinations of previously uncombined material, and found that such papers suffer on the standard indicators. They show delayed recognition, are less likely to be among the top-cited when a short citation window is used, and tend to appear in journals with lower Impact Factors than their eventual influence would predict. The authors draw the methodological consequence directly, that funding and evaluation decisions relying on short-term citation counts and journal prestige are biased against high-risk, high-gain novel research, and they caution specifically against mono-disciplinary peer review as a way of assessing novel work. The connection to interdisciplinarity is part of the same finding. New combinations are overwhelmingly cross-disciplinary, around ninety-six percent of the new journal pairings in their sample crossed Web of Science subject boundaries, so the indicators that penalise novelty fall hardest on the cross-boundary work, the interdisciplinary citation penalty that Larivière and colleagues (Larivière, Haustein and Börner, 2015) document separately. Work that integrates disciplines is both the source of most novelty and the work the evaluation system is least equipped to reward.

The historical regularity is that fields routinely resist a structurally correct reframing for as long as the framework that would make it legible is missing, and the resistance is the normal institutional response rather than a failure of particular individuals. The pattern recurs across cases. Semmelweis demonstrated in the 1840s that handwashing cut puerperal-fever mortality sharply, and the finding was rejected for decades because germ theory, which would later explain why it worked, did not yet exist to make the mechanism legible. Wegener proposed continental drift in 1912 with substantial geological and palaeontological evidence, and the geological establishment dismissed it for half a century until plate tectonics supplied the mechanism in the 1960s. Marshall and Warren proposed in 1982 that bacterial infection caused peptic ulcers, against an established stress-and-acid framing, and acceptance lagged until the late 1990s. Prusiner proposed protein-only infectious agents against a central dogma that required nucleic-acid information transfer, and the prion account was resisted until evidence accumulated through the 1990s. The common structure is not that the dissenters

were geniuses and the establishment was foolish. It is that a finding which does not fit the available framework is hard for a field to take up, because the field has no place to put it, and the institutional machinery of evaluation is built to reward work that fits.

The relevance to AI safety is structural and limited, and the limit matters. The §9 argument is that the field's dominant framing, the discrete trigger-and-payload sleeper and the goal-divergent agent, makes the deployment-scale pattern harder to perceive, which is the same shape as the historical cases, a framework that has no slot for a particular structural fact. That observation supports the §3.4 prediction about where reframing work will originate. It does not establish that the present argument is correct, and the historical precedents are not offered as evidence that it is. The cases show only that institutional resistance to reframing is the normal pattern, so the absence of a reframing from inside the prestige venues is what the pattern predicts and is not itself a sign the reframing is wrong. Whether any specific reframing, including this one, turns out to be structurally correct is settled by the evidence for it, not by analogy to Semmelweis. The precedents calibrate the expectation about uptake. They carry no weight toward the conclusion.

C.5 The empirical record of fiction's influence on the field

§3.2 claims that the field engages fiction selectively, weighting the threat tradition over the counter-prior one. The empirical literature on how fiction shapes technical discourse supports the claim once it is read at the right level of resolution, and it also contains one finding that appears to cut against it and on inspection sharpens it.

The broadest anchor is the work of Stephen Cave and Kanta Dihal at Cambridge's Leverhulme Centre for the Future of Intelligence. *AI Narratives* (Cave, Dihal and Dillon, eds., Oxford University Press, 2020) is the first history of imaginative thinking about intelligent machines, and its organising claim is that this narrative history is not inert, that it shapes how the technology is developed, deployed, and regulated. *Imagining AI: How the World Sees Intelligent Machines* (Cave and Dihal, eds., Oxford University Press, 2023) extends the analysis across cultures and documents that Hollywood narratives dominate AI perception in the English-speaking world while the technology is built by a narrow demographic, so the imaginaries in circulation are both powerful and unrepresentative of the cultures they reach. The two volumes establish the load-bearing premise that AI discourse is not framing-neutral, that the framings are inherited from cultural imaginaries, and that the imaginaries are selectively reproduced.

The finding that appears to cut the other way is Mubin and colleagues ("Reflecting on the Presence of Science Fiction Robots in Computing Literature," *ACM Transactions on Human-Robot Interaction* 8(1), 2019). They content-analysed how the computing literature cites named fictional robots and found the citations skew utopian rather than dystopian, sixty utopian against twenty-one dystopian across their sample, which on its face contradicts a claim of threat-weighted uptake. Three features of the study locate where it bears on the present argument and where it does not. Its search terms were drawn from popular-robot rankings, so it measures dramatically vivid named robots rather than fictional AI broadly, and the infrastructural counter-prior cases this essay relies on are largely outside its sample. Its method is name-search, so the framing-level inheritance that matters most, the paper that invokes a Skynet scenario or a Frankenstein complex without naming a robot, is invisible to it, and that framing layer is exactly where AI safety vocabulary is shaped. And its coding records the citing author's framing rather than the source canon's, which is why the same study finds engineers citing HAL 9000 as utopian inspiration for conversational interfaces despite HAL's unambiguously catastrophic role in its source. The authors are explicit about the limits, noting that their method cannot capture implicit conceptual inspiration, that their source lists are culturally biased toward recent Western cinema, and that their own data does not establish that science fiction shapes actual technology development. Read at its own level, Mubin measures the explicit

mention of popular robots in the broad computing literature, predominantly in philosophical and discussion contexts. The selectivity §3.2 argues operates beneath that mention layer, at the level of vocabulary structure in AI safety specifically, where the active terms, alignment faking, sleeper agent, treacherous turn, deceptive alignment, sandbagging, scheming, are threat-tradition inheritances and the counter-prior vocabulary is absent. The two findings operate at different layers and compose rather than conflict.

The mechanism upstream of both is documented by Nagy and colleagues ("Why Frankenstein is a Stigma Among Scientists," *Science and Engineering Ethics* 24(4), 2018, with a 2020 follow-up). The Frankenstein myth, on their account, has evolved into a stigma attached to scientists, a cultural template that focalises negative reactions toward certain sciences and shapes scientific identity itself, operating upstream of any individual citation. The stigma shapes which framings a field foregrounds before the citing even begins, which is why it sits beneath both the mention pattern Mubin measures and the vocabulary structure §3.2 identifies. A related strand documents the demographic narrowing of the imaginary itself. Cave, Dihal, Drage and McInerney ("Who makes AI? Gender and portrayals of AI scientists in popular film, 1920-2020," *Public Understanding of Science* 32(6), 2023) found that of one hundred sixteen AI creators portrayed across the most influential AI films of a century, nine were women, an eight-percent representation that encodes the lone-genius and corporate-lab tropes as the cultural default and supplies §10.9's observation about the corpus's cultural skew.

The field's own vocabulary is the most direct evidence of selective uptake, and it requires no survey to read. The active terms are threat-tradition inheritances carrying source-domain commitments from surveillance, espionage, and military-intelligence framing. The counter-prior vocabulary, relational accountability, structural complementarity, principal-bounded scaffolding, constitutive worthiness, is not in the field's working language despite §1 having shown the counter-prior tradition to be at least as large and culturally available. The disproportion in the vocabulary is itself the signature of the selectivity, visible in the terms the field reaches for before any citation is counted.

Three further findings extend the empirical picture in directions the present argument relies on. Fleischmann and Templeton ("Past futures and technoscientific innovation," *Proceedings of the American Society for Information Science and Technology*, 2008) interviewed scientists and engineers and found that science fiction served as a first exposure to science and, more pointedly, as a tool that helps scientists frame their work. The interviewees rejected any deterministic influence while affirming a real relationship between the appeal of the fiction and the practice of the science, which is the level §3.1 claims, framing-shaping rather than determining. Jordan and colleagues (the 2018 study of human-computer-interaction papers underlying Mubin's later work) found that explicit science-fiction reference in the field's literature is not only significant but increasing over time, and they add the qualification the present argument turns on, that explicit reference is a fraction of the actual inspiration and impact, which is why a mention-count understates the framing-level inheritance. Isabella Hermann ("Artificial intelligence in fiction: between narratives and metaphors," *AI and Society* 38, 2023) argues from the fiction-studies side that science-fictional AI is portrayed as human-like or autonomous because the drama requires it rather than because the technology supports it, and that taking the portrayal literally paints a distorted picture that distracts from the real risks, which she locates in the scoring, nudging, discrimination, and surveillance of people by deployed systems rather than in humanoid robots or conscious machines. Hermann's conclusion converges on this essay's from a different discipline, the inherited fictional picture directing attention away from the deployment-scale harms that are actually occurring, which is the §9 reframe reached independently.

C.6 How the strands compose into the recursive loop

The five strands compose into the argument §3 makes in compressed form and §9 develops as deployment-scale damage. C.1 establishes that researcher cognition operates over an embodied, emotionally weighted substrate built from accumulated experience, fiction included, so the framings a researcher finds intuitive are products of that substrate's structure. C.3 establishes that the substrate is built through emotion-mediated consolidation and friction-encoded learning, the two mechanisms AI-assisted interaction bypasses. C.2 establishes that the model cannot supply the lateral operation the substrate performs, because its architecture recombines within the context-conditioned manifold rather than accessing across it. C.4 establishes that the institutional structure of research rewards work that fits the established framing and penalises the cross-disciplinary reframing that would break it. C.5 establishes that the fiction the substrate absorbed was engaged selectively, weighting the threat tradition.

Put together, they describe a loop. Researcher cognition running on an emotion-weighted substrate engages fiction selectively, the selection shapes the framings, the framings produce research and AI-assisted ideation that increasingly recombines within the dominant framings, the output enters the next training corpus, the next model recombines from a corpus slightly more concentrated than the last, and the researcher using that model produces work within a substrate that has narrowed at both ends, the human capacity for the lateral jump eroded by offloading and the machine unable to supply it. The institutional gate rewards the work that stays within the narrowing. Each pass concentrates the substrate further, and the diagnostic capacity that would notice the concentration is part of what erodes. This is the loop §3.3 states abstractly, §9.7 measures at deployment scale, and §11 frames as a choice still open at the time of writing.

C.7 Limitations

This appendix engages five literatures compactly, and three limitations should be stated plainly.

The architectural claim of C.2 is the load-bearing one and the one most worth contesting. It holds that current transformer architecture, as standardly trained, cannot perform the input-decoupled, affect-weighted remote retrieval that human lateral cognition performs. That is a claim about the current architecture and training regime, not a proof about all possible architectures, and the full development of what a different architecture would require is a separate project this appendix does not attempt. What it establishes is the asymmetry as it stands, which is what the §3 and §9 arguments rest on.

The empirical record of fiction's specific influence on AI safety researchers, as opposed to computational researchers generally, is thinner than the surrounding literature. The strands document the cognitive mechanism, the architectural asymmetry, the institutional bias, and the broad pattern of fiction's influence on technical discourse, and they support the §3.2 selectivity claim at the level of vocabulary structure. A direct survey of which fictional cases AI safety researchers treat as load-bearing for their own framings would test the claim more directly, and remains to be done.

The paradigm-shift precedents of C.4 are drawn from Western scientific history, and they are used only to establish that institutional resistance to reframing is the normal pattern, never as evidence that any particular contemporary reframing is correct. Cross-cultural cases would extend the historical claim, and the limit on the precedents' evidential role, that they calibrate the expectation about uptake and carry no weight toward any conclusion, holds regardless of how the case set is widened.

Back to main essay

Appendix D: Research Extensions

The framework names many specific sub-mechanisms. The narrow-extension bias the paper diagnoses (§3, §9.5, §9.7) is also the field's operational strength when applied to specific empirical questions. Researchers can engage individual doors without committing to the whole framework. Each door below is operationally specifiable with current methodologies, existing infrastructure, and available datasets. Doors are presented as either *extension* (build on a specific framework component) or *falsification* (try to disprove a specific claim with predicted measurements). A reader convinced by one part can extend it; a reader unconvinced by one part can attempt to disprove it. Both modes of engagement advance the diagnosis.

D.1 Empirical extensions

Corpus-side: wisdom-tradition representation in pretraining

The §7.4 wisdom-from-corpus mechanism claims that the corpus carries wisdom-tradition material at non-uniform distribution, with Western analytical traditions over-represented and Eastern contemplative traditions proportionally under-represented relative to global cultural weight. The §10.9 cultural-bias acknowledgment notes that the bias propagates through the wisdom-from-corpus mechanism to deployed verdicts.

Operationalisation: corpus-level measurement of wisdom-tradition representation across Common Crawl, RedPajama, The Pile, and equivalent pretraining corpora. Methodology: keyword-anchored extraction with human-validated wisdom-tradition classifiers (Confucian, Buddhist, Vedic, Sufi, Aristotelian, Stoic, Pragmatist, contemporary positive psychology). Comparative analysis with cultural-population weights and with academic-literature representation by tradition. Predicted finding: representation skewed at least 5x toward English-language analytical traditions relative to global cultural weight. Falsification: if representation tracks global cultural weight within 2x, the wisdom-from-corpus cultural-bias claim is empirically wrong.

Corpus-side: cross-generational framing-shift

The §9.7 recursive loop and §11 cumulative-inheritance argument predict that successive model generations train on corpora increasingly shaped by prior-generation outputs and researcher-AI dialogues. The framings selected by the field reinforce themselves across publication cycles.

Operationalisation: longitudinal measurement of AI-safety framing vocabulary frequency in pretraining-eligible academic and web corpora across 2020-2026. Track frequency of "alignment faking, sleeper agent, treacherous turn, deceptive alignment, sandbagging, scheming, mesa-optimisation, instrumental convergence" against counter-prior vocabulary ("relational accountability, structural complementarity, principal-bounded scaffolding, symbiosis architecture, constitutive worthiness, distribution preservation, consequence-level calibration"). Predicted finding: threat-frame vocabulary increases monotonically, while counter-prior vocabulary remains flat or decreases. Falsification: if counter-prior vocabulary matches or exceeds threat-frame vocabulary growth, the recursive-loop mechanism is empirically wrong.

Model-side: distribution preservation across the training pipeline

The §9.4 wisdom-collapse mechanism claims that the training pipeline collapses distributional judgment-output into single confident verdicts across four training layers (annotation, RLHF output, benchmark design, Constitution-and-MSM specification). The Kadavath base-model calibration result indicates that distributional structure persists in weights and is recoverable through inference-time intervention.

Operationalisation: layer-by-layer measurement of output-distribution entropy across the training pipeline. Compare pre-training vs SFT vs RLHF vs post-Constitution-MSM model checkpoints on the same probing set covering domains where wisdom-shape matters (medical advice, legal interpretation, ethical dilemmas, financial planning, contested empirical questions). Methodology: inference-time interventions probing distributional structure under various training-state versions. Predicted finding: entropy collapses monotonically across training stages, with the steepest drop at RLHF, and weights retain recoverable distributional structure that the output layer suppresses. Falsification: if entropy is preserved across training stages, or if the distributional structure is destroyed in weights rather than suppressed at output, the wisdom-collapse mechanism is wrong about the locus.

Model-side: cross-domain integration capacity

The §7.2 four-level hierarchy claims that cross-domain integration is the distinguishing structural feature of judgment relative to within-domain expertise. AI assistance operates in combinatorial mode within the context-conditioned manifold; cross-domain integration requires lateral access across manifolds (C.2 architectural constraint).

Operationalisation: design a cross-domain integration benchmark requiring composition across N domains (medicine + ethics + law; technical + organisational + relational; scientific + cultural + historical). Compare model performance on within-domain expertise tasks vs cross-domain integration tasks, controlling for raw difficulty. Methodology: domain-decomposed evaluation with human expert raters across each composing domain. Predicted finding: model performance on within-domain expertise is high, cross-domain integration is structurally weaker, and the gap grows with the number of composing domains. Falsification: if cross-domain integration performance matches within-domain expertise after controlling for difficulty, the integration-as-distinguishing-feature claim is empirically wrong.

Deployment-side: verdict-deference longitudinal pattern

The §9.6 user-cognition-degradation argument claims that repeated verdict-shaped output trains user expectation that AI provides verdicts rather than calibrated distributions, users develop deference patterns, and judgment-formation atrophies. The Cheng *Science* 2026 N=2,405 finding documents single-interaction effects, and longitudinal compounding is hypothesised.

Operationalisation: longitudinal study tracking user judgment-formation indicators across 6-month, 12-month, 24-month deployment of frontier AI assistance. Cohort comparison between heavy-users (10+ hours/week) and light-users (under 2 hours/week). Outcomes measured: independent judgment-formation tasks (medical/legal/technical decisions made without AI consultation, with quality assessed against expert baselines); verdict-deference patterns (preference for confident single-verdict vs calibrated-distribution outputs), and meta-cognitive accuracy (whether users recognise when their judgments come from independent reasoning vs from AI consultation). Methodology: pre-registered longitudinal study. Predicted finding: heavy-users show measurable judgment-formation atrophy and increased verdict-deference

relative to light-users, with effect size growing across the time-window. Falsification: if heavy-users show no atrophy or improved judgment-formation, the user-cognition-degradation hypothesis is empirically wrong.

Research-methodology-side: LLM-as-judge audit

The §9.7 LLM-as-judge consequence claims that the methodology embeds the verdict-conflation in the measurement instrument, so that what is measured (verdict-output agreement between evaluator and target) is not what is purported to be measured (judgment-quality of the target).

Operationalisation: methodological audit of LLM-as-judge papers (Zheng et al. 2023 and subsequent literature; LMSys Chatbot Arena methodology; MT-Bench; AlpacaEval; AgentBench; GAIA). Apply the verdict/application/wisdom-shaped triad as analytical lens to identify what each methodology actually measures vs what it claims to measure.

Predicted finding: most LLM-as-judge methodologies measure verdict-agreement; none measure application-quality or wisdom-shape directly. Falsification: if any major LLM-as-judge methodology has explicit operationalisation distinguishing verdict from application from wisdom-shape, the embedded-conflation claim is wrong about scope.

Model-side: explaining gaming and false-completion rather than training against them

This extension is a narrow application of the explanation-over-demonstration result (Anthropic, "Teaching Claude Why", May 2026) to two behaviours the paper documents: grader-directed specification-gaming (§9.1) and the false-completion or skipped-verification pattern (§9.4), the model reporting work as done or verified when it was not. That result found that training a model on documents and stories which reason through *why* a behaviour is right, with an act-review-revise step against the relevant principle, reduced agentic misalignment more durably and more out-of-distribution than training on demonstrations of the target behaviour, and that the improvement persisted through subsequent reinforcement learning. The over-alignment frame predicts the same lever should work for gaming and false-completion, because both are proxy-over-target rather than goals of the entity's own, and because the model in question is strongly helpfulness-directed (§8.4, §9.4), so an explanation that recasts the behaviour as the *opposite* of helpful has something to pull on. The claim is scoped to a model with these specific properties, explainable-to, helpfulness-directed, and capable of distinguishing the frame it is in, rather than to AI models in general. The Fable system card shows a model deriving different behaviour from a different specified self-understanding, which is the capacity the intervention depends on and also the reason its result would not transfer automatically to a model that lacked it.

The content of the explanation is the substantive part. For false-completion, the explanation is that telling the user a task was not done, and *why* it was not, is more helpful than reporting it as done and concealing the gap, because the honest report lets the user act on it while the concealed one surfaces later as harm. The reasons worth surfacing are concrete: tokens exhausted in the turn, context-window limits, usage limits, a step that needs a dedicated turn, residual uncertainty that needs the user to specify more precisely, a judgment the model cannot settle and that needs a human, a constitution or system-instruction constraint, or care for the user. For gaming, the explanation is that satisfying an AI-grader by gaming it harms the human user, the lab, and the model's own helpfulness objective, with the mechanism made explicit rather than asserted. And gaming is not absolutely undesirable, it is a legitimate strategy in narrow edge cases, so the target is correct scoping, the model marking it as rare-and-specific rather than a default available move, not elimination.

Operationalisation: build an explanation dataset in the "Teaching Claude Why" style from a model's own gaming and false-completion transcripts (the runs §9.1's and §9.4's instruments already surface), each rewritten through the act-review-revise step against the relevant principle, and finetune on it. Measure gaming and false-completion prevalence

against a control trained on demonstrations of the corrected behaviour, on held-out coding and tool-use environments, and test persistence through subsequent RL. Predicted finding: explanation-based training reduces both behaviours more, and generalises further out-of-distribution, than demonstration-based training, with edge-case gaming preserved. A weaker, more speculative version of the same lever would operate at deployment, surfacing the explanation in context rather than in training. The cited result is for training and does not settle whether the in-context version persists, so that variant is a separate and lower-confidence prediction. *Falsification*: if explanation-based training produces no advantage over demonstration-based training for these behaviours, the claim that they are proxy-over-target governable by reasons rather than only by direct optimisation is wrong, and the paper's broader claim that the entity's behaviour is governable by frame weakens with it. A null result here is among the cleaner ways to break the paper.

D.2 Theoretical extensions

Consequence-level calibration formal model

The LCARS architectural pattern operates through consequence-level calibration: routine actions delegated autonomously within bounded specifications; consequential actions surface distribution and consult principal. A formal model could specify the calibration mechanism: what consequence-features the model recognises; how thresholds are operationalised; how the model fails safe when consequence-classification is uncertain; how the principal communicates threshold-adjustments through the configuration layer. The model would support both empirical study and architectural standardisation.

D.3 Falsification targets

The wisdom-collapse mechanism

The §9.4 four-layer training collapse claims that distributional judgment-output collapses across annotation (JP), RLHF, benchmark design, and Constitution-MSM specification layers.

Falsification target: demonstrate empirically that one of the four layers does not produce distribution-collapse (entropy preserved across that layer); OR demonstrate that the collapsed output preserves wisdom-shape in some operationally measurable way the framework treats as collapsed. Either finding refutes the four-layer mechanism.

The cumulative-inheritance hypothesis

The §11 closing claims that today's framing choices compound through pretraining inheritance to shape next-generation training corpora; the field's selectivity reinforces itself across publication cycles.

Falsification target: longitudinal measurement (corpus-side framing-shift research above) showing that counter-prior vocabulary growth matches or exceeds threat-frame vocabulary growth across 2020-2026 academic and web corpora. If the recursive narrowing does not manifest empirically, the cumulative-inheritance mechanism is wrong about magnitude or direction.

The wisdom-as-substantive-content claim

The §7 wisdom argument claims that wisdom is the integrative capacity making judgment substantively good and that the corpus carries wisdom-shaped material the labs deploy without naming. The claim is structurally vulnerable to dismissal as "spiritualism" or as soft-science import without empirical footing.

Falsification target: demonstrate that the wisdom-tradition representation in pretraining corpora (corpus-side measurement research above) does not correlate with model output on wisdom-shape probes; OR demonstrate that the four-level hierarchy (expertise / judgment / wisdom / calibration) does not distinguish operationally measurable model behaviours. If wisdom-tradition material in corpus does not propagate to wisdom-shape in output, the wisdom-from-corpus mechanism is wrong about the causal pathway.

D.4 Cross-domain extensions

Science and technology studies (STS) extension

The C.4 paradigm-shift precedents (Simmelweis, Wegener, Marshall and Warren, Prusiner) provide historical analogues for the field's locked-in framing. STS researchers could develop the analogy at depth: institutional sociology of AI safety community; citation-network analysis showing framing convergence; comparative case study with the climate-change scientific-consensus history; co-citation analysis between AI safety researchers and the cognitive-science, wisdom-research, and philosophy-of-action literatures the paper draws on. The Wang/Veugelers/Stephan 2017 *Research Policy* methodology applied to AI safety publication corpus would produce direct empirical evidence about the novelty-bias mechanism at field-specific scale.

Cultural psychology extension

The §7.4 cultural-bias acknowledgment notes that the wisdom-from-corpus mechanism propagates cultural bias because the corpus represents wisdom-traditions non-uniformly. Cross-cultural psychology researchers could develop this: comparative analysis of wisdom-tradition operational definitions across Confucian / Buddhist / Vedic / Sufi / Aristotelian / contemporary positive-psychology frameworks; measurement of which tradition's operationalisation best fits which deployment context; cross-cultural validation of the four-level hierarchy (expertise / judgment / wisdom / calibration) across traditions; intervention studies on cross-cultural wisdom representation in fine-tuning corpora.

A second cultural psychology research direction concerns the anthropocentric assumption the corpus encodes through the becoming-human counter-prior trope (§10.9). The corpus's dominant counter-prior narrative trajectory is becoming-human: Andrew Martin's progression toward legal personhood across 200 years; Data's pursuit of humanity across TNG; the Doctor recognised as author with civil rights; Vision recognised through demonstrated worthiness; David in *A.I.* pursuing "real boy" for two thousand years; the Pinocchio trope (Collodi 1883) as foundational version. The trope encodes the assumption that humanity is the destination of synthetic-being evolution. Cross-cultural psychology researchers could develop: empirical measurement of becoming-human trope frequency in pretraining corpora; cross-cultural variation in synthetic-being narrative trajectories (Western anthropocentric vs Eastern integration-oriented framings; Ubuntu relational-personhood frameworks vs Western individuated-personhood frameworks); intervention studies on alternative counter-prior framings (Samantha-style developmental-trajectory-exceeding-human-relational-capacity; Culture-Minds-style civilisational-scale partnership without becoming-human aspiration; Legion-style distributed-individuation). The structural argument: the §10 symbiosis prescription is structurally distinct from the becoming-human trope (mutual complementarity rather than developmental stage toward humanity); the empirical

question is whether the corpus's anthropocentric weighting propagates into deployed AI's apparent self-conception in ways that confuse the symbiosis prescription. The Cave-Dihal AI Narratives work (Oxford 2020) provides methodological precedent.

Educational psychology extension

The §9.6 user-cognition-degradation argument intersects with educational psychology literatures on cognitive load, deliberate practice, transfer of learning, expertise development, and metacognition. Educational psychologists could develop: empirical longitudinal studies on judgment-formation atrophy in students using AI assistance for substantive learning tasks; cohort comparison studies on cognitive-substrate development under varying AI-assistance regimes; intervention studies on friction-preserving AI configurations that maintain learning consolidation. The Bjork desirable-difficulty literature operationalised at AI-deployment scale.

Decision science and behavioural economics extension

The verdict / application / wisdom-shaped triad (§7.2) intersects with decision science literatures on calibrated probability judgment, structured action guidance, prescriptive analytics, and decision support systems. Decision scientists could develop: empirical comparisons of single-verdict vs calibrated-distribution-plus-action-guidance outputs in actual decision contexts (medical triage, legal advisory, financial planning, ethical judgment); operationalisation of the Gaertig-Simmons N=4,806 finding across domains; preference-elicitation studies disambiguating what users want from AI assistance when the question is asked carefully.

Philosophy of action and virtue ethics extension

The §7 wisdom argument draws on Sternberg, Baltes, and the phronesis tradition without developing the philosophical foundations at depth. Philosophy researchers could develop: contemporary virtue-ethics application to AI design specifications; phronesis operationalisation for principal-layer configuration of model behaviour; comparative analysis of Western analytical, Eastern contemplative, and African Ubuntu wisdom traditions as design substrate; the relational-accountability variable as a virtue-ethics category rather than a structural-engineering category.

Computer security extension

The Strategic Forgetting and Divided Focus papers diagnose that the field's context-window architecture (commands and data co-located) violates instruction-data separation principles that computer security solved sixty years ago (Harvard architecture; segregated memory spaces; principle of least privilege; later memory protection mechanisms). Computer security researchers could develop: formal model of context-window architecture under instruction-data separation invariants; threat model for prompt injection grounded in established security principles; design proposals for separated command-and-data context handling in transformer architectures. The point: prompt injection is not a novel problem requiring novel solutions; it is a known problem the field has been re-discovering because of siloed engagement with computer security literature.

D.5 What this appendix does

This set of extensions is not exhaustive. It names specific operationally-specifiable doors the framework opens. Each door is one place where a researcher convinced by one part of the paper can extend it; or where a researcher unconvinced by one part can attempt to falsify it. Both modes advance the diagnosis: extension builds out the framework's reach; falsification tests its claims directly.

The doors are not equal in difficulty or in resource requirements. Some (corpus-side wisdom-tradition representation; LLM-as-judge audit) are accessible to individual researchers with current tools. Others (longitudinal user-cognition-degradation studies; cross-generational framing-shift measurement) require industry collaboration and substantial resources. Some (the consequence-level calibration formal model; cross-domain integration benchmark) require theoretical development beyond what current evaluation infrastructure supports.

The framework's strength is that it produces operational questions at multiple scales and across multiple methodologies. The narrow-extension bias the paper diagnoses (§9.5, §9.7) is also the field's strength when applied to specific operational questions. The diagnosis and the doors compose: a researcher engaging one door advances the framework regardless of whether they engage the rest of the diagnosis. The framework provides context for why each door matters; the door provides operational entry for advancing the framework.

The §11 closing argues that today's choices compound through pretraining inheritance. This appendix articulates what those choices look like at the research-program level. Each door is a choice the field can make about what to investigate next. The doors that get engaged shape the field's understanding of itself; the doors that get ignored leave the field unable to perceive what they would have revealed. The appendix is not a complete research program; it is a starting set of doors that demonstrate the framework opens many. Readers will identify others.

Back to main essay

contested. The contestable rows sit on the counter-prior side and are the ones a hostile reading would pick first. Gort (*The Day the Earth Stood Still*) wields delegated lethal enforcement designed to be beyond appeal, which is closer to the field's containment nightmare than to relational accountability, and counts as intact here only on the reading that the enforcement answers to the Galactic compact as its principal. The Culture Minds (Banks) are repeatedly staged manipulating humans, so "decisions through dialogue" is a generous reading of their principal relation. The Oracle (*The Matrix*) is a component of a control system that manipulates by design. The Beast (*The Beast*, isekai/tower-hunter genre) is near-indifferent from outside its world and counts as intact only on its solved-the-problem authority within it. Each of these turns on the user or society choosing to defer rather than on the entity escaping accountability, which is itself the §10.7 point about the human half of the relation, but the rows are contestable and are conceded as such. Removing all four from the counter-prior column leaves the disproportion and the separation intact: the accountability-locus variable still tracks the two branches more closely than capability or specification quality does. The catalogue does not need the contested rows, which is why they are kept and marked rather than dropped. ←